




Fig. 1. **A taxonomy of robot-evaluation benchmarks for predictive embodied intelligence** (86 representative benchmarks of the 160 in Table ??). Each sub-family is one cell listing its benchmarks, each with a short statement of what it measures. Benchmarks are grouped by evaluation lane (§3.1–§3.4) and then by the capability they stress. The lanes run from closed-loop task-success suites to the *bridge* frontier (shaded), the only lane that turns prediction into executed action, which is the survey’s novelty axis. Full per-benchmark detail: Tables ??, ??.