

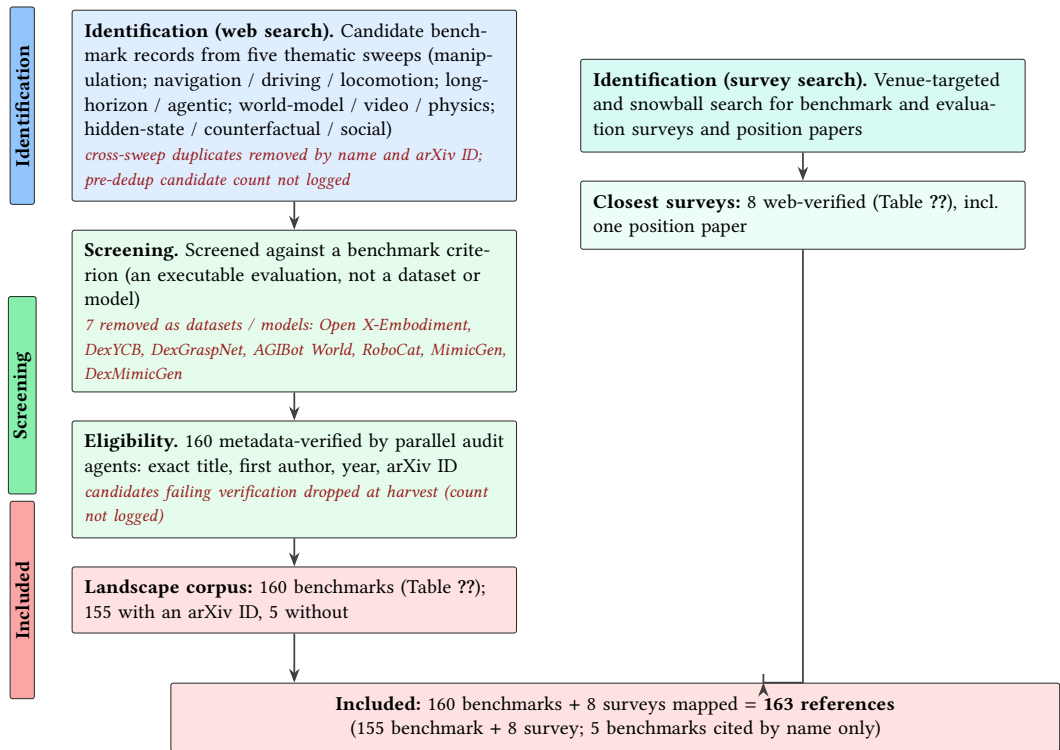


Fig. 1. **Corpus construction** in the PRISMA 2020 style [?]. The left column is the web-search stream behind the landscape corpus (Table ??): five thematic sweeps, de-duplicated by name and arXiv ID, screened to drop datasets and models, and metadata-verified by parallel audit agents. The right column is the venue-targeted and snowball search behind the closest surveys (Table ??). Both streams feed the included set. Candidates whose metadata could not be verified were discarded at harvest; their count was not logged, which we state rather than estimate.