



CLOSED-LOOP task-success suites



OPEN-LOOP world-model & video-generation evaluation



PREDICTION-TO-ACTION bridges

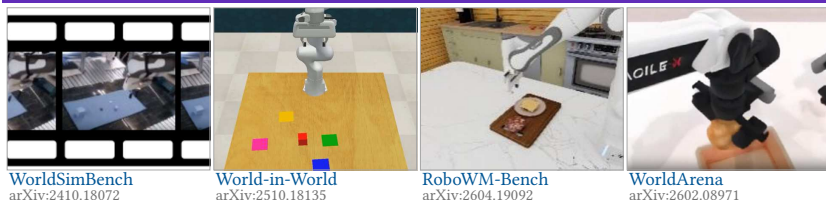


Fig. 1. Benchmarks across the surveyed corpus. Representative subject figures from the catalogue, grouped by evaluation lane rather than embodiment: *closed-loop* task-success suites, *open-loop* world-model and video-generation evaluation, and the *prediction-to-action* bridges. Each panel is a live hyperlink to the source paper on arXiv and is annotated with its arXiv id (per-panel figure provenance in Table 1); the figure is vector with photographs embedded at full resolution, so it can be zoomed without pixelation. Only figures that depict the benchmark’s task or scene are shown; results and leaderboard figures appear in Figure 2. Benchmarks with no locatable subject figure are omitted, not fabricated. Images © their respective authors, reproduced for scholarly review.