



Results reported across the surveyed corpus

vector, zoomable, per-panel arXiv links


Leaderboard and model-comparison charts




World-in-World
arXiv:2510.18135



EvalCrafter
arXiv:2310.11440




VideoScore
arXiv:2406.15252




WorldModelBench
arXiv:2502.20694


Per-dimension radar & multi-metric




VBench
arXiv:2311.17982



EWMBench
arXiv:2505.09694




WorldArena
arXiv:2602.08971




EvalCrafter
arXiv:2310.11440


Distributions & composition




EvalCrafter
arXiv:2310.11440



EvalCrafter
arXiv:2310.11440



VBench
arXiv:2311.17982



VideoPhy
arXiv:2406.03520

Fig. 2. Results reported across the surveyed corpus. Representative results figures from the catalogue, grouped by type of evidence: *leaderboard and model-comparison charts*, *per-dimension radar and multi-metric plots*, and *distributions and composition*. Each panel is a live hyperlink to the source paper on arXiv and is annotated with its arXiv id (per-panel provenance in Table 2); the figure is vector with plots embedded at full resolution. Only results/leaderboard/metric figures are shown; task and scene figures appear in Figure 1. As a benchmark survey the reported evidence is dominated by rankings, radars and distributions rather than training curves, and is sparser than the subject figures, so this gallery is smaller; benchmarks that publish no results figure are omitted, not fabricated. A benchmark may appear more than once when it reports several evidence types. Images © their respective authors, reproduced for scholarly review.