




Fig. 1. **The evaluation loop that isolates the advantage of prediction.** A benchmark that can answer “does world modelling help a robot act” must run a direct Vision-Language-Action policy *and* a world-model policy (predict then plan) under one closed loop, then report the closed-loop gain (δ) sliced by capability (β). Today’s suites break this loop in one of two ways: model-agnostic policy suites host a single policy and never build the contrast, while world-model suites score the prediction open-loop and never execute it (§??).