

Construct		Counts (a benchmark has it if and only if...)	Does not count
Open-loop (open)	eval	it scores prediction or generation <i>quality</i> on fixed inputs without executing the prediction in a feedback loop: video fidelity (FVD), physics or rollout consistency, or frame-level accuracy	any protocol where an action is executed and the environment responds; offline scoring of a policy’s success
Closed-loop (closed)	eval	success is measured by running a policy <i>in</i> the environment with state feedback and scoring task outcome: success rate, sub-goal completion, or reward under interaction	scoring a generated video or a static prediction; a metric computed without stepping the environment
Bridge (bridge)		a prediction (imagined rollout, video, or latent plan) is <i>converted into executed actions</i> and scored on the resulting task success	a benchmark that stops at prediction quality; one that never turns a prediction into a control signal
Native VLA-vs-WM contrast (γ)		the benchmark <i>itself</i> runs both a direct Vision-Language-Action policy and a world-model policy under one protocol and reports the head-to-head	a model-agnostic suite that merely <i>hosts</i> either family; a suite evaluated with only one family in the original paper
Advantage-of- prediction (δ)		a metric quantifies the closed-loop <i>gain</i> of a predictive policy over a matched direct-VLA baseline on the same tasks (assessed, not adopted)	absolute success with no matched baseline; open-loop quality scores; a ranking that never isolates the predictive component

Table 1. Operational definitions used to place each benchmark on the evaluation axes (§??). They resolve the ambiguous boundary cases: hosting a policy is not the same as *building* a VLA-vs-world-model contrast (γ), and an absolute success number is not an advantage-of-prediction measurement (δ) without a matched baseline. open open-loop, closed closed-loop, bridge bridge.