

	Policy suites [10]	Embodied nav / social [14]	World-model eval [8]	Bridges [18]
Loop closure	✓ closed-loop; steps the environment [11]	✓ closed-loop; long-horizon rollouts [7]	✗ open-loop; scores generation, no action [12]	~ converts a prediction into an executed action [16]
Task horizon	~ short reactive skills; long by chaining [10]	✓ long-horizon household and dialog [13]	~ short clips; temporal but not task-level [5]	~ a single executed rollout [6]
Capability breadth	~ manipulation + transfer; hidden-state rare [15]	✓ navigation, social, some hidden-state [4]	~ physical / temporal; counterfactual barely [2]	✗ narrow; executability only [17]
Model-family contrast	✗ model-agnostic host of any policy [1]	~ memory-vs-Markovian in one suite [3]	✗ world-model-only scores [8]	~ WM as policy substrate; no VLA pole [18]
Reproducibility	~ sim; some real-vs-sim transfer [9]	~ sim-heavy; a sim-to-real gap [4]	✓ fixed offline sets; deterministic scoring [19]	✗ new, still-unstandardised protocols [17]
Structural blind spot	hosts a world-model policy but never builds the VLA contrast	capability cuts exist, but no per-capability family ablation	visual quality is not task success; no behavioural utility	reaches action, but omits the per-capability VLA head-to-head

Table 1. **Evaluation-lane comparison matrix:** the four benchmark lanes on six evaluation axes (loop closure, task horizon, capability breadth, model-family contrast, reproducibility, and characteristic blind spot). Marks grade each lane on the axis (✓ favourable, ~ mixed, ✗ weak), with an anchoring benchmark per cell. No lane closes the loop, slices capability, *and* builds a native VLA-vs-world-model contrast at once; the bridge lane comes closest yet still omits the per-capability VLA head-to-head, which is the empty cell this survey names (§??).