


Fig. 1. **Results reported across the surveyed corpus.** Representative results figures from the catalogue, grouped by type of evidence: *leaderboard and model-comparison charts, per-dimension radar and multi-metric plots, and distributions and composition*. Each panel is a live hyperlink to the source paper on arXiv and is annotated with its arXiv id (per-panel provenance in Table ??); the figure is vector with plots embedded at full resolution. Only results/leaderboard/metric figures are shown; task and scene figures appear in Figure ??. As a benchmark survey the reported evidence is dominated by rankings, radars and distributions rather than training curves, and is sparser than the subject figures, so this gallery is smaller; benchmarks that publish no results figure are omitted, not fabricated. A benchmark may appear more than once when it reports several evidence types. Images © their respective authors, reproduced for scholarly review.