




Fig. 1. **Predict or Act?** The survey’s organising question, posed with real benchmarks. *Closed-loop* (left): policy and embodied suites score task success by running a policy in the environment. *Open-loop* (right): world-model and video-generation benchmarks score prediction or generation quality with no action. The centre axis is the question neither pole answers, whether prediction earns a closed-loop advantage, built explicitly by only 11 of 160 benchmarks (Table ??). Every tile is a representative figure from the benchmark’s paper and links to it; per-tile provenance is in Table 1. Images © their respective authors, reproduced for scholarly review.

Benchmark	arXiv	Source	Benchmark	arXiv	Source
Closed-loop (task-success suites)			Open-loop (world-model & video eval)		
LIBERO	2306.03310	Fig. 1 (p.2)	WorldModelBench	2502.20694	rep. figure (pp.1–3)
CALVIN	2112.03227	rep. figure (pp.1–3)	Physics-IQ	2501.09038	rep. figure (pp.1–3)
Meta-World	1910.10897	rep. figure (pp.1–3)	WorldScore	2504.00983	rep. figure (pp.1–3)
ManiSkill2	2302.04659	rep. figure (pp.1–3)	EWMBench	2505.09694	rep. figure (pp.1–3)
RoboCasa	2406.02523	rep. figure (pp.1–3)	EVA-Bench	2410.15461	rep. figure (pp.1–3)
Habitat 3.0	2310.13724	rep. figure (pp.1–3)	VideoPhy	2406.03520	rep. figure (pp.1–3)
ALFRED	1912.01734	rep. figure (pp.1–3)	Physion++	2306.15668	rep. figure (pp.1–3)
GemBench	2410.01345	rep. figure (pp.1–3)	Physion	2106.08261	Fig. 1 (p.3)
VLABench	2412.18194	Fig. 1 Overview (p.1)	EvalCrafter	2310.11440	rep. figure (pp.1–3)

Table 1. **Per-tile provenance for Figure 1.** Every tile is a representative figure cropped from the benchmark’s own arXiv paper (page/figure noted; “rep. figure” marks the paper’s teaser-region figure on pp.1–3), reproduced for scholarly review and hyperlinked to its source. Images © their respective authors.

REFERENCES