




Fig. 1. **Profile of the 160 catalogue benchmarks** (Table ??), re-counted directly from the catalogue. (a) publication year, showing the 2023–2024 surge; (b) evaluation lane (embodied navigation dominates, then policy suites); (c) evaluation mode (closed-loop task success dominates, open-loop prediction is second, only four bridges reach action); (d) whether the benchmark builds a native Vision-Language-Action vs world-model contrast, the survey’s organising question: **138 of 160 are model-agnostic**, and only **11 (7%)** build the contrast explicitly. Counts are the exact column tallies of Table ??.




Fig. 2. **The world-model-evaluation wave is recent.** Per-year counts of closed-loop task-success benchmarks (green) and open-loop world-model / video-generation benchmarks (magenta) in the catalogue. Closed-loop suites dominate through 2024; in **2025 open-loop evaluation overtakes them** (13 vs 6) as the world-model-generation literature arrives. The four prediction-to-action bridges are newer still, all appearing in 2024 or later. Years are the benchmarks’ first-release years (2026 is partial and omitted); counts are the tallies of Table ??.

Temporary page!

$\LaTeX$ was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away, because $\LaTeX$ now knows how many pages to expect for this document.