




Fig. 1. **The evaluation loop that isolates the advantage of prediction.** A benchmark that can answer “does world modelling help a robot act” must run a direct Vision-Language-Action policy *and* a world-model policy (predict then plan) under one closed loop, then report the closed-loop gain (δ) sliced by capability (β). Today’s suites break this loop in one of two ways: model-agnostic policy suites host a single policy and never build the contrast, while world-model suites score the prediction open-loop and never execute it (§??).




Fig. 2. **An actionable evaluation protocol for the future agenda** (Table ??). The panels are *schematic*: they show the axes and the *target* shape of each measurement, not measured results. (a) a world-model policy should beat a matched direct-VLA baseline in closed-loop success, and the per-capability gap is the advantage of prediction (δ), on LIBERO / CALVIN-style suites run with both families; (b) counterfactual accuracy is near chance today and should rise, on WorldPrediction and occlusion splits; (c) a world model’s open-loop quality far exceeds its closed-loop task success, and that gap must be reported, on bridge suites; (d) the fraction of capabilities with a native VLA-vs-world-model contrast is far below target across the landscape.