

Benchmark	arXiv	Panel source (subject scene)
Closed-loop task-success suites		
ALFRED	1912.01734	sim task scene (rep. figure, pp.1–3)
BEHAVIOR-1K	2403.09227	household sim scene (rep. figure, pp.1–3)
CALVIN	2112.03227	tabletop manipulation scene (rep. figure, pp.1–3)
GemBench	2410.01345	manipulation scene (rep. figure, pp.1–3)
ManiSkill2	2302.04659	multi-task sim scenes (rep. figure, pp.1–3)
Meta-World	1910.10897	task-suite thumbnail grid (rep. figure, pp.1–3)
RoboCasa	2406.02523	kitchen sim scene (rep. figure, pp.1–3)
LIBERO	2306.03310	manipulation task scene (embedded raster, pp.1–8)
VLABench	2412.18194	mahjong-tile manipulation scene (embedded raster, pp.1–8)
Habitat 3.0	2310.13724	human–robot collaboration scene (embedded raster, pp.1–8)
Open-loop world-model & video-generation evaluation		
EVA-Bench	2410.15461	embodied-anticipation scene (rep. figure, pp.1–3)
EWMBench	2505.09694	generated-video sample (rep. figure, pp.1–3)
Physics-IQ	2501.09038	physical-realism sample (rep. figure, pp.1–3)
Physion	2106.08261	Fig. 1 example-scenario frames (p.3)
Physion++	2306.15668	physical-prediction scene (rep. figure, pp.1–3)
VideoPhy	2406.03520	generated-video sample (rep. figure, pp.1–3)
WorldModelBench	2502.20694	generated-video sample (rep. figure, pp.1–3)
WorldScore	2504.00983	generated world sample (rep. figure, pp.1–3)
VBench-2.0	2503.21755	generated-video sample (embedded raster, pp.1–8)
WorldPrediction	2506.04363	procedural-action video frame (embedded raster, pp.1–8)
Prediction-to-action bridges		
WorldSimBench	2410.18072	video-prediction filmstrip, interactive-eval figure (embedded raster, pp.1–8)
World-in-World	2510.18135	closed-loop manipulation scene (embedded raster, pp.1–8)
RoboWM-Bench	2604.19092	manipulation rollout scene (embedded raster, pp.1–8)
WorldArena	2602.08971	dual-arm manipulation scene (embedded raster, pp.1–8)

Table 1. **Per-panel provenance for Figure 1.** Every panel is a real subject/scene figure cropped from the benchmark’s own arXiv paper and hyperlinked to its source. Exact figure numbers are given where verified from the paper (e.g. Physion Fig. 1); otherwise the teaser-region basis is noted (“rep. figure, pp.1–3” for tiles reused from the collage pool; “embedded raster, pp.1–8” for the nine gap tiles hand-picked from per-paper candidate contact sheets after a subject/results/neither classification pass). No figure number is invented. Images © their respective authors, reproduced for scholarly review.




Fig. 1. **Benchmarks across the surveyed corpus.** Representative subject figures from the catalogue, grouped by evaluation lane rather than embodiment: *closed-loop* task-success suites, *open-loop* world-model and video-generation evaluation, and the *prediction-to-action* bridges. Each panel is a live hyperlink to the source paper on arXiv and is annotated with its arXiv id (per-panel figure provenance in Table 1); the figure is vector with photographs embedded at full resolution, so it can be zoomed without pixelation. Only figures that depict the benchmark’s task or scene are shown; results and leaderboard figures appear in Figure ?? . Benchmarks with no locatable subject figure are omitted, not fabricated. Images © their respective authors, reproduced for scholarly review.