



Confidence estimated after the fact, and confidence trained into the model


Estimated after the fact: read-off, post-hoc maps, prompts, extra inference (S0 to S3)




Stengel-Eskin and Van Durme 2023
Fig. 1, p. 6



Lu 2026
Fig. 1, p. 2



Marashian et al. 2026
Fig. 1, p. 1




Doula et al. 2025
Fig. 1, p. 3


Trained into the model: a calibration term in the training loss or reward (S4)



Yang et al. 2026b
Fig. 1, p. 3



Guo et al. 2026
Fig. 1, p. 2



Gao et al. 2026
Fig. 4, p. 4

Fig. 1. **Confidence estimated after the fact, and confidence trained into the model.** Panels are drawn from Figures 8 and 9 under the same licence rule: one per stage for S0, S1, S2, S3, and one per training family for S4 (T1, T3, T4; T2, T5 not shown). Each label links to the paper; Table B1 gives the figure, page and licence.