




Fig. 11. A measurable agenda, partly measured. Decision ECE is the expected calibration error of the probability a model places on its chosen option (Table 4). Panels (a) and (d) come from one community leaderboard, the Jev Decision Index [multimodalart 2026] (release 0.2.1, a single snapshot of 2026-09-27; not peer reviewed): skill is its chance-normalised score over 38 decision benchmarks (0 to 100), ECE is computed on a 1-in-6 sample of 32 of them, and all values are point estimates without uncertainty intervals. (a) Reliability measured on the Jev Decision Index (Table 2) for Jev and the three open systems with the highest skill; each point is a confidence bin, and bins holding under 1% of answers are omitted. (b) Selective risk and (c) escalation regret are drawn as target shapes, not measured results: purple marks the target, grey dashed a typical baseline, and the named testbeds are in Table 4. (d) Decision ECE against skill for Jev and the 67 open systems on the board; lower right is better. The board records no cost per decision, so skill is used as the horizontal axis.