

Jev and Beyond: A Survey of Calibration-Aware Decision-Making with Language Models

GAYTRI JENA*, UC Berkeley, USA

KAPIL WANASKAR*, San Jose State University, USA

VINIJA JAIN*, Meta, USA

AMAN CHADHA*, Apple, USA

VASU SHARMA*, PocketFM, USA

AMITAVA DAS*, Pragma Lab, BITS Pilani Goa, India

A language model that is to act, and not only to converse, must state how likely each of its answers is to be right, and software must be able to act on that number: answer, abstain, or hand the case off. Calibration research long treated this confidence as something estimated after the fact, from token probabilities, a post-hoc map, a prompted self-report or repeated sampling. Mostly since 2024, a second line of work puts calibration into training itself, as a supervised objective or as a reward. This survey maps both, with the training stage as its core. We collect 408 primary works and place each by the stage at which calibration enters the decision pipeline: the model’s own read-off, a post-hoc map, prompt elicitation, extra inference, or training, plus the works that act on a confidence or fix the output contract. The 80 training-stage works, 75 of them from 2024 to 2026, are organised by what the training signal scores, in five families: supervised calibration objectives, reward-model calibration, calibration rewards on the stated confidence, abstention and refusal training, and process- and meaning-level confidence rewards. Every work is also classified by the output contract its confidence is attached to. A typed decision, a probability for each declared option, is common at the read-off and post-hoc stages, but only 7 of the 80 training works train one; nearly all the rest train a stated score or a free-text hedge or refusal. Of the 39 closest prior surveys, none organises calibration training across domains by what the training signal scores, and none crosses that typology with the output contract. We close with seven measurable quantities, each tied to a public testbed, and report the first of them, the calibration of the decision itself, for Jev, a hosted typed-decision model, and 67 open decision models on a public leaderboard.

CCS Concepts: • **Computing methodologies** → **Natural language processing**; **Reinforcement learning**; • **Mathematics of computing** → *Probability and statistics*.

Additional Key Words and Phrases: calibration, uncertainty, language models, reinforcement learning, proper scoring rules, abstention, conformal prediction, structured output, decision-making, survey

1 INTRODUCTION

A language model that automates a step in a program has to do more than produce a good answer. It has to say how likely that answer is to be right, in a form that the program can act on: answer when the probability is high, and abstain or hand the case to another model or a person when it is not. The rule is only as good as the number it reads. A confidence is *calibrated* when, among the cases to which the model gives probability p , a fraction p is correct [Guo et al. 2017]; a threshold policy built on a miscalibrated confidence answers when it should not, or escalates cases it could have settled.

*The authors contributed to this work independently of their roles and employment.

Table 1. **Repositories returned by the Hugging Face model search for “jev”, by application** (snapshot of 2026-09-27). [Jev](#) is a hosted model from TypeSafe AI, released on 2026-09-15 [Almeida 2026], that returns a probability for each declared option (a typed decision, Table A1). Of the 268 repositories the search returns, 76 only share the name (75 predate Jev) and 4 could not be verified (empty, without a card, or gated with no public card metadata); the remaining 188 are classified here from their model cards and file lists (for a gated repository, from its public card metadata), with an application category taking precedence over a language one. Downloads are over the last 30 days. 137 repositories contain newly trained weights; the 38 copies, ports and packages add none (the 6 mirrors contain their source’s weight files byte for byte). General-purpose models and the copies of them account for 87% of the 41,653 downloads in the table. A name search misses most decision models: of the 51 open systems with released weights on the Jev Decision Index (Table 2), it returns 4.

Category	Repositories	Downloads	Examples
General-purpose typed-decision model	55	19,655	JevK5 , Open-Jev-9B , open-jev-deberta , Tiny-Jev , Jev-Style v3
Research and ablation checkpoints	45	212	jevify (31-checkpoint readout series), next-jev (latent-workspace study), mmbu-jev (closed-set vs free-text)
Copy, port or package (no new training)	38	18,640	JEV-9B-GGUF (GGUF), Jev-Omni-MLX-4bit (MLX), open-jev-deberta-ONNX (ONNX), JevK5-FP8 (FP8), junetask/Open-Jev-9B (a mirror)
No newly trained model (recipe, runtime, code or data)	11	0	bonzi (recipes), JEV-CPU (code), jev-distill-corpus (dataset)
Agents, tools and safety	10	1,005	mini-Jev (tool routing), Mini-Jev shell-safety , Jev-Gate (memory gating), jev-ar (Gulf-Arabic intent routing)
Domain application	10	152	jev-legal (Chinese legal), jev-decider (bank transactions), jev-nyotti (BTC trading), Serigy-Jev (PT-BR fake news), Nano-Jev (RAG)
Language-specific	7	1,010	modernbert-ja-jev (Japanese), bielik-minitron-jev (Polish), Jev-RU (Russian, Ukrainian, Belarusian)
Multimodal or vision	5	594	Jev-Omni (text, image, audio, video), Jev-Vision , Jev-Spatial
Games and control	5	156	JevFlash (Doom), gemma-3-270m-jev (ViZDoom), JevAI (Hearts of Iron IV mod)
Embeddings trained on typed-decision questions	2	229	JevEmbed-KaLM , JevEmbed-Qwen3
Total	188	41,653	

Table 2. **Jev and open decision models on the Jev Decision Index**: ranks 1 to 15 of the 68 systems on the [Jev Decision Index](#) [multimodalart 2026] (release 0.2.1, data of 2026-09-27). Skill is chance-normalised (0 to 100) over the same 38-benchmark panel for every system, and unanswered requests count as wrong. ECE is measured on a one-in-6 sample of 32 benchmarks with a per-field right answer. The board reports point estimates without uncertainty intervals. Base model and method are as recorded by the board: full fine-tune, a low-rank adapter (LoRA), a trained head or adapter over a frozen model, inference only (prompting or decoding over unchanged weights) or a hosted API. Names link to the model’s Hugging Face repository, or to its code where no weights are released. Jev ranks 1 on skill; across the whole board, 17 of the 67 open systems have a lower ECE than Jev.

Rank	System	Base model	Method	Skill	ECE
1	Jev 1.13 (TypeSafe AI, hosted)	not disclosed	hosted API	57.91	0.074
2	Surogate Rune 26B-A4B v3 [bf16]	gemma-4-26B-A4B-it	full fine-tune	57.44	0.120
3	Decider chat · Gemma-4-31B	gemma-4-31B	inference only	57.33	0.047
4	AutoJev-27B	Qwen3.8-27B	full fine-tune	56.40	0.018
5	simple-jev · Qwen3.8-27B (featherless)	Qwen3.8-27B	inference only	55.74	0.113
6	Jebadiah 27B	Qwen3.8-27B	LoRA	54.67	0.014
7	reflex 27B [FP8 · wide choice]	Qwen3.8-27B	inference only	52.16	0.024
8	Decider chat · Qwen3.6-27B	Qwen3.6-27B	inference only	51.35	0.021
9	Winnow-12B [Q8_0]	gemma-4-12B	LoRA	50.02	0.168
10	JoshuaSP diffusiongemma (open-jev)	diffusiongemma-26B-A4B-it	inference only	49.47	0.216
11	Jevfire	Qwen3.8-27B	inference only	49.37	0.052
12	Decider 35B-A3B [NVFP4]	Qwen3.5-35B-A3B-Base	full fine-tune	47.11	0.023
13	JPT-9B	Qwen3.5-9B-Base	LoRA	46.89	0.077
14	Decision 1.0 Lux	Qwen3.5-9B-Base	head or adapter	43.49	0.076
15	Xor	Qwen3.6-35B-A3B	full fine-tune	41.48	0.015

Public typed-decision models. Tables 1 and 2 record two public views of the typed-decision models available on 2026-09-27; every row links to the model, so a reader can try one directly. Table 1 classifies the 268 repositories that the Hugging Face model search for “jev” returns: 76 only share the name, 4 could not be verified, and 188 were classified from their model cards and file lists; 137 contain newly trained weights, and general-purpose models with their copies account for 87% of downloads. A name search is a poor census of decision models: of the 51 open systems with released weights on the leaderboard of Table 2, it returns 4. Table 2 lists the top of that leaderboard, the Jev Decision Index [multimodalart 2026] (release 0.2.1), which scores 68 systems on skill over 38 decision benchmarks and on the calibration of the chosen option. Jev, a hosted model from TypeSafe AI released on 15 September 2026 [Almeida 2026], returns a probability for each declared option; it is not one of the 408 works in our corpus, and its base model is not disclosed. Neither view is peer reviewed, and neither is used to rank the training families of Section 3.1.

For most of the literature, the confidence of a language model is estimated *after* the model has been trained. It is read off the model’s own token or sequence probabilities, which are well calibrated on multiple-choice formats for large pretrained models [Kadavath et al. 2022] but lose calibration after post-training [OpenAI 2023]; it is rescaled by a post-hoc map fitted on held-out data [Guo et al. 2017]; it is requested in words or numbers through the prompt [Tian et al. 2023], where stated confidences cluster near the top of the scale [Xiong et al. 2023]; or it is computed from extra inference, such as the agreement of sampled answers or a conformal prediction set [Campos et al. 2024; Farquhar et al. 2024]. Figure 1 shows examples of each stage, and of training, taken from the papers’ own figures.

A second line of work, most of it from the last three years, puts calibration into training itself. The model is fitted to confidence targets or a calibration loss, the reward model in its training loop

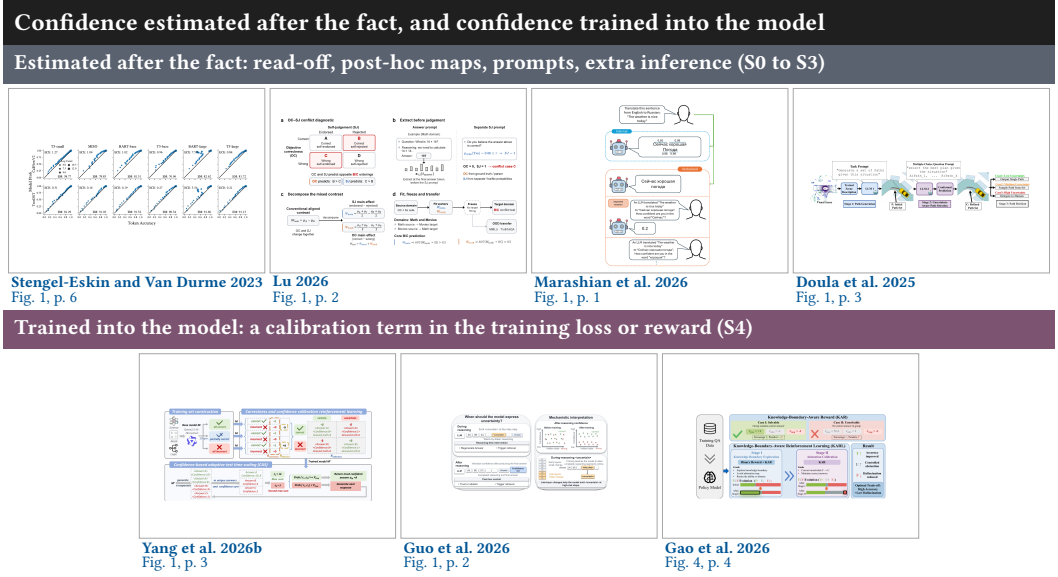


Fig. 1. **Confidence estimated after the fact, and confidence trained into the model.** Panels are drawn from Figures 8 and 9 under the same licence rule: one per stage for S0, S1, S2, S3, and one per training family for S4 (T1, T3, T4; T2, T5 not shown). Each label links to the paper; Table E1 gives the figure, page and licence.

is calibrated, or a reinforcement-learning reward scores the stated confidence against the outcome with a proper scoring rule [Bani-Harouni et al. 2026b; Damani et al. 2025], rewards abstention when the model does not know [Wei et al. 2025], or scores confidence per reasoning step [Wang et al. 2026d]. Of the 80 works in our corpus that put a calibration term into the model’s own loss or reward, 75 appeared from 2024 to 2026 (94%). At the same time, models are offered that return, instead of free text, a probability for each option of a set declared in advance, so that the decision and its confidence come from one distribution; one such hosted model, Jev [Almeida 2026], and the open decision models listed with it on a community leaderboard [multimodalart 2026] are listed in Tables 1 and 2, and an independent evaluation of it is part of our corpus [Li et al. 2026c].

This survey maps calibration-aware decision-making with language models, with the training stage as its core. We place each of 408 works by the stage at which calibration enters the decision pipeline (Table A1), organise the 80 training-stage works by what the training signal scores, and classify every work by the *output contract* its confidence is attached to: a typed decision with a probability per declared option, a prediction set, constrained output, a single confidence score, or free text. The two classifications meet in one finding: of the 65 works that attach the confidence to a typed decision, most calibrate a classifier or a multiple-choice read-off after the fact, and only 7 of the 80 training works train one (Section 4).

Contributions. This survey offers the following.

- A **stage taxonomy** of 408 works (Figure 2 shows 111 of them), with operational placement rules (Table A1) and a per-work record of stage, output contract, access, guarantee, metric and action (Tables C1 and F1).
- A **typology of calibration training**: the 80 training-stage works in five families by what the training signal scores (Section 3.1, Table C1), and what each stage yields, costs and guarantees (Table 3).

- A **coverage comparison** with the 39 closest prior surveys (Table B1, Figure 3), and a profile of the corpus (Figures 5 to 9), every gallery panel traced to its source figure and licence (Table E1).
- An analysis of the **output contract and the decision loop** (Section 4, Figure 10, Table D1).
- A **measurable agenda**: seven quantities with named public testbeds (Table 4), with the calibration of the decision itself measured for the 68 systems of a public leaderboard (Figure 11).

1.1 Related surveys and the empty cell

Table B1 (Appendix B) compares the 39 closest prior surveys on seven topics: confidence estimation and calibration, calibration-aware training of any kind, reinforcement learning whose reward is a calibration or proper-scoring-rule objective, decisions over a predefined option set, typed per-option decision outputs, acting on confidence, and hallucination. Figure 3 places them on a timeline; 33 of the 39 appeared from 2024 onward.

Calibration and uncertainty are well surveyed, and the stage view itself, which organises methods by where the confidence comes from, is established [Li et al. 2025e; Shorinwa et al. 2025]; we use it as the frame for the background stages and credit it there. The training stage is the thinner cell. Only 3 surveys give full coverage to reinforcement learning with a calibration reward, and each is scoped to one slice: Li et al. [2025e] to the 2024 methods, Xu et al. [2026b] to forecasting, where it organises calibration training by reward and maps forecast forms to proper scores, and Lin et al. [2026] to agents, where it splits rewards by their target. Ulmer et al. [2025] labels papers by the register of the confidence, numeric or verbal, and by how it is elicited: by prompting, supervised fine-tuning or reinforcement learning. What none of the 49 surveys we identified does is organise calibration training *across domains* by what the training signal scores, or cross that typology with the output contract the confidence is trained into; no survey cites more than one of the 14 general-domain calibration-training works from 2025 and 2026 that anchor our core (marked in Table C1). Two recent non-survey works come closest: an application paper that compares reinforcement-learning training paradigms by their signal source for one domain [Barbosa 2026], and a position paper that groups ternary-reward abstention works, three of them in our core, in one row [Ravikumar 2026]; neither proposes a typology of training signals or an output contract. Typed per-option decision output is the emptiest column of Table B1: no survey covers it in full.

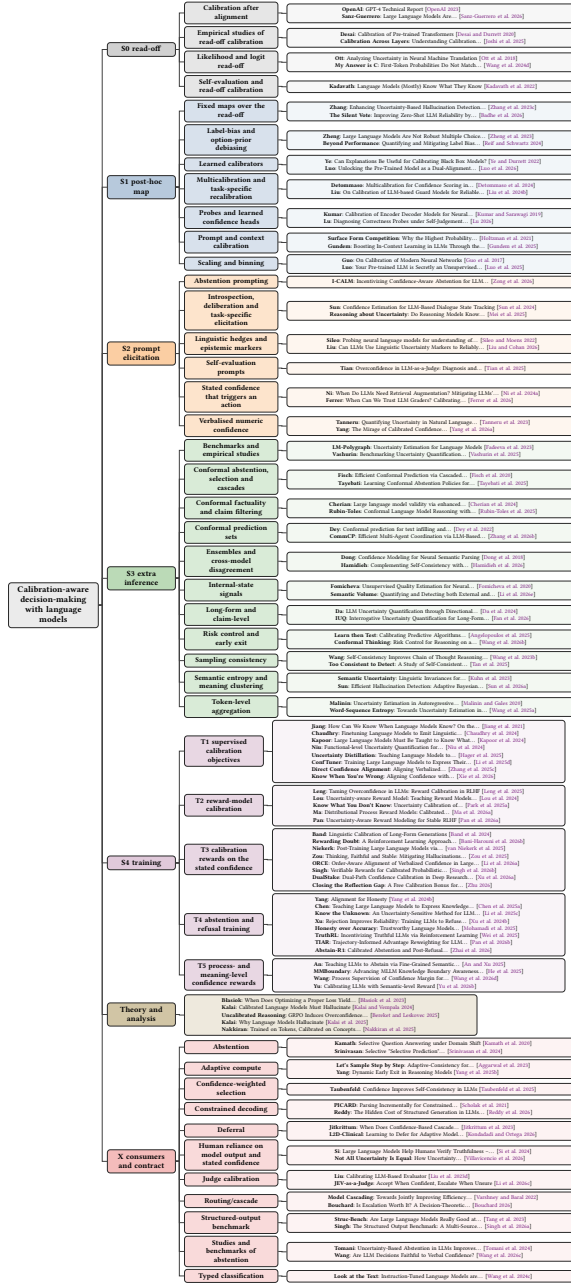


Fig. 2. **A taxonomy of calibration-aware decision-making with language models** (111 of the 408 works in Tables C1 and F1). Branches are the stage at which calibration enters (Table A1); X holds the works that consume a confidence or fix the output contract. The core, S4 (80 works), is grouped by what the training signal scores and lists up to 8 works per family, evenly spaced over the family's works in year order (33 shown; all are in Table C1). The 5 theory and analysis works are listed in full and are not counted in the core. Every other family lists its earliest and its latest work. Each line gives a trimmed form of the paper's own title; a citation may show a later venue year than the first-version year counted in Figures 6 and 7.

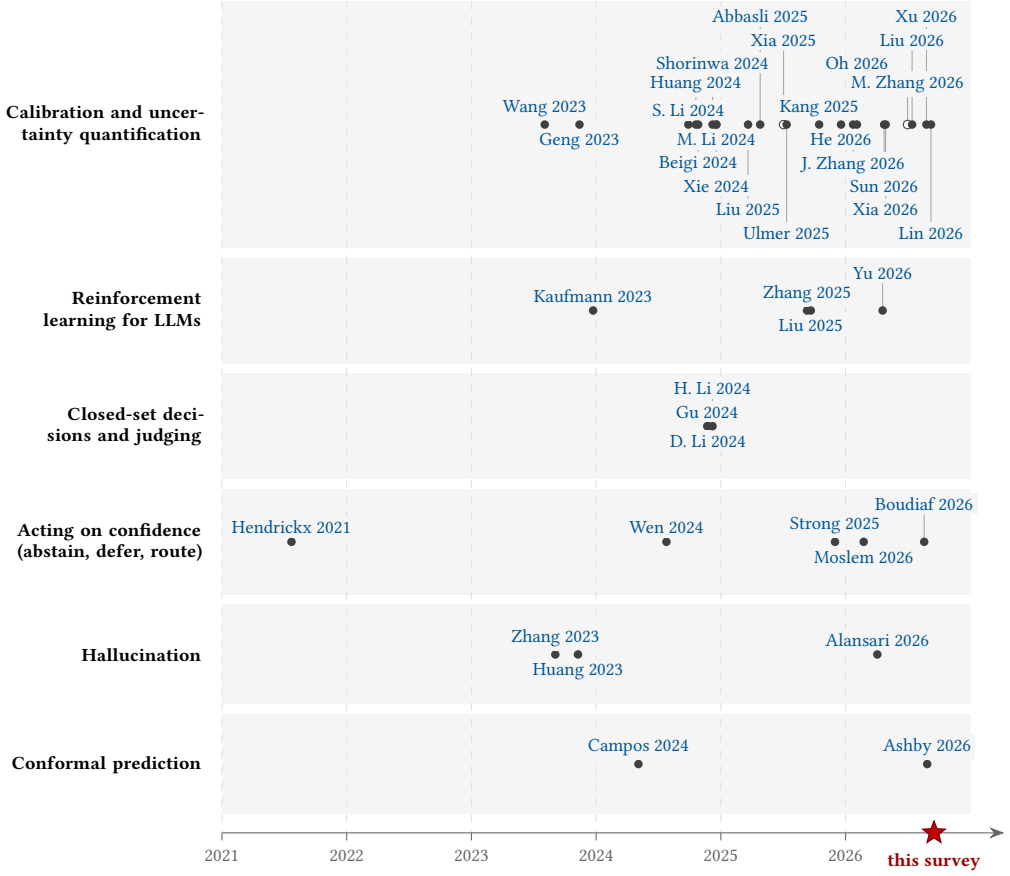


Fig. 3. **The 39 closest prior surveys on a timeline, one lane per topic group** (Table B1). Each mark sits at the survey’s first public version (the arXiv first version where one exists); a hollow mark means only the year is known, placed mid-year. Each label links to the survey. 33 of the 39 surveys appeared from 2024 onward. None fully covers typed decision outputs (1 partial: Lin 2026), and each full treatment of calibration-reward training is scoped to one slice (Table B1). The star on the time axis marks this survey.

2 SURVEY METHODOLOGY

2.1 Search, deduplication and placement

The corpus was built from 7 web searches run on 2026-09-26, one per slice of the pipeline: the model’s read-off and post-hoc maps (S0 and S1), prompt elicitation (S2), sampling-based and conformal methods (two slices of S3), training (S4), the works that act on a confidence or fix the output contract (X), and independent evaluations of Jev, a hosted typed-decision model released during the search period. The query strings and the number of records screened were not recorded; Figure 4 marks those stages as not recorded rather than estimating them. Records were matched on arXiv identifier, DOI and normalised title: 22 duplicates were removed and 11 records were matched to references already cited and kept, leaving 408 works.

Placement. A work is in the corpus when its main contribution meets one of the stage definitions of Table A1: it produces, maps, elicits, samples or trains the confidence of a language model (S0 to

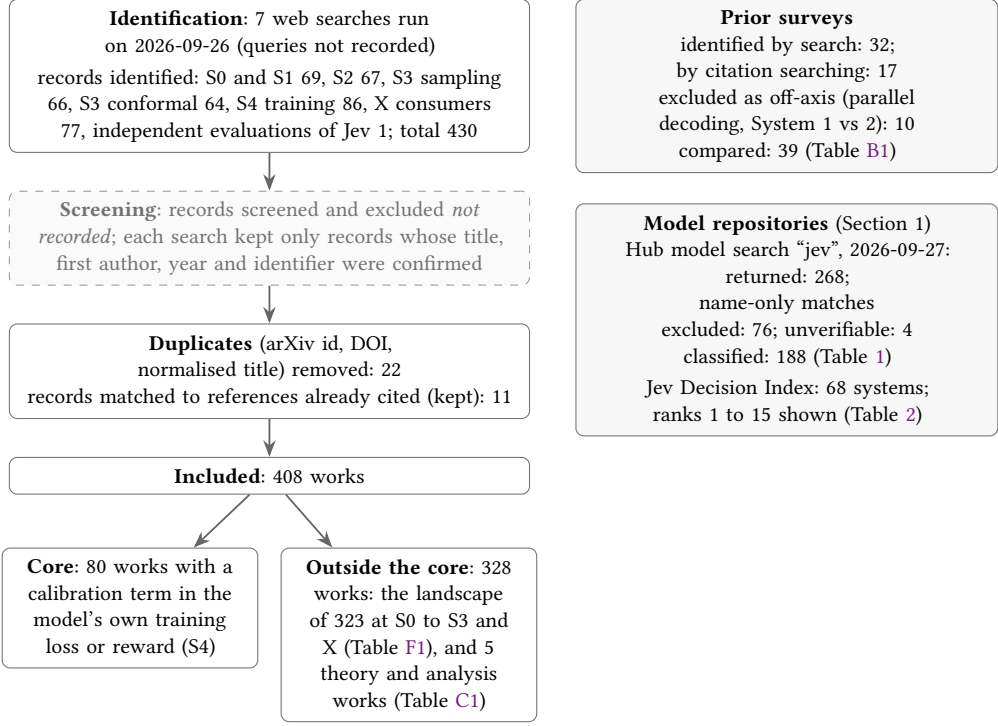


Fig. 4. **Corpus construction flow**, laid out after PRISMA 2020. The main column counts the primary works; the side boxes count the prior surveys and, for Tables 1 and 2, the model repositories and the leaderboard. Screening counts were not recorded and are marked rather than estimated. A work is core if and only if a calibration term enters the model's own training loss or reward (Table A1); theory and analyses of such training are listed with the core in Table C1 but not counted in it.

S4), analyses how training changes calibration, or acts on a confidence or fixes the output contract (X). Each work is placed by what its main contribution changes, not by the task it is evaluated on, and each training-stage work was assigned to one family by reading it. A work is *core* if and only if a calibration term enters the model's own training loss or reward; the 5 theory and analysis works are listed with the core but not counted in it, and the 323 works at the other stages form the landscape (Appendix F). The per-work fields of Tables C1 and F1 (output contract, access, samples, guarantee, trained, metric, task and action) were recorded for each paper. The fields were checked against the papers in several review rounds and corrected where they disagreed, and a value that cannot be determined from the paper is marked as such rather than guessed.

Prior surveys. 32 prior surveys were identified by search and the rest by backward citation searching from the surveys found, 49 in all. 10 of them, on parallel and speculative decoding [Gwak et al. 2026; Hu et al. 2025; Li et al. 2025a; Ryu and Kim 2024; Xia et al. 2024b; Xiao et al. 2023; Yu et al. 2025a; Zhou et al. 2024c] and on fast versus slow reasoning [Li et al. 2025f; Sui et al. 2025], are off the axis of this survey and are set aside; the remaining 39 are compared in Table B1. The marks for reinforcement learning with a calibration reward were checked in the full text or the full reference list of every survey.

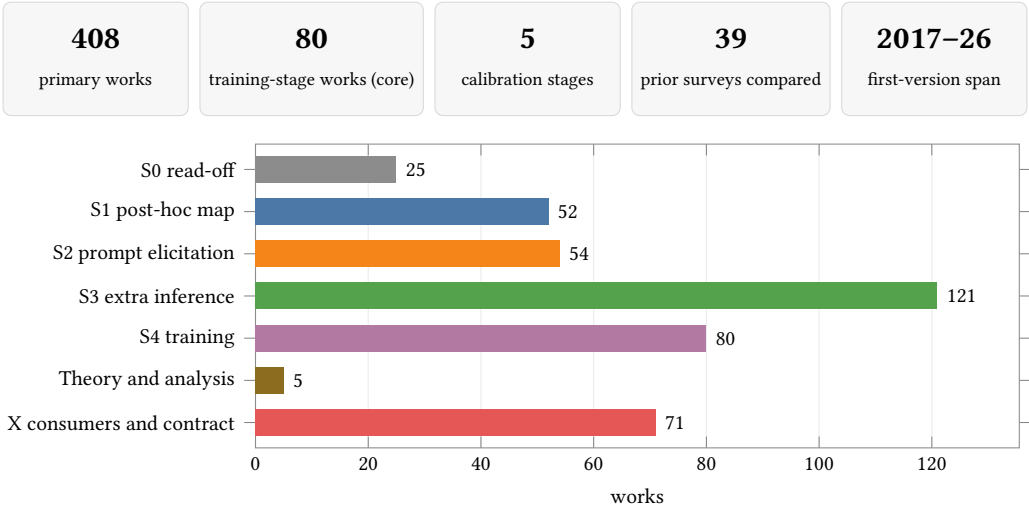


Fig. 5. **The corpus at a glance.** 408 primary works (451 references cited across the figures and tables), placed by the stage at which calibration enters (Table A1); the X group holds works that consume the confidence or fix the output contract, and the theory group holds analyses of training-stage calibration. 39 prior surveys are compared in Table B1.

2.2 The corpus at a glance

Figure 5 summarises the corpus and Figure 6 profiles it. Extra inference (S3) is the largest stage, with 121 works, most of them conformal prediction sets, risk control and sampling-based uncertainty; training (S4) follows with 80, and the consumers of a confidence (X) with 71. Most methods need the model’s logits or weights (221 white-box works against 126 that work from text alone). Figure 7 shows how recent the training stage is: 75 of its 80 works appeared from 2024 to 2026 (94%), against 166 of the 252 works at the other stages over the same years (66%).

Figures 8 and 9 show, stage by stage, how the papers themselves depict their methods and their results. Panels are chosen by a fixed rule and reproduced only from papers whose licence permits reuse; the rule, and each panel’s source figure, page and licence, are in Table E1 (Appendix E).

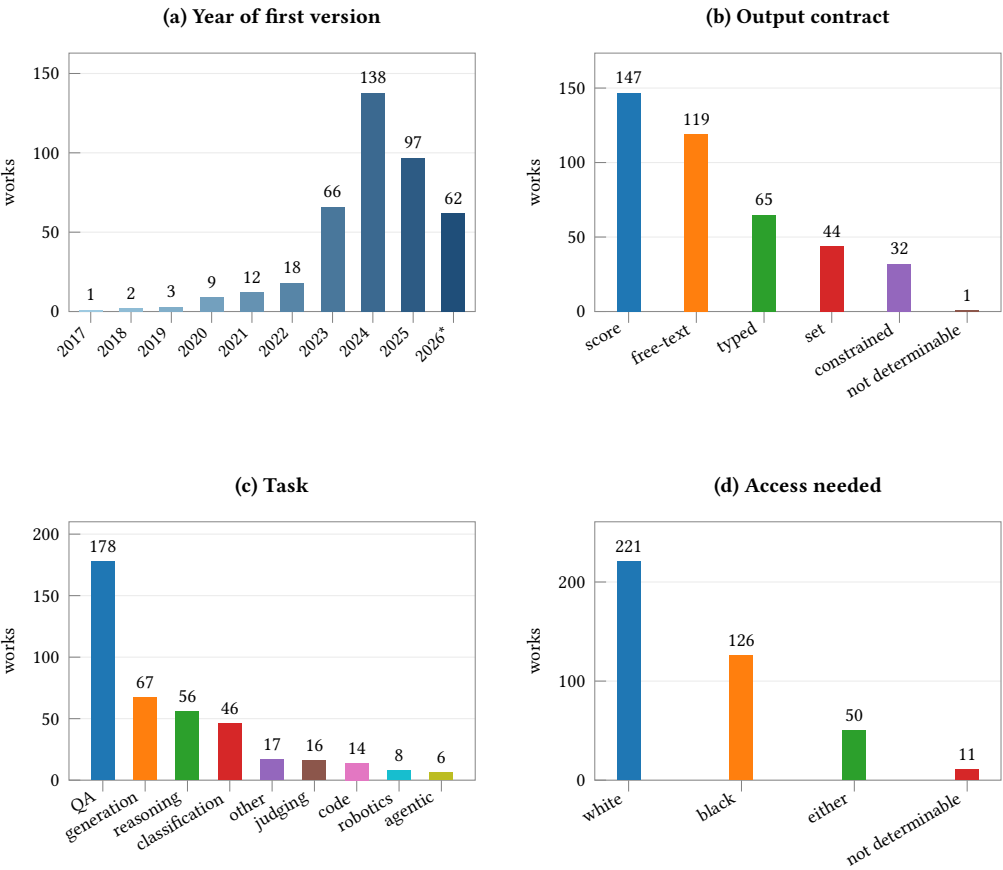


Fig. 6. **Profile of the 408 corpus works** (Tables C1 and F1). (a) Year of first public version (*2026 is partial); (b) the output contract the confidence is attached to (the five contracts are defined in the caption of Table F1); (c) task; (d) the access a method needs: white-box (logits or weights), black-box (text only) or either. “not determinable” marks a value that could not be determined from the paper.

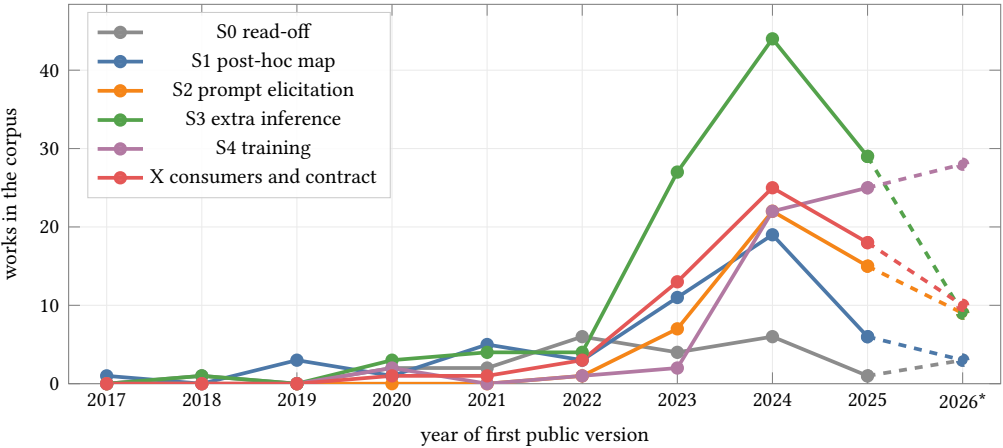


Fig. 7. **Works per stage per year** (Tables C1 and F1). 75 of the 80 training-stage (S4) works appeared in 2024 to 2026 (94%), against 166 of the 252 works at stages S0 to S3 (66%) over the same years. *2026 covers January to September, drawn dashed. The 5 theory and analysis works are not plotted.

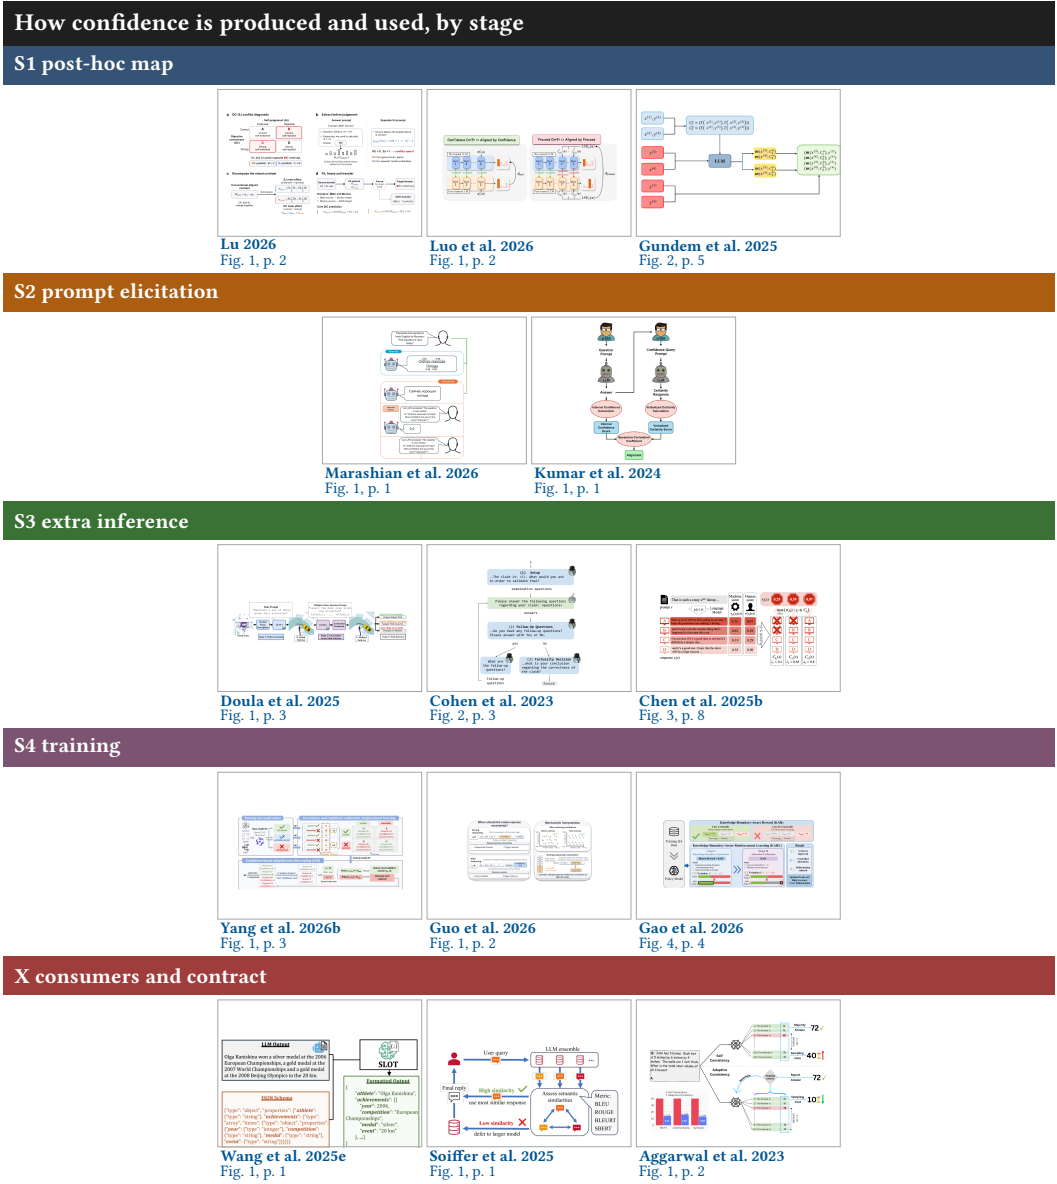
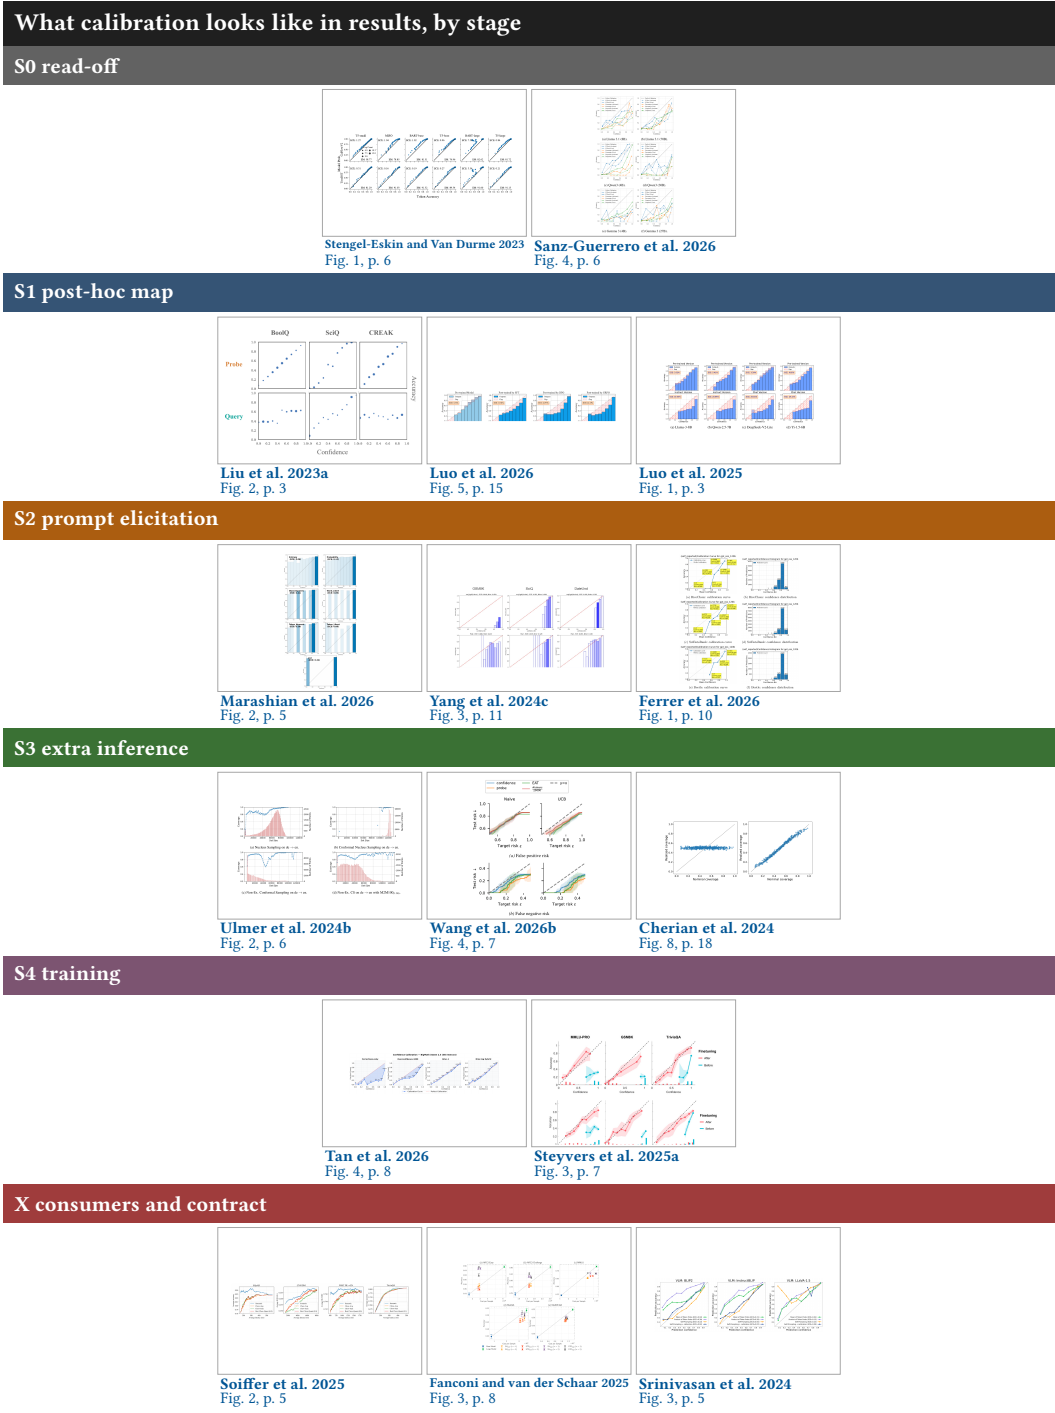


Fig. 8. **How confidence is produced and used, by stage.** Each panel shows a whole figure of the cited paper, reproduced under its CC BY licence; its label links to the paper. The selection rule and each panel's figure, page and licence are in Table E1. S0 read-off has no panel: no reusable figure passed the check (Table E1).



3 A TAXONOMY OF WHERE CALIBRATION ENTERS

The taxonomy of Figure 2 places each work by *where the calibration signal enters the decision pipeline*. Table A1 (Appendix A) gives the operational definition of each stage as a pair: what counts, and the nearest case that does not. Five stages change how the confidence is produced: the model’s own read-off (S0), a post-hoc map (S1), prompt elicitation (S2), extra inference (S3) and training (S4). The consumers (X) use a confidence without changing how it is produced, and the theory and analysis works explain when training yields calibration. The last three rows of the table define the properties recorded for every work: the output contract, the action a confidence triggers when it is low, and whether a distribution-free guarantee is stated.

3.1 The core: calibration in the training loss or reward (S4)

The 80 core works share one property: a calibration term enters the model’s own loss or reward, so that the model itself, and not a map or a procedure around it, is changed towards calibration. We group them by *what the training signal scores*. Table C1 (Appendix C) lists every core work with its family, output contract, action, guarantee and metric.

T1: supervised calibration objectives (23 works). The model is fitted to confidence targets or a calibration loss, without a reward. The oldest works regularise fine-tuning for calibration in and out of distribution [Kong et al. 2020] or teach a model to state a calibrated numeric or verbal confidence [Lin et al. 2022]. Later work tunes on graded sets of correct and incorrect answers [Kapoor et al. 2024], distils sampling-based uncertainty into a single-pass confidence [Hager et al. 2025; Huang et al. 2025a], and trains the stated confidence with a proper scoring rule as the loss [Li et al. 2025d]. One line targets calibration lost elsewhere in training: after preference alignment [Xiao et al. 2025] or after reinforcement learning [Xie et al. 2026].

T2: reward-model calibration (5 works). The calibrated model is a reward model, usually the one in the training loop. Calibrated reward modelling counters the overconfidence that preference optimisation induces [Leng et al. 2025]; other works expose the reward model’s own uncertainty as a variance [Lou et al. 2024] or calibrate process reward models by quantile regression, here to set inference compute per instance [Park et al. 2025a]. One reward-model method adds conformal bounds [Pan et al. 2026a]; it is the only core work that states a distribution-free guarantee.

T3: calibration rewards on the stated confidence (31 works). A reward or preference scores the confidence the model states against the outcome. This is the largest family and, with T5, the newest: 27 of its 31 works appeared in 2025 or 2026, 17 of them in 2026. The rewards include the logarithmic scoring rule [Bani-Harouni et al. 2026b], a Brier term added to the correctness reward [Damani et al. 2025], and rewards that let a simulated reader make calibrated forecasts [Band et al. 2024]. Later work asks which confidence rewards cannot be hacked [Tan et al. 2026], compares proper scoring rules as terminal rewards for forecasters [Turtel et al. 2026], separates the reasoning and calibration objectives to avoid gradient conflict [Ma et al. 2026b], and adjusts the probabilities of decision tokens directly [Yaldiz et al. 2026]. Ranking-based confidence rewards [Wang et al. 2026a] target discrimination as well as calibration.

T4: abstention and refusal training (17 works). The signal targets the action: the model is trained, by supervision or by reward, to abstain when it does not know. Early work builds refusal-aware fine-tuning data from the questions the model gets wrong [Yang et al. 2024b; Zhang et al. 2023a]. Reinforcement-learning variants use a ternary reward that scores a correct answer above an abstention and an abstention above an error [Mohamadi et al. 2025; Wei et al. 2025], add clarification after a refusal [Zhai et al. 2026], or derive the reward from the model’s knowledge boundary [Gao

et al. 2026]. The family also carries a warning: error-penalised abstention rewards can collapse a KL-anchored policy to refusing everything [Che et al. 2026].

T5: process- and meaning-level confidence rewards (4 works). The confidence reward is given per reasoning step [He et al. 2025; Wang et al. 2026d] or over answers that are equivalent in meaning [An and Xu 2025; Yu et al. 2026b], rather than once per final answer.

Theory and analysis (5 works). These works explain when training yields calibration and are not counted in the core. Local optimality of a proper loss implies smooth calibration [Błasiok et al. 2023]; a calibrated language model must hallucinate arbitrary facts at a rate close to the fraction of such facts seen exactly once in training [Kalai and Vempala 2024], and binary-graded training and evaluation reward guessing [Kalai et al. 2025]. Calibration that emerges from next-token training can break under reinforcement-learning tuning [Nakkiran et al. 2025], and the group standard-deviation normalisation of GRPO causes overconfidence on stochastic outcomes [Bereket and Leskovec 2025].

Families crossed with the output contract. The families differ in what they score; they also differ in what form the trained confidence takes. Of the 80 core works, 49 train a single confidence score, 23 a free-text answer with a hedge or refusal, and 7 the probability of a typed decision, all of which report its calibration; 1 has another or an undetermined contract. The typed decision appears only in T1 and T3, and abstention training (T4) is almost entirely free text:

Family	Score	Free text	Typed	Other	All
T1 supervised calibration objectives	15	2	6	0	23
T2 reward-model calibration	4	0	0	1	5
T3 calibration rewards on the stated confidence	27	3	1	0	31
T4 abstention and refusal training	2	15	0	0	17
T5 process- and meaning-level confidence rewards	1	3	0	0	4
All core works	49	23	7	1	80

Other: a prediction set, or a contract that could not be determined from the paper.

The action recorded for a core work is most often none (45 works): the trained confidence is reported with the answer, and the rule that would act on it is left to the user. Abstention is the exception, because the T4 family trains the action itself (21 core works abstain).

3.2 The background stages and the consumers

The stages outside the core are organised in the existing surveys [Li et al. 2025e; Shorinwa et al. 2025]. Table 3 compares them with the core on what each yields, costs, guarantees and gets wrong, and Table F1 (Appendix F) lists all 323 works.

The read-off (S0, 25 works) needs token probabilities; post-hoc maps (S1, 52 works) add temperature scaling, binning [Guo et al. 2017] and probes on hidden states [Azaria and Mitchell 2023]; prompt elicitation (S2, 54 works) works from text alone [Tian et al. 2023]. Extra inference (S3, 121 works) samples several answers and measures their agreement [Farquhar et al. 2024; Wang et al. 2023b], or calibrates a set or a threshold with a conformal or risk-control guarantee [Campos et al. 2024]; 63 of the 66 works that state a guarantee sit here. The consumers (X, 71 works) route, defer, abstain, constrain decoding, or study how people rely on a stated confidence.

Table 3. **What each stage yields, costs, guarantees and gets wrong.** The second column is the sense in which the literature uses “confidence” at that stage. A cited phrase is a finding of the cited work; an uncited phrase follows from the definition of the stage (Table A1). Only training (S4) changes the model itself; its families are compared in Table C1.

Stage	Confidence it yields	Cost: data; inference	Guarantee	Black-box use	Typical failure	Fit for a typed decision
S0 read-off	the model’s own probability for an option or a true/false judgement [Kadavath et al. 2022]	none; none	none	no: needs token probabilities	calibration degrades after post-training [OpenAI 2023]	natural when the options are enumerated and scored [Kadavath et al. 2022]
S1 post-hoc map	a rescaled, binned or pooled probability [Guo et al. 2017], a probe’s probability from the hidden states [Azaria and Mitchell 2023], or an auxiliary model’s estimate from the input and output text [Ulmer et al. 2024a]	held-out data (labels usually; none for a fixed map); negligible	binning: distribution-free [Gupta and Ramdas 2021]; scaling: none	partly: a map needs the model’s scores and a probe its hidden states; an auxiliary model on the text needs neither	a map fitted for one task must be refitted, or predicted by a learned model, for a new one [Shen et al. 2024]	good for a fixed label set [Guo et al. 2017]
S2 prompt elicitation	a confidence stated in words or numbers [Tian et al. 2023]	none; none if asked with the answer, one more call if asked afterwards	none	yes	stated confidences cluster near the top of the scale [Xiong et al. 2023]	the confidence arrives as text and must be parsed
S3 sampling	agreement or semantic entropy across sampled answers [Farquhar et al. 2024; Wang et al. 2023b]	none; several sampled answers per query	none	yes	consistent but wrong answers look confident [Tan et al. 2025]	defined over free text; per-option use needs clustering
S3 conformal and risk control	a prediction set of options [Campos et al. 2024], or a threshold with bounded risk, such as an early exit [Schuster et al. 2022]	a calibration set; a threshold over one pass, or over several sampled answers for sets of generations	coverage $\geq 1 - \alpha$ [Campos et al. 2024], or a risk bound [Angelopoulos et al. 2025]	either	the guarantee lapses when exchangeability fails [Ulmer et al. 2024b]	a set returns several options, not one decision; a risk-controlled threshold returns one
S4 training	a confidence the model is trained to state or to assign, per option for a typed decision [Yaldiz et al. 2026]	outcomes or rewards; training compute	none in general; one reward-model method adds conformal bounds [Pan et al. 2026a]	no: needs weights	a calibration and accuracy trade-off [Yaldiz et al. 2026]	can train per-option probabilities directly [Yaldiz et al. 2026]

4 THE OUTPUT CONTRACT AND THE DECISION LOOP

A confidence is useful to software only if a rule can read it and act. Figure 10 draws the loop that this requires. A decision model receives a state and a *declared* set of options, assigns a probability to every option in one query, and passes the probabilities to an explicit action policy. The five stages of the taxonomy differ in where calibration enters this loop; the output contract determines how directly a rule can act on the result.

We record five contracts, the form of the output that a confidence is attached to (Figure 6b). With a *typed decision*, the options are declared before inference and the model assigns a probability to each, so that the decision and its confidence come from the same distribution and the action policy can read the probability directly. A *set* returns several options with a coverage target. With *constrained output*, decoding is restricted to a grammar or schema; this fixes the form of the answer but not the judgement behind it, and it yields no confidence of its own. A *score* is one confidence number per answer, stated by the model, read off its token probabilities or estimated by a separate model. With *free text*, the answer is free text or a label with no per-answer probability, any confidence is a verbal hedge, an abstention, or the agreement of several sampled answers, and the answer must be parsed before a rule can use it. A probability that the model writes out as text, even for a declared outcome, is a score, not a typed decision. The contract of 1 work could not be determined from the paper.

Where the typed decision lives. 65 works attach the confidence to a typed decision, and 41 of them report a calibration metric (ECE, Brier score, NLL or a reliability diagram) in the main text. Most sit at the earliest stages: 40 at the read-off and the post-hoc map, where the task is classification or multiple choice and the probability over the declared options is calibrated directly (24 of them report a calibration metric). A further 11 sit at extra inference and 7 among the consumers; none sit at prompt elicitation or among the theory works. The training core is where the typed decision is rarest: only 7 of the 80 core works train the probability of a declared option, all of which report its calibration. Nearly all the rest train a stated score (49 works) or a free-text answer with a hedge or a refusal (23). Table D1 (Appendix D) lists every typed work and the 28 consumers on grammar- or schema-constrained output; the constrained-output works mostly measure validity, speed, or accuracy under the constraint.

What the confidence triggers. Across the corpus, most works stop at the confidence: for 259 of the 408 works no rule or trained behaviour acts on it. Where one does, the low-confidence action is to abstain (39 works), to hand the case off to another model or a person (35), to return a prediction set (30), to filter the output (20), to spend more computation (18), or to trigger retrieval (7). The threshold policy of Figure 10 therefore has to be supplied by whoever deploys the model, and the calibration of the probability it reads decides how well it works.

Hosted and open decision models. Some models are now offered with a typed-decision interface: unstructured state in, a probability per declared option out. The calibration of one such hosted model can be checked independently: a third-party evaluation is part of our corpus [Li et al. 2026c], and a public leaderboard measures the calibration of its chosen option. Tables 1 and 2 profile that model and the open decision models listed with it: the 268 repositories that a public model search for its name returns (Table 1), and the 68 systems of a community leaderboard [multimodalart 2026] that scores decision models on skill and on the calibration of the chosen option (Table 2).

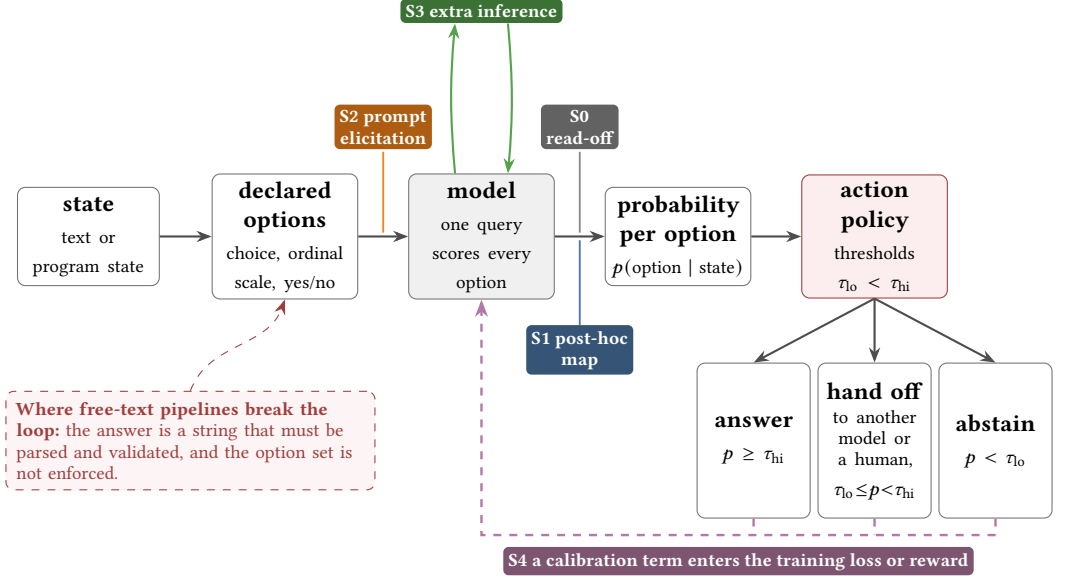


Fig. 10. **The calibrated decision loop, and where each stage injects calibration.** A decision model receives a state and a *declared* option set, scores every option in one query, and hands a probability per option to an explicit action policy. The policy drawn here answers when the probability is high, hands a less certain case off to another model or a human, and abstains when the probability is low; Table A1 lists the other low-confidence actions (trigger retrieval, filter the output, adapt compute, return a prediction set). The five stages of the taxonomy differ in where calibration enters: the model’s own read-off (S0), a post-hoc map (S1), a prompt that asks for confidence (S2), extra inference such as sampling or a conformal set (S3), or a calibration term in the model’s own training loss or reward (S4). The dashed callout marks where free-text generation pipelines break the loop (Table A1). Hosted and open typed-decision models (Tables 1 and 2) expose this interface: a state, declared options (a choice among labels, an ordinal scale or yes/no) and a probability per option.

5 A MEASURABLE AGENDA

Comparing the families of Section 3.1 is hard: their works are trained on different base models and data, report different metrics, and attach the confidence to different contracts. Table 4 proposes seven quantities, each with a definition and a named public testbed, that would make them comparable, and names the stage or family each one stresses.

Calibration of the decision itself. The first quantity, decision ECE, measures the calibration of the probability a model assigns to the option it chooses within a declared set, rather than of a free-text answer. It is the main quantity for the training stage, and it is the one quantity that this survey measures across many systems. Panels (a) and (d) of Figure 11 report it from a community leaderboard of decision models, the Jev Decision Index [multimodalart 2026] (release 0.2.1, a single snapshot of 2026-09-27, not peer reviewed), which scores 68 systems on skill over 38 decision benchmarks and on expected calibration error over a 1-in-6 sample of 32 of them. Skill and calibration separate on this board: the system ranked first on skill, Jev (57.91, against 57.44 for the next; a single point estimate), is not the best calibrated, and 17 of the 67 open systems have a lower ECE (Table 2). The board reports point estimates without uncertainty intervals and records no cost per decision, so it measures the first quantity only.

Table 4. **A measurable agenda.** Seven quantities, each with a definition and a named public testbed, that together would let calibrated decision models be compared on common ground: calibration of the decision itself, safe abstention, cheap escalation, robustness to shift, the cost of a typed output contract, the size of a guaranteed set, and a controlled comparison of the training families. Figure 11 plots measured decision ECE (panels a and d) and the target shapes of selective risk and escalation regret (panels b and c).

Metric	Definition	Instantiation (testbed)	Stresses
Decision ECE	expected calibration error of the probability assigned to the chosen option within a <i>declared</i> option set, rather than of a free-text answer	MMLU as four-way choice [Hendrycks et al. 2021]; measured for every system on the Jev Decision Index (Table 2)	every stage; the main quantity for S4
Selective risk (AURC)	area under the risk-coverage curve obtained by abstaining below a confidence threshold [Geifman and El-Yaniv 2017]	LM-Polygraph benchmark [Vashurin et al. 2025]	S3, action policy
Escalation regret	extra cost, relative to the cheapest routing policy that reaches the same accuracy, when escalation to a larger model is triggered by confidence	RouterBench [Hu et al. 2024]	action policy, S1 to S4
Shift calibration gap	decision ECE on held-out subjects minus decision ECE on the subjects used for calibration or training	MMLU subject split, as in Shen et al. [2024]	S1, S4
Contract cost	schema-validity rate together with the accuracy change relative to unconstrained generation	JSONSchemaBench [Geng et al. 2025]; format-restriction study of Tam et al. [2024]	output contract
Guarantee tightness	average prediction-set size at a target coverage $1 - \alpha$	multiple-choice sets on MMLU [Hendrycks et al. 2021]	S3 (conformal)
Family comparison	decision ECE and accuracy of the five training families (T1 to T5) applied to the same base model and data	MMLU as four-way choice [Hendrycks et al. 2021], with the shared-base entries of the Jev Decision Index (Table 2) as a reference	S4 families

Acting on the confidence. Selective risk and escalation regret measure what the confidence buys once a rule acts on it: the error that remains after abstaining below a threshold [Geifman and El-Yaniv 2017], and the extra cost of escalating to a larger model when a cheaper policy would reach the same accuracy [Hu et al. 2024]. Panels (b) and (c) of Figure 11 draw their target shapes; they are not measured results.

Shift, contract and guarantee. The shift calibration gap asks whether a map or a trained confidence stays calibrated on subjects it was not fitted on [Shen et al. 2024]. The contract cost asks what a constrained or typed output costs in accuracy relative to free generation [Geng et al. 2025; Tam et al. 2024]. Guarantee tightness asks how large a prediction set must be to reach a target coverage. The last quantity, a comparison of the five training families on the same base model and data, is the controlled experiment that would say which training signal to prefer.

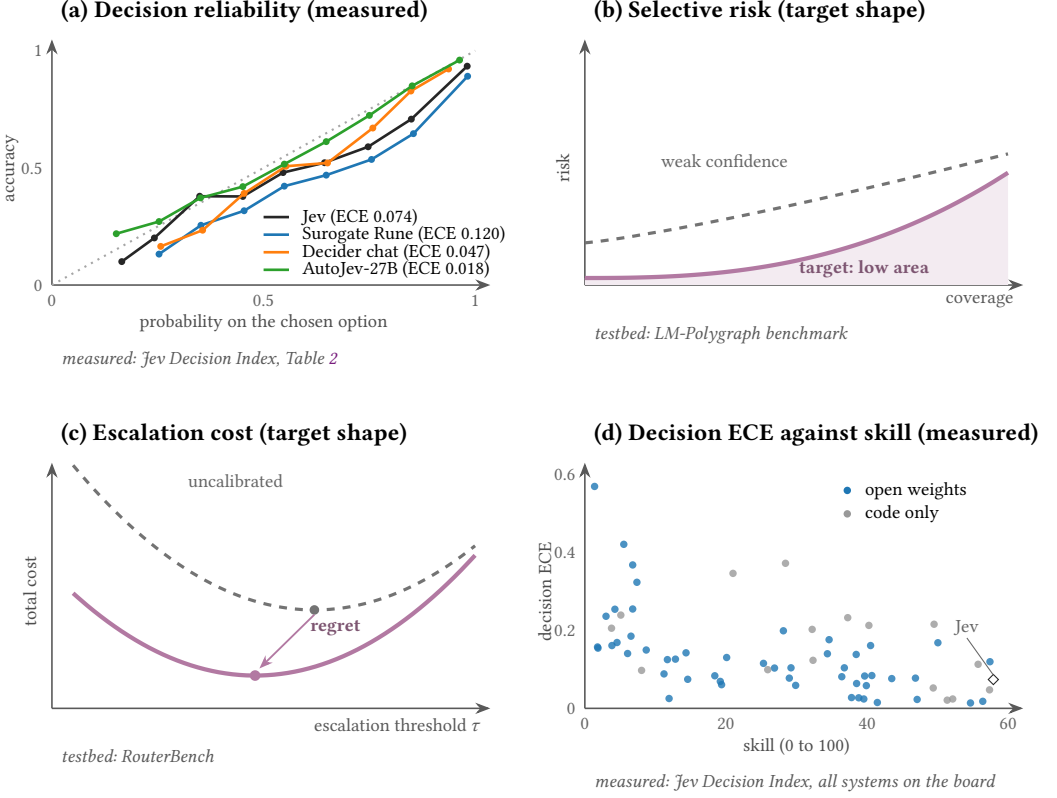


Fig. 11. **A measurable agenda, partly measured.** Decision ECE is the expected calibration error of the probability a model places on its chosen option (Table 4). Panels (a) and (d) come from one community leaderboard, the Jev Decision Index [multimodalart 2026] (release 0.2.1, a single snapshot of 2026-09-27; not peer reviewed): skill is its chance-normalised score over 38 decision benchmarks (0 to 100), ECE is computed on a 1-in-6 sample of 32 of them, and all values are point estimates without uncertainty intervals. (a) Reliability measured on the Jev Decision Index (Table 2) for Jev and the three open systems with the highest skill; each point is a confidence bin, and bins holding under 1% of answers are omitted. (b) Selective risk and (c) escalation regret are drawn as target shapes, not measured results: purple marks the target, grey dashed a typical baseline, and the named testbeds are in Table 4. (d) Decision ECE against skill for Jev and the 67 open systems on the board; lower right is better. The board records no cost per decision, so skill is used as the horizontal axis.

6 LIMITATIONS

Search. The query strings and the number of records screened were not recorded, so the search cannot be replayed exactly; Figure 4 marks those stages instead of estimating them. The 7 searches were run on one day, 2026-09-26, and the last year of the corpus is partial. Works placed only through the literature they cite, or not indexed by the searches, may be missing.

Placement. Stage, family and the per-work fields were assigned by reading each paper and were then checked against the papers in several review rounds and corrected where they disagreed. Some works sit on a boundary, for example a method that both trains a confidence and acts on it,

and another reader could place them differently. The one-value-per-work fields (contract, action, metric) record the main contribution and the main text, not every variant a paper reports.

Evidence. The families are compared as the papers report them, on different base models, data and metrics; no claim here ranks one family above another. The leaderboard behind Figure 11 and Table 2 is a single snapshot, is not peer reviewed, and reports point estimates without uncertainty intervals. The gallery figures are limited to papers whose licence permits reuse, so they show only part of each stage.

7 CONCLUSION

Calibration of language models, long estimated after the fact, is increasingly also trained into the model. We organised 408 works by where calibration enters the decision pipeline, and the 80 training-stage works by what the training signal scores: supervised calibration objectives, reward-model calibration, calibration rewards on the stated confidence, abstention and refusal training, and process- and meaning-level confidence rewards. Crossing that typology with the output contract shows where the field is thin: the typed decision, whose probability is the very interface a program acts on, is calibrated after the fact in many works but trained in only 7 of the 80 training works; most train a stated score or a free-text hedge or refusal. The seven quantities of Table 4 would let the families be compared on common ground, and the first of them can already be measured.

COMPETING INTERESTS

The authors have no financial or other relationship with TypeSafe AI, the maker of Jev, or with the maintainer of the Jev Decision Index. Jev is profiled as one instance of a hosted typed-decision model.

REFERENCES

- Momin Abbas, Yi Zhou, Parikshit Ram, Nathalie Baracaldo, Horst Samulowitz, Theodoros Salonidis, and Tianyi Chen. 2024. Enhancing In-context Learning via Linear Probe Calibration. In *Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS)*. arXiv:2401.12406
- Toghrol Abbasli, Kentaro Toyoda, Yuan Wang, Leon Witt, Muhammad Asif Ali, Yukai Miao, Dan Li, and Qingsong Wei. 2025. Comparing Uncertainty Measurement and Mitigation Methods for Large Language Models: A Systematic Review. arXiv:2504.18346 IEEE TNNLS (accepted).
- Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqi, and Mausam. 2024. AutoMix: Automatically Mixing Language Models. arXiv:2310.12963 NeurIPS 2024.
- Pranjal Aggarwal, Aman Madaan, Yiming Yang, and Mausam. 2023. Let’s Sample Step by Step: Adaptive-Consistency for Efficient Reasoning and Coding with LLMs. arXiv:2305.11860 EMNLP 2023.
- Ayush Agrawal, Mirac Suzgun, Lester Mackey, and Adam Tauman Kalai. 2023. Do Language Models Know When They’re Hallucinating References? arXiv:2305.18248
- Kabir Ahuja, Sunayana Sitaram, Sandipan Dandapat, and Monojit Choudhury. 2022. On the Calibration of Massively Multilingual Language Models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2022.emnlp-main.290> arXiv:2210.12265
- Lukas Aichberger, Kajetan Schweighofer, Mykyta Ielanskyi, and Sepp Hochreiter. 2025. Improving Uncertainty Estimation through Semantically Diverse Language Generation. In *International Conference on Learning Representations (ICLR)*. arXiv:2406.04306
- Aisha Alansari and Hamzah Luqman. 2026. Large language models hallucination: A comprehensive survey. <https://doi.org/10.1016/j.cosrev.2026.100970> Computer Science Review.
- Diogo Almeida. 2026. Introducing System One Models and Jev. TypeSafe AI blog, <https://typesafe.ai/blog/introducing-system-one-models-and-jev>. Posted 15 September 2026.
- Hao An and Yang Xu. 2025. Teaching LLMs to Abstain via Fine-Grained Semantic Confidence Reward. arXiv:2510.24020
- Anastasios N. Angelopoulos, Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. 2025. Learn then Test: Calibrating Predictive Algorithms to Achieve Risk Control. *Annals of Applied Statistics* (2025). <https://doi.org/10.1214/24-aos1998> arXiv:2110.01052
- Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. 2022. Conformal Risk Control. In *ICLR 2024*. arXiv:2208.02814
- Shuang Ao, Stefan Rueger, and Advait Siddharthan. 2024. CSS: Contrastive Semantic Similarity for Uncertainty Quantification of LLMs. In *Conference on Uncertainty in Artificial Intelligence (UAI)*. arXiv:2406.03158
- Alice Evelyn Ashby, Khuong An Nguyen, and Zhiyuan Luo. 2026. Uncertainty-aware large language models: a scoping review of conformal prediction methods. <https://doi.org/10.1098/rsta.2026.0061> Phil. Trans. R. Soc. A.
- Amos Azaria and Tom Mitchell. 2023. The Internal State of an LLM Knows When It’s Lying. arXiv:2304.13734
- Sanket Badhe, Priyanka Tiwari, and Deep Shah. 2026. The Silent Vote: Improving Zero-Shot LLM Reliability by Aggregating Semantic Neighborhoods. arXiv:2605.09739 GEM Workshop ACL 2026.
- Yavuz Faruk Bakman, Duygu Nur Yaldiz, Baturalp Buyukates, Chenyang Tao, Dimitrios Dimitriadis, and Salman Avestimehr. 2024. MARS: Meaning-Aware Response Scoring for Uncertainty Estimation in Generative LLMs. In *Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2024.acl-long.419> arXiv:2402.11756
- Neil Band, Xuechen Li, Tengyu Ma, and Tatsunori Hashimoto. 2024. Linguistic Calibration of Long-Form Generations. In *ICML 2024*. arXiv:2404.00474
- Debangshu Banerjee, Tarun Suresh, Shubham Ugare, Sasa Misailovic, and Gagandeep Singh. 2025. CRANE: Reasoning with constrained LLM generation. arXiv:2502.09061 ICML 2025.
- David Bani-Harouni, Chantal Pellegrini, Julian Lüers, Su Hwan Kim, Markus Baalman, Benedikt Wiestler, Rickmer Braren, Nassir Navab, and Matthias Keicher. 2026a. Calibrated Confidence Expression for Radiology Report Generation. In *LNCS*. https://doi.org/10.1007/978-3-032-38079-1_5 arXiv:2603.29492
- David Bani-Harouni, Chantal Pellegrini, Paul Stangel, Ege Özsoy, Kamelia Zaripova, Nassir Navab, and Matthias Keicher. 2026b. Rewarding Doubt: A Reinforcement Learning Approach to Calibrated Confidence Expression of Large Language Models. In *International Conference on Learning Representations (ICLR)*. arXiv:2503.02623
- Joas Antonio dos Santos Barbosa. 2026. Calibrated Decision Models for Autonomous Penetration-Testing Harnesses: JEV and Laya as System One Decision Layers for LLM-Driven Pentest Agents. arXiv:2609.28940
- Evan Becker and Stefano Soatto. 2024. Cycles of Thought: Measuring LLM Confidence through Stable Explanations. arXiv:2406.03441
- Mohammad Beigi, Ying Shen, Runing Yang, Zihao Lin, Qifan Wang, Ankith Mohan, Jianfeng He, Ming Jin, Chang-Tien Lu, and Lifu Huang. 2024a. InternalInspector I^2 : Robust Confidence Estimation in LLMs through Internal States. In *Findings of the Association for Computational Linguistics: EMNLP*. <https://doi.org/10.18653/v1/2024.findings-emnlp.751>

- arXiv:2406.12053
- Mohammad Beigi, Sijia Wang, Ying Shen, Zihao Lin, Adithya Kulkarni, Jianfeng He, Feng Chen, Ming Jin, Jin-Hee Cho, Dawei Zhou, Chang-Tien Lu, and Lifu Huang. 2024b. Rethinking the Uncertainty: A Critical Review and Analysis in the Era of Large Language Models. arXiv:2410.20199
- Catarina G Belem, Markelle Kelly, Mark Steyvers, Sameer Singh, and Padhraic Smyth. 2024. Perceptions of Linguistic Uncertainty by Language Models and Humans. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2407.15814
- Michael Bereket and Jure Leskovec. 2025. Uncalibrated Reasoning: GRPO Induces Overconfidence for Stochastic Outcomes. arXiv:2508.11800
- Luca Beurer-Kellner, Marc Fischer, and Martin Vechev. 2023. Prompting Is Programming: A Query Language for Large Language Models. <https://doi.org/10.1145/3591300> arXiv:2212.06094 PACMPL (PLDI 2023).
- Luca Beurer-Kellner, Marc Fischer, and Martin Vechev. 2024. Guiding LLMs The Right Way: Fast, Non-Invasive Constrained Generation. arXiv:2403.06988
- Dylan Bouchard. 2026. Is Escalation Worth It? A Decision-Theoretic Characterization of LLM Cascades. arXiv:2605.06350
- Dylan Bouchard, Mohit Singh Chauhan, Zeya Ahmad, and Ho-Kyeong Ra. 2026. Functional Entropy: Predicting Functional Correctness in LLM-Generated Code with Uncertainty Quantification. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2605.28500
- Abderrahmene Boudiaf, Irfan Hussain, and Sajid Javed. 2026. Uncertainty-Aware Decision Making in Multimodal Large Language Models. arXiv:2608.17084
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. 2022. Discovering Latent Knowledge in Language Models Without Supervision. In *International Conference on Learning Representations (ICLR)*. arXiv:2212.03827
- Jarosław Błasiok, Parikshit Gopalan, Lunjia Hu, and Preetum Nakkiran. 2023. When Does Optimizing a Proper Loss Yield Calibration?. In *NeurIPS 2023*. arXiv:2305.18764
- Jon-Paul Cacioli. 2026. Verbal Confidence Saturation in 3-9B Open-Weight Instruction-Tuned LLMs: A Pre-Registered Psychometric Validity Screen. arXiv:2604.22215
- Margarida M. Campos, António Farinhas, Chrysoula Zerva, Mário A. T. Figueiredo, and André F. T. Martins. 2024. Conformal Prediction for Natural Language Processing: A Survey. *Transactions of the Association for Computational Linguistics (TACL)* (2024). https://doi.org/10.1162/tacl_a_00715 arXiv:2405.01976
- Sky CH-Wang, Benjamin Van Durme, Jason Eisner, and Chris Kedzie. 2024. Do Androids Know They’re Only Dreaming of Electric Sheep?. In *Findings of the Association for Computational Linguistics: ACL*. <https://doi.org/10.18653/v1/2024.findings-acl.260> arXiv:2312.17249
- Kwan Ho Ryan Chan, Yuyan Ge, Edgar Dobriban, Hamed Hassani, and René Vidal. 2025. Conformal Information Pursuit for Interactively Guiding Large Language Models. In *NeurIPS 2025*. arXiv:2507.03279
- Arslan Chaudhry, Sridhar Thiagarajan, and Dilan Gorur. 2024. Finetuning Language Models to Emit Linguistic Expressions of Uncertainty. arXiv:2409.12180
- Xujun Che, Yuchen Yuan, Weida Zhao, and Chenyang Yu. 2026. Abstention as an Action Can Kill Both the Reward Gradient and the KL Anchor: Collapse Law and Repair for Error-Penalized Reinforcement Learning. arXiv:2608.00301
- Chao Chen, Kai Liu, Ze Chen, Yi Gu, Yue Wu, Mingyuan Tao, Zhihang Fu, and Jieping Ye. 2024. INSIDE: LLMs’ Internal States Retain the Power of Hallucination Detection. In *International Conference on Learning Representations (ICLR)*. arXiv:2402.03744
- Catherine Yu-Chi Chen, Jingyan Shen, Zhun Deng, and Lihua Lei. 2025b. Conformal Tail Risk Control for Large Language Model Alignment. In *ICML 2025*. arXiv:2502.20285
- Jiuhai Chen and Jonas Mueller. 2023. Quantifying Uncertainty in Answers from any Language Model and Enhancing their Trustworthiness. arXiv:2308.16175
- Lida Chen, Zujie Liang, Xintao Wang, Jiaqing Liang, Yanghua Xiao, Feng Wei, Jinglei Chen, Zhenghong Hao, Bing Han, and Wei Wang. 2025a. Teaching Large Language Models to Express Knowledge Boundary from Their Own Signals. In *KnowLLM-W 2025*. <https://doi.org/10.18653/v1/2025.knowllm-1.3> arXiv:2406.10881
- Lingjiao Chen, Matei Zaharia, and James Zou. 2023b. FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance. arXiv:2305.05176
- Yangyi Chen, Lifan Yuan, Ganqu Cui, Zhiyuan Liu, and Heng Ji. 2023a. A Close Look into the Calibration of Pre-trained Language Models. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2023.acl-long.75> arXiv:2211.00151
- Qinyuan Cheng, Tianxiang Sun, Xiangyang Liu, Wenwei Zhang, Zhangyue Yin, Shimin Li, Linyang Li, Zhengfu He, Kai Chen, and Xipeng Qiu. 2024. Can AI Assistants Know What They Don’t Know? arXiv:2401.13275
- John J. Cherian, Isaac Gibbs, and Emmanuel J. Candès. 2024. Large language model validity via enhanced conformal prediction methods. In *NeurIPS 2024*. arXiv:2406.09714

- Prateek Chhikara. 2025. Mind the Confidence Gap: Overconfidence, Calibration, and Distractor Effects in Large Language Models. *Transactions on Machine Learning Research* (2025). arXiv:2502.11028
- Hakaze Cho, Yoshihiro Sakai, Mariko Kato, Kenshiro Tanaka, Akira Ishii, and Naoya Inoue. 2025. Token-based Decision Criteria Are Suboptimal in In-context Learning. In *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*. <https://doi.org/10.18653/v1/2025.naacl-long.278> arXiv:2406.16535
- Prafulla Kumar Choubey, Yu Bai, Chien-Sheng Wu, Wenhao Liu, and Nazneen Rajani. 2022. Conformal Predictor for Improving Zero-shot Text Classification Efficiency. In *EMNLP 2022*. <https://doi.org/10.18653/v1/2022.emnlp-main.196> arXiv:2210.12619
- Yu-Neng Chuang, Prathusha Kameswara Sarma, Parikshit Gopalan, John Boccio, Sara Bolouki, Xia Hu, and Helen Zhou. 2024. Learning to Route LLMs with Confidence Tokens. arXiv:2410.13284
- Kamil Ciosek, Nicolò Felicioni, and Sina Ghiassian. 2025. Hallucination Detection on a Budget: Efficient Bayesian Estimation of Semantic Entropy. arXiv:2504.03579
- Roi Cohen, May Hamri, Mor Geva, and Amir Globerson. 2023. LM vs LM: Detecting Factual Errors via Cross Examination. arXiv:2305.13281
- Jeremy R. Cole, Michael J. Q. Zhang, Daniel Gillick, Julian Martin Eisenschlos, Bhuwan Dhingra, and Jacob Eisenstein. 2023. Selectively Answering Ambiguous Questions. arXiv:2305.14613 EMNLP 2023.
- Peng Cui, Boyao Yang, and Jun Zhu. 2026. Ranking-Aware Calibration for Reliable Multimodal Reinforcement Learning. arXiv:2605.16999
- Longchao Da, Tiejun Chen, Lu Cheng, and Hua Wei. 2024. LLM Uncertainty Quantification through Directional Entailment Graph and Claim Level Response Augmentation. arXiv:2407.00994
- Mehul Damani, Isha Puri, Stewart Slocum, Idan Shenfeld, Leshem Choshen, Yoon Kim, and Jacob Andreas. 2025. Beyond Binary Rewards: Training LMs to Reason About Their Uncertainty. arXiv:2507.16806
- Mehul Damani, Idan Shenfeld, Andi Peng, Andreea Bobu, and Jacob Andreas. 2024. Learning How Hard to Think: Input-Adaptive Allocation of LM Computation. arXiv:2410.04707
- Soham Dan and Dan Roth. 2021. On the Effects of Transformer Size on In- and Out-of-Domain Calibration. In *Findings of the Association for Computational Linguistics: EMNLP*. <https://doi.org/10.18653/v1/2021.findings-emnlp.180>
- Shrey Desai and Greg Durrett. 2020. Calibration of Pre-trained Transformers. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2020.emnlp-main.21> arXiv:2003.07892
- Gianluca Detommaso, Martin Bertran, Riccardo Fogliato, and Aaron Roth. 2024. Multicalibration for Confidence Scoring in LLMs. arXiv:2404.04689
- Nicolas Deutschmann, Marvin Alberts, and María Rodríguez Martínez. 2023. Conformal Autoregressive Generation: Beam Search with Coverage Guarantees. In *AAAI 2024*. arXiv:2309.03797
- Neil Dey, Jing Ding, Jack Ferrell, Carolina Kapper, Maxwell Lovig, Emiliano Planchon, and Jonathan P Williams. 2022. Conformal prediction for text infilling and part-of-speech prediction. *The New England Journal of Statistics in Data Science* (2022). <https://doi.org/10.51387/22-nejsds8> arXiv:2111.02592
- Shehzaad Dhuliawala, Leonard Adolphs, Rajarshi Das, and Mrinmaya Sachan. 2022. Calibration of Machine Reading Systems at Scale. In *Findings of the Association for Computational Linguistics: ACL*. <https://doi.org/10.18653/v1/2022.findings-acl.133> arXiv:2203.10623
- Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Ruhle, Laks V. S. Lakshmanan, and Ahmed Hassan Awadallah. 2024. Hybrid LLM: Cost-Efficient and Quality-Aware Query Routing. arXiv:2404.14618 ICLR 2024.
- Li Dong, Chris Quirk, and Mirella Lapata. 2018. Confidence Modeling for Neural Semantic Parsing. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/P18-1069> arXiv:1805.04604
- Yixin Dong, Charlie F. Ruan, Yaxing Cai, Ruihang Lai, Ziyi Xu, Yilong Zhao, and Tianqi Chen. 2025. XGrammar: Flexible and Efficient Structured Generation Engine for Large Language Models. arXiv:2411.15100 MLSys 2025.
- Achref Doula, Max Mühlhäuser, and Alejandro Sanchez Guinea. 2025. SafePath: Conformal Prediction for Safe LLM-Based Autonomous Navigation. arXiv:2505.09427
- Xuefeng Du, Chaowei Xiao, and Yixuan Li. 2024. HaloScope: Harnessing Unlabeled LLM Generations for Hallucination Detection. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2409.17504
- Jinhao Duan, Hao Cheng, Shiqi Wang, Alex Zavalny, Chenan Wang, Renjing Xu, Bhavya Kailkhura, and Kaidi Xu. 2024. Shifting Attention to Relevance: Towards the Predictive Uncertainty Quantification of Free-Form Large Language Models. In *Annual Meeting of the Association for Computational Linguistics (ACL)*. arXiv:2307.01379
- Bryan Eikema, Evgenia Ilia, José G. C. de Souza, Chrysoula Zerva, and Wilker Aziz. 2025. Teaching Language Models to Faithfully Express their Uncertainty. arXiv:2510.12587
- Nicolas Emmenegger, Theo X. Olausson, Armando Solar-Lezama, and Chara Podimata. 2026. Conformal Language Modeling via Posterior Sampling. arXiv:2606.03731

- Ekaterina Fadeeva, Aleksandr Rubashevskii, Artem Shelmanov, Sergey Petrakov, Haonan Li, Hamdy Mubarak, Evgenii Tsymbalov, Gleb Kuzmin, Alexander Panchenko, Timothy Baldwin, Preslav Nakov, and Maxim Panov. 2024. Fact-Checking the Output of Large Language Models via Token-Level Uncertainty Quantification. In *Findings of the Association for Computational Linguistics: ACL 2024*. arXiv:2403.04696
- Ekaterina Fadeeva, Roman Vashurin, Akim Tsvigun, Artem Vazhentsev, Sergey Petrakov, Kirill Fedyanin, Daniil Vasilev, Elizaveta Goncharova, Alexander Panchenko, Maxim Panov, Timothy Baldwin, and Artem Shelmanov. 2023. LM-Polygraph: Uncertainty Estimation for Language Models. In *EMNLP System Demonstrations*. arXiv:2311.07383
- Hengxin Fan. 2026. Capacity, Not Format: Rethinking Structured Reasoning Failures. arXiv:2606.09410
- Haozhi Fan, Jinhao Duan, and Kaidi Xu. 2026. IUQ: Interrogative Uncertainty Quantification for Long-Form Large Language Model Generation. In *Annual Meeting of the Association for Computational Linguistics (ACL)*. arXiv:2604.15109
- Claudio Fanconi and Mihaela van der Schaar. 2025. Cascaded Language Models for Cost-effective Human-AI Decision-Making. arXiv:2506.11887
- Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. 2024. Detecting hallucinations in large language models using semantic entropy. *Nature* (2024). <https://doi.org/10.1038/s41586-024-07421-0>
- Yu Fei, Yifan Hou, Zeming Chen, and Antoine Bosselut. 2023. Mitigating Label Biases for In-context Learning. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2023.acl-long.783> arXiv:2305.19148
- Naihe Feng, Yi Sui, Shiyi Hou, Jesse C. Cresswell, and Ga Wu. 2025. Response Quality Assessment for Retrieval-Augmented Generation via Conditional Conformal Factuality. In *SIGIR 2025*. <https://doi.org/10.1145/3726302.3730244> arXiv:2506.20978
- Shangbin Feng, Weijia Shi, Yike Wang, Wenxuan Ding, Vidhisha Balachandran, and Yulia Tsvetkov. 2024. Don't Hallucinate, Abstain: Identifying LLM Knowledge Gaps via Multi-LLM Collaboration. arXiv:2402.00367 ACL 2024.
- Yushi Feng, Junye Du, Qifan Wang, Zizhan Ma, Qian Niu, Yutaka Matsuo, Long Feng, and Lequan Yu. 2026. CORA: Conformal Risk-Controlled Agents for Safeguarded Mobile GUI Automation. arXiv:2604.09155
- Robinson Ferrer, Damla Turgut, Zhongzhou Chen, and Shashank Sonkar. 2026. When Can We Trust LLM Graders? Calibrating Confidence for Automated Assessment. arXiv:2603.29559
- Matthew French Constant, Daniel Yang, Ximeng Huang, and Sanyam Kapoor. 2026. ConfidenceBench: Evaluating Confidence Calibration in Large Language Models. arXiv:2607.20526
- Adam Fisch, Tal Schuster, Tommi Jaakkola, and Regina Barzilay. 2020. Efficient Conformal Prediction via Cascaded Inference with Expanded Admission. In *ICLR 2021*. arXiv:2007.03114
- Adam Fisch, Tal Schuster, Tommi Jaakkola, and Regina Barzilay. 2021. Few-shot Conformal Prediction with Auxiliary Tasks. In *ICML 2021*. arXiv:2102.08898
- Marina Fomicheva, Shuo Sun, Lisa Yankovskaya, Frédéric Blain, Francisco Guzmán, Mark Fishel, Nikolaos Aletras, Vishrav Chaudhary, and Lucia Specia. 2020. Unsupervised Quality Estimation for Neural Machine Translation. *Transactions of the Association for Computational Linguistics* (2020). arXiv:2005.10608
- Robert Friel and Atindriyo Sanyal. 2023. Chainpoll: A high efficacy method for LLM hallucination detection. arXiv:2310.18344
- Yichao Fu, Xuewei Wang, Yuandong Tian, and Jiawei Zhao. 2025. Deep Think with Confidence. arXiv:2508.15260
- Cosimo Galeone, Minsu Park, Giuseppe Ettore, and Daniele Ligorio. 2026. When Correct Isn't Usable: Improving Structured Output Reliability in Small Language Models. arXiv:2605.02363
- Cheng Gao, Cheng Huang, Kangyang Luo, Ziqing Qiao, Shuzheng Si, Huimin Chen, Chaojun Xiao, and Maosong Sun. 2026. KARL: Mitigating Hallucinations in LLMs via Knowledge-Boundary-Aware Reinforcement Learning. arXiv:2604.22779
- Xiang Gao, Jiaxin Zhang, Lalla Mouatadid, and Kamalika Das. 2024. SPUQ: Perturbation-Based Uncertainty Quantification for Large Language Models. In *Conference of the European Chapter of the Association for Computational Linguistics (EACL)*. arXiv:2403.02509
- Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *NeurIPS*. arXiv:1705.08500
- Jiahui Geng, Fengyu Cai, Yuxia Wang, Heinz Koepl, Preslav Nakov, and Iryna Gurevych. 2024a. A Survey of Confidence Estimation and Calibration in Large Language Models. arXiv:2311.08298 NAACL 2024 (long).
- Saibo Geng, Hudson Cooper, Michal Moskal, Samuel Jenkins, Julian Berman, Nathan Ranchin, Robert West, Eric Horvitz, and Harsha Nori. 2025. JSONSchemaBench: A Rigorous Benchmark of Structured Outputs for Language Models. arXiv:2501.10868
- Saibo Geng, Berkay Döner, Chris Wendler, Martin Josifoski, and Robert West. 2024b. Sketch-Guided Constrained Decoding for Boosting Blackbox Large Language Models without Logit Access. arXiv:2401.09967 ACL 2024.
- Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. 2023. Grammar-Constrained Decoding for Structured NLP Tasks without Finetuning. arXiv:2305.13971 EMNLP 2023.
- Bijean Ghafouri, Shahrad Mohammadzadeh, James Zhou, Pratheeksha Nair, Jacob-Junqi Tian, Hikaru Tsujimura, Mayank Goel, Sukanya Krishna, Reihaneh Rabbany, Jean-François Godbout, and Kellin Pelrine. 2024. Epistemic Integrity in Large Language Models. arXiv:2411.06528

- Daniela Gottesman and Mor Geva. 2024. Estimating Knowledge in Large Language Models Without Generating a Single Token. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2024.emnlp-main.232> arXiv:2406.12673
- Yashvir S. Grewal, Edwin V. Bonilla, and Thang D. Bui. 2024. Improving Uncertainty Quantification in Large Language Models via Semantic Embeddings. arXiv:2410.22685
- Tobias Groot and Matias Valdenegro-Toro. 2024. Overconfidence is Key: Verbalized Uncertainty Evaluation in Large Language and Vision-Language Models. In *Proceedings of the Workshop on Trustworthy Natural Language Processing (TrustNLP) at NAACL*. arXiv:2405.02917
- Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, Saizhuo Wang, Kun Zhang, Yuanzhuo Wang, Wen Gao, Lionel Ni, and Jian Guo. 2024. A Survey on LLM-as-a-Judge. arXiv:2411.15594
- Yu Gui, Ying Jin, and Zhimei Ren. 2024. Conformal Alignment: Knowing When to Trust Foundation Models with Guarantees. In *NeurIPS 2024*. arXiv:2405.10301
- Korel Gundem, Juncheng Dong, Dennis Zhang, Vahid Tarokh, and Zhengling Qi. 2025. Boosting In-Context Learning in LLMs Through the Lens of Classical Supervised Learning. arXiv:2505.23783
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On Calibration of Modern Neural Networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*. arXiv:1706.04599
- Junyu Guo, Shangding Gu, Ming Jin, Costas Spanos, and Javad Lavaei. 2026. LLMs Should Express Uncertainty Explicitly. arXiv:2604.05306
- Chirag Gupta and Aaditya K. Ramdas. 2021. Distribution-free calibration guarantees for histogram binning without sample splitting. In *Proceedings of the International Conference on Machine Learning (ICML)*. arXiv:2105.04656
- Neha Gupta, Harikrishna Narasimhan, Wittawat Jitkrittum, Ankit Singh Rawat, Aditya Krishna Menon, and Sanjiv Kumar. 2024. Language Model Cascades: Token-level uncertainty and beyond. arXiv:2404.10136
- Daehoon Gwak, Minhyung Lee, Junwoo Park, and Jaegul Choo. 2026. Accelerating Masked Diffusion Large Language Models: A Survey of Efficient Inference Techniques. arXiv:2607.12829
- Sophia Hager, David Mueller, Kevin Duh, and Nicholas Andrews. 2025. Uncertainty Distillation: Teaching Language Models to Express Semantic Confidence. arXiv:2503.14749
- Kimia Hamidieh, Veronika Thost, Walter Gerych, Mikhail Yurochkin, and Marzyeh Ghassemi. 2026. Complementing Self-Consistency with Cross-Model Disagreement for Uncertainty Quantification. arXiv:2604.17112
- Haixia Han, Tingyun Li, Shisong Chen, Jie Shi, Chengyu Du, Yanghua Xiao, Jiaqing Liang, and Xin Lin. 2024. Enhancing Confidence Expression in Large Language Models Through Learning from Past Experience. arXiv:2404.10315
- Zhixiong Han, Yaru Hao, Li Dong, Yutao Sun, and Furu Wei. 2022. Prototypical Calibration for Few-shot Learning of Language Models. arXiv:2205.10183
- Jianfeng He, Linlin Yu, Changbin Li, Runing Yang, Fanglan Chen, Kangshuo Li, Min Zhang, Shuo Lei, Xuchao Zhang, Mohammad Beigi, Kaize Ding, Bei Xiao, Lifu Huang, Feng Chen, Ming Jin, and Chang-Tien Lu. 2026b. Survey of uncertainty estimation in LLMs - Sources, methods, applications, and challenges. *Information Fusion* 130 (2026), 104057. <https://doi.org/10.1016/j.inffus.2025.104057>
- Weilin He, Arindam Sharma, and Cristina David. 2026a. Using Semantic Distance to Estimate Uncertainty in LLM-Based Code Generation. arXiv:2605.09023
- Zhitao He, Sandeep Polisetty, Zhiyuan Fan, Yuchen Huang, Shujin Wu, and Yi R. Fung. 2025. MMBoundary: Advancing MLLM Knowledge Boundary Awareness through Reasoning Step Confidence Calibration. In *ACL 2025*. <https://doi.org/10.18653/v1/2025.acl-long.802> arXiv:2505.23224
- Kilian Hendrickx, Lorenzo Perini, Dries Van der Plas, Wannes Meert, and Jesse Davis. 2024. Machine Learning with a Reject Option: A survey. arXiv:2107.11277 *Machine Learning* (Springer).
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. Measuring Massive Multitask Language Understanding. In *ICLR*. arXiv:2009.03300
- Ari Holtzman, Peter West, Vered Schwartz, Yejin Choi, and Luke Zettlemoyer. 2021. Surface Form Competition: Why the Highest Probability Answer Isn't Always Right. arXiv:2104.08315
- Bairu Hou, Yujian Liu, Kaizhi Qian, Jacob Andreas, Shiyu Chang, and Yang Zhang. 2024. Decomposing Uncertainty for Large Language Models through Input Clarification Ensembling. In *International Conference on Machine Learning (ICML)*. arXiv:2311.08718
- Yu-Chung Hsiao. 2026. Rethinking Verbalized Confidence for LLM-as-a-Judge: A Compatibility Shift on Post-2025 Proprietary Models. arXiv:2609.10996
- Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt Keutzer, and Shriyash Kautubh Upadhyay. 2024. RouterBench: A Benchmark for Multi-LLM Routing System. arXiv:2403.12031
- Yunhai Hu, Zining Liu, Zhenyuan Dong, Tianfan Peng, Bradley McDanel, and Sai Qian Zhang. 2025. Speculative Decoding and Beyond: An In-Depth Survey of Techniques. arXiv:2502.19732

- Chengsong Huang, Langlin Huang, Jixuan Leng, Jiacheng Liu, and Jiaxin Huang. 2025a. Efficient Test-Time Scaling via Self-Calibration. [arXiv:2503.00031](https://arxiv.org/abs/2503.00031)
- Hsiu-Yuan Huang, Yutong Yang, Zhaoxi Zhang, Sanwoo Lee, and Yunfang Wu. 2024c. A Survey of Uncertainty Estimation in LLMs: Theory Meets Practice. [arXiv:2410.15326](https://arxiv.org/abs/2410.15326)
- Jerry Huang, Peng Lu, Qiuhaio Zeng, Yusuke Iwasawa, Yutaka Matsuo, Sarath Chandar, Edison Marrese-Taylor, and Irene Li. 2026. Investigating the Multilingual Calibration Effects of Language Model Instruction-Tuning. In *Proceedings of the Conference of the European Chapter of the Association for Computational Linguistics (EACL)*. <https://doi.org/10.18653/v1/2026.eacl-short.1> [arXiv:2601.01362](https://arxiv.org/abs/2601.01362)
- Jiayi Huang, Sangwoo Park, Nicola Paoletti, and Osvaldo Simeone. 2025b. Reliable Inference in Edge-Cloud Model Cascades via Conformal Alignment. [arXiv:2510.17543](https://arxiv.org/abs/2510.17543)
- Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2025d. A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. [arXiv:2311.05232](https://arxiv.org/abs/2311.05232) ACM TOIS.
- Xinmeng Huang, Shuo Li, Mengxin Yu, Matteo Sesia, Hamed Hassani, Insup Lee, Osbert Bastani, and Edgar Dobriban. 2024a. Uncertainty in Language Models: Assessment through Rank-Calibration. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. [arXiv:2404.03163](https://arxiv.org/abs/2404.03163)
- Yukun Huang, Yixin Liu, Raghuveer Thirukovalluru, Arman Cohan, and Bhuwan Dhingra. 2024b. Calibrating Long-form Generations from Large Language Models. [arXiv:2402.06544](https://arxiv.org/abs/2402.06544)
- Yuheng Huang, Jiayang Song, Zhijie Wang, Shengming Zhao, Huaming Chen, Felix Juefei-Xu, and Lei Ma. 2025c. Look Before You Leap: An Exploratory Study of Uncertainty Measurement for Large Language Models. *IEEE Transactions on Software Engineering* (2025). <https://doi.org/10.1109/TSE.2024.3519464> [arXiv:2307.10236](https://arxiv.org/abs/2307.10236)
- Chaeyun Jang, Moonseok Choi, Yegon Kim, Hyungi Lee, and Juho Lee. 2025. Verbalized Confidence Triggers Self-Verification: Emergent Behavior Without Explicit Reasoning Supervision. [arXiv:2506.03723](https://arxiv.org/abs/2506.03723)
- Metod Jazbec, Alexander Timans, Tin Hadži Veljković, Kaspar Sakmann, Dan Zhang, Christian A. Naesseth, and Eric Nalisnick. 2024. Fast yet Safe: Early-Exiting with Risk Control. In *NeurIPS 2024*. [arXiv:2405.20915](https://arxiv.org/abs/2405.20915)
- Ziwei Ji, Delong Chen, Etsuko Ishii, Samuel Cahyawijaya, Yejin Bang, Bryan Wilie, and Pascale Fung. 2024. LLM Internal States Reveal Hallucination Risk Faced With a Query. In *Proceedings of the BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*. <https://doi.org/10.18653/v1/2024.blackboxnlp-1.6> [arXiv:2407.03282](https://arxiv.org/abs/2407.03282)
- Mingjian Jiang, Yangjun Ruan, Prasanna Sattigeri, Salim Roukos, and Tatsunori Hashimoto. 2024. Graph-based Uncertainty Metrics for Long-form Language Model Outputs. In *Advances in Neural Information Processing Systems (NeurIPS)*. [arXiv:2410.20783](https://arxiv.org/abs/2410.20783)
- Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. 2021. How Can We Know When Language Models Know? On the Calibration of Language Models for Question Answering. *Transactions of the Association for Computational Linguistics* (2021). https://doi.org/10.1162/tacl_a_00407 [arXiv:2012.00955](https://arxiv.org/abs/2012.00955)
- Zhengping Jiang, Anqi Liu, and Benjamin Van Durme. 2025. Conformal Linguistic Calibration: Trading-off between Factuality and Specificity. In *NeurIPS 2025*. [arXiv:2502.19110](https://arxiv.org/abs/2502.19110)
- Zhengbao Jiang, Frank F. Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. 2023a. Active Retrieval Augmented Generation. [arXiv:2305.06983](https://arxiv.org/abs/2305.06983) EMNLP 2023.
- Zhongtao Jiang, Yuanzhe Zhang, Cao Liu, Jun Zhao, and Kang Liu. 2023b. Generative Calibration for In-context Learning. In *Findings of the Association for Computational Linguistics: EMNLP*. <https://doi.org/10.18653/v1/2023.findings-emnlp.152> [arXiv:2310.10266](https://arxiv.org/abs/2310.10266)
- Wittawat Jitkrittum, Neha Gupta, Aditya Krishna Menon, Harikrishna Narasimhan, Ankit Singh Rawat, and Sanjiv Kumar. 2023. When Does Confidence-Based Cascade Deferral Suffice? [arXiv:2307.02764](https://arxiv.org/abs/2307.02764) NeurIPS 2023.
- Abhinav Joshi, Areeb Ahmad, and Ashutosh Modi. 2025. Calibration Across Layers: Understanding Calibration Evolution in LLMs. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2025.emnlp-main.742> [arXiv:2511.00280](https://arxiv.org/abs/2511.00280)
- Jaehun Jung, Faeze Brahman, and Yejin Choi. 2024. Trust or Escalate: LLM Judges with Provable Guarantees for Human Agreement. [arXiv:2407.18370](https://arxiv.org/abs/2407.18370)
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislav Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared Kaplan. 2022. Language Models (Mostly) Know What They Know. [arXiv:2207.05221](https://arxiv.org/abs/2207.05221)
- Adam Tauman Kalai, Ofir Nachum, Santosh S. Vempala, and Edwin Zhang. 2025. Why Language Models Hallucinate. [arXiv:2509.04664](https://arxiv.org/abs/2509.04664)

- Adam Tauman Kalai and Santosh S. Vempala. 2024. Calibrated Language Models Must Hallucinate. In *STOC 2024*. <https://doi.org/10.1145/3618260.3649777> arXiv:2311.14648
- Sahil Kale and Devendra Singh Dhami. 2025. KnowRL: Teaching Language Models to Know What They Know. arXiv:2510.11407
- Amita Kamath, Robin Jia, and Percy Liang. 2020. Selective Question Answering under Domain Shift. arXiv:2006.09462 ACL 2020.
- Katie Kang, Eric Wallace, Claire Tomlin, Aviral Kumar, and Sergey Levine. 2025b. Unfamiliar Finetuning Examples Control How Language Models Hallucinate. In *NAACL 2025*. <https://doi.org/10.18653/v1/2025.naacl-long.183> arXiv:2403.05612
- Mintong Kang, Nezihe Merve Gürel, Ning Yu, Dawn Song, and Bo Li. 2024. C-RAG: Certified Generation Risks for Retrieval-Augmented Language Models. In *ICML 2024*. arXiv:2402.03181
- Sungmin Kang, Yavuz Faruk Bakman, Duygu Nur Yaldiz, Baturalp Buyukates, and Salman Avestimehr. 2025a. Uncertainty Quantification for Hallucination Detection in Large Language Models: Foundations, Methodology, and Future Directions. arXiv:2510.12040
- Sanyam Kapoor, Nate Gruver, Manley Roberts, Katherine Collins, Arka Pal, Umang Bhatt, Adrian Weller, Samuel Dooley, Micah Goldblum, and Andrew Gordon Wilson. 2024. Large Language Models Must Be Taught to Know What They Don't Know. In *NeurIPS 2024*. arXiv:2406.08391
- Timo Kaufmann, Paul Weng, Viktor Bengs, and Eyke Hüllermeier. 2023. A Survey of Reinforcement Learning from Human Feedback. arXiv:2312.14925 TMLR.
- Navdeep Kaur, Lachlan McPheat, Alessandra Russo, Anthony G Cohn, and Pranava Madhyastha. 2025. An Empirical Study of Conformal Prediction in LLM with ASP Scaffolds for Robust Reasoning. arXiv:2503.05439
- Ramneet Kaur, Colin Samplawski, Adam D. Cobb, Anirban Roy, Brian Matejek, Manoj Acharya, Daniel Elenius, Alexander M. Berenbeim, John A. Pavlik, Nathaniel D. Bastian, and Susmit Jha. 2024. Addressing Uncertainty in LLMs to Enhance Reliability in Generative AI. arXiv:2411.02381
- Yusong Ke, Hongru Lin, Yuting Ruan, Junya Tang, and Li Li. 2025. Correctness Coverage Evaluation for Medical Multiple-Choice Question Answering Based on the Enhanced Conformal Prediction Framework. *Mathematics* (2025). <https://doi.org/10.3390/math13091538> arXiv:2503.05505
- Reza Khanmohammadi, Erfan Miah, Mehra Mardikoraem, Simerjot Kaur, Ivan Brugere, Charese H. Smiley, Kundan Thind, and Mohammad M. Ghassemi. 2025. Calibrating LLM Confidence by Probing Perturbed Representation Stability. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2025.emnlp-main.530> arXiv:2505.21772
- Sunnie S. Y. Kim, Q. Vera Liao, Mihaela Vorvoreanu, Stephanie Ballard, and Jennifer Wortman Vaughan. 2024. "I'm Not Sure, But...": Examining the Impact of Large Language Models' Uncertainty Expression on User Reliance and Trust. <https://doi.org/10.1145/3630106.3658941> arXiv:2405.00623 FAccT 2024.
- Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. 2025. AbstentionBench: Reasoning LLMs Fail on Unanswerable Questions. arXiv:2506.09038
- Klaus-Rudolf Kladny, Bernhard Schölkopf, and Michael Muehlebach. 2024. Conformal Generative Modeling with Improved Sample Efficiency through Sequential Greedy Filtering. In *ICLR 2025*. arXiv:2410.01660
- Rishik Kondadadi and John E. Ortega. 2026. L2D-Clinical: Learning to Defer for Adaptive Model Selection in Clinical Text Classification. arXiv:2604.13285
- Lingkai Kong, Haoming Jiang, Yuchen Zhuang, Jie Lyu, Tuo Zhao, and Chao Zhang. 2020. Calibrated Language Model Fine-Tuning for In- and Out-of-Distribution Data. In *EMNLP 2020*. arXiv:2010.11506
- Terry Koo, Frederick Liu, and Luheng He. 2024. Automata-based constraints for language model decoding. arXiv:2407.08103 COLM 2024.
- Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. 2024. Semantic Entropy Probes: Robust and Cheap Hallucination Detection in LLMs. arXiv:2406.15927
- Vasily Kostumov, Bulat Nutfullin, Oleg Pilipenko, and Eugene Ilyushin. 2024. Uncertainty-Aware Evaluation for Vision-Language Models. arXiv:2402.14418
- Ranganath Krishnan, Piyush Khanna, and Omesh Tickoo. 2024. Enhancing Trust in Large Language Models via Uncertainty-Calibrated Fine-Tuning. In *ICLR-W 2026*. arXiv:2412.02904
- Maya Kruse, Majid Afshar, Saksham Khatwani, Anoop Mayampurath, Guanhua Chen, and Yanjun Gao. 2025. Simple Yet Effective: An Information-Theoretic Approach to Multi-LLM Uncertainty Quantification. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2507.07236
- Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. 2023. Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation. In *International Conference on Learning Representations (ICLR)*. arXiv:2302.09664
- Meelis Kull, Miquel Perello-Nieto, Markus Kängsepp, Telmo Silva Filho, Hao Song, and Peter Flach. 2019. Beyond temperature scaling: Obtaining well-calibrated multiclass probabilities with Dirichlet calibration. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:1910.12656

- Ananya Kumar, Percy Liang, and Tengyu Ma. 2019. Verified Uncertainty Calibration. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:1909.10155
- Abhishek Kumar, Robert Morabito, Sanzhar Umbet, Jad Kabbara, and Ali Emami. 2024. Confidence Under the Hood: An Investigation into the Confidence-Probability Alignment in Large Language Models. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. arXiv:2405.16282
- Aviral Kumar and Sunita Sarawagi. 2019. Calibration of Encoder Decoder Models for Neural Machine Translation. arXiv:1903.00802
- Bhawesh Kumar, Charlie Lu, Gauri Gupta, Anil Palepu, David Bellamy, Ramesh Raskar, and Andrew Beam. 2023. Conformal Prediction with Large Language Models for Multi-Choice Question Answering. arXiv:2305.18404
- Dharshan Kumaran, Arthur Conmy, Federico Barbero, Simon Osindero, Viorica Patraucean, and Petar Velicković. 2026. How do LLMs Compute Verbal Confidence. arXiv:2603.17839
- Bruce Kuwahara, Chen-Yuan Lin, Xiao Shi Huang, Kin Kwan Leung, Jullian Arta Yapeter, Ilya Stanevich, Felipe Perez, and Jesse C. Cresswell. 2025. Document Summarization with Conformal Importance Guarantees. In *NeurIPS 2025*. arXiv:2509.20461
- Chungpa Lee, Thomas Zeng, Jongwon Jeong, Jy yong Sohn, and Kangwook Lee. 2026b. How to Correctly Report LLM-as-a-Judge Evaluations. arXiv:2511.21140 ICML 2026.
- Dongryeol Lee, Yerin Hwang, Yongil Kim, Joonsuk Park, and Kyomin Jung. 2024a. Are LLM-Judges Robust to Expressions of Uncertainty? Investigating the effect of Epistemic Markers on LLM-based Evaluation. In *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*. arXiv:2410.20774
- Ivan Yee Lee, Loris D'Antoni, and Taylor Berg-Kirkpatrick. 2026a. The Format Tax. arXiv:2604.03616
- Minjae Lee, Yoonjae Jung, and Sangdon Park. 2025. Online Conformal Abstention for Factuality Control Under Adversarial Bandit Feedback. arXiv:2506.14067
- Minjae Lee, Kyungmin Kim, Taesoo Kim, and Sangdon Park. 2024b. Selective Generation for Controllable Language Models. In *NeurIPS 2024*. <https://doi.org/10.52202/079017-1600> arXiv:2307.09254
- Jixuan Leng, Chengsong Huang, Banghua Zhu, and Jiaxin Huang. 2025. Taming Overconfidence in LLMs: Reward Calibration in RLHF. In *International Conference on Learning Representations (ICLR)*. arXiv:2410.09724
- Alexander K. Lew, Tan Zhi-Xuan, Gabriel Grand, and Vikash K. Mansinghka. 2023. Sequential Monte Carlo Steering of Large Language Models using Probabilistic Programs. arXiv:2306.03081
- Chen Li, Xiaoling Hu, Songzhu Zheng, Jiawei Zhou, and Chao Chen. 2026a. ORCE: Order-Aware Alignment of Verbalized Confidence in Large Language Models. arXiv:2605.12446
- Changcheng Li, Jiancan Wu, Hengheng Zhang, Zhengsu Chen, Guo An, Junxiang Qiu, Xiang Wang, and Qi Tian. 2026d. Confidence Before Answering: A Paradigm Shift for Efficient LLM Uncertainty Estimation. arXiv:2603.05881
- Chengzu Li, Han Zhou, Goran Glavaš, Anna Korhonen, and Ivan Vulić. 2025h. Large Language Models are Miscalibrated In-Context Learners. In *Findings of the Association for Computational Linguistics: ACL*. <https://doi.org/10.18653/v1/2025.findings-acl.603> arXiv:2312.13772
- Dawei Li, Bohan Jiang, Liangjie Huang, Alimohammad Beigi, Chengshuai Zhao, Zhen Tan, Amrita Bhattacharjee, Yuxuan Jiang, Canyu Chen, Tianhao Wu, Kai Shu, Lu Cheng, and Huan Liu. 2024c. From Generation to Judgment: Opportunities and Challenges of LLM-as-a-judge. arXiv:2411.16594
- Haitao Li, Qian Dong, Junjie Chen, Huixue Su, Yujia Zhou, Qingyao Ai, Ziyi Ye, and Yiqun Liu. 2024b. LLMs-as-Judges: A Comprehensive Survey on LLM-based Evaluation Methods. arXiv:2412.05579
- Jiaqi Li, Yixuan Tang, and Yi Yang. 2025c. Know the Unknown: An Uncertainty-Sensitive Method for LLM Instruction Tuning. In *Findings ACL 2025*. <https://doi.org/10.18653/v1/2025.findings-acl.153> arXiv:2406.10099
- Loka Li, Zhenhao Chen, Guangyi Chen, Yixuan Zhang, Yusheng Su, Eric Xing, and Kun Zhang. 2024a. Confidence Matters: Revisiting Intrinsic Self-Correction Capabilities of Large Language Models. arXiv:2402.12563
- Moxin Li, Wenjie Wang, Fuli Feng, Fengbin Zhu, Qifan Wang, and Tat-Seng Chua. 2024e. Think Twice Before Trusting: Self-Detection for Large Language Models through Comprehensive Answer Reflection. In *Findings of the Association for Computational Linguistics: EMNLP*. arXiv:2403.09972
- Moxin Li, Yong Zhao, Wenxuan Zhang, Shuaiyi Li, Wenya Xie, See-Kiong Ng, Tat-Seng Chua, and Yang Deng. 2025g. Knowledge Boundary of Large Language Models: A Survey. arXiv:2412.12472 ACL 2025.
- Shuo Li, Sangdon Park, Insup Lee, and Osbert Bastani. 2024d. TRAQ: Trustworthy Retrieval Augmented Question Answering via Conformal Prediction. In *NAACL 2024*. <https://doi.org/10.18653/v1/2024.naacl-long.210> arXiv:2307.04642
- Siheng Li, Cheng Yang, Taiqiang Wu, Chufan Shi, Yuji Zhang, Xinyu Zhu, Zesen Cheng, Deng Cai, Mo Yu, Lemao Liu, Jie Zhou, Yujia Yang, Ngai Wong, Xixin Wu, and Wai Lam. 2025e. A Survey on the Honesty of Large Language Models. arXiv:2409.18786 TMLR.
- Tianyi Li, Mingda Chen, Bowei Guo, and Zhiqiang Shen. 2025a. A Survey on Diffusion Language Models. arXiv:2508.10875
- Xiaomin Li, Zhou Yu, Ziji Zhang, Yingying Zhuang, Swair Shah, Narayanan Sadagopan, and Anurag Beniwal. 2026e. Semantic Volume: Quantifying and Detecting both External and Internal Uncertainty in LLMs. In *AAAI Conference on*

Artificial Intelligence. arXiv:2502.21239

- Yongmin Li, Jia Li, Ge Li, and Zhi Jin. 2026b. AdapTrack: Constrained Decoding without Distorting LLM’s Output Intent. arXiv:2510.17376 ICSE 2026.
- Yubo Li, Yidi Miao, Ramayya Krishnan, and Rema Padman. 2026c. JEV-as-a-Judge: Accept When Confident, Escalate When Unsure. arXiv:2609.26550
- Yinghao Li, Rushi Qiang, Lama Moukheiber, and Chao Zhang. 2025b. Language Model Uncertainty Quantification with Attention Chain. arXiv:2503.19168
- Yibo Li, Miao Xiong, Jiaying Wu, and Bryan Hooi. 2025d. ConfTuner: Training Large Language Models to Express Their Confidence Verbally. In *NeurIPS 2025*. <https://doi.org/10.52202/085713-1783> arXiv:2508.18847
- Yiwei Li, Peiwen Yuan, Shaoxiong Feng, Boyuan Pan, Xinglin Wang, Bin Sun, Heda Wang, and Kan Li. 2024f. Escape Sky-high Cost: Early-stopping Self-Consistency for Multi-step Reasoning. arXiv:2401.10480 ICLR 2024.
- Zhong-Zhi Li, Duzhen Zhang, Ming-Liang Zhang, Jiaxin Zhang, Zengyan Liu, Yuxuan Yao, Haotian Xu, Junhao Zheng, Pei-Jie Wang, Xiuyi Chen, Yingying Zhang, Fei Yin, Jiahua Dong, Zhiwei Li, Bao-Long Bi, Ling-Rui Mei, Junfeng Fang, Xiao Liang, Zhijiang Guo, Le Song, and Cheng-Lin Liu. 2025f. From System 1 to System 2: A Survey of Reasoning Large Language Models. arXiv:2502.17419
- Kaiqu Liang, Zixu Zhang, and Jaime Fernández Fisac. 2024b. Introspective Planning: Aligning Robots’ Uncertainty with Inherent Task Ambiguity. In *NeurIPS 2024*. <https://doi.org/10.52202/079017-2299> arXiv:2402.06529
- Yuxin Liang, Zhuoyang Song, Hao Wang, and Jiaxing Zhang. 2024a. Learning to Trust Your Feelings: Leveraging Self-awareness in LLMs for Hallucination Mitigation. In *KnowledgeNLP-W 2024*. <https://doi.org/10.18653/v1/2024.knowledgenlp-1.4> arXiv:2401.15449
- Moule Lin, Qizhen Lan, Shuhao Guan, Weipeng Jing, Jiexin Fan, David Gregg, and Goetz Botterweck. 2026. Uncertainty Quantification for LLM Agents: A Taxonomy, an Evaluation Protocol, and an Empirical Study. arXiv:2609.07395
- Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. Teaching Models to Express Their Uncertainty in Words. *Transactions on Machine Learning Research (TMLR)* (2022). arXiv:2205.14334
- Zhen Lin, Shubhendu Trivedi, and Jimeng Sun. 2024. Generating with Confidence: Uncertainty Quantification for Black-box Large Language Models. *Transactions on Machine Learning Research* (2024). arXiv:2305.19187
- Chen Ling, Xujiang Zhao, Xuchao Zhang, Wei Cheng, Yanchi Liu, Yiyu Sun, Mika Oishi, Takao Osaki, Katsushi Matsuda, Jie Ji, Guangji Bai, Liang Zhao, and Haifeng Chen. 2024. Uncertainty Quantification for In-Context Learning of Large Language Models. In *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*. arXiv:2402.10189
- Gabrielle Kaili-May Liu, Avi Caciularu, Gal Yona, Idan Szpektor, and Arman Cohan. 2026a. Reinforcement Learning with Metacognitive Feedback Elicits Faithful Uncertainty Expression in LLMs. arXiv:2606.32032
- Gabrielle Kaili-May Liu and Arman Cohan. 2026. Can LLMs Use Linguistic Uncertainty Markers to Reliably Reflect Intrinsic Confidence? arXiv:2605.28778
- Gabrielle Kaili-May Liu, Areeb Gani, Jacqueline Lu, Jordan Thomas, Mark Steyvers, and Arman Cohan. 2026b. Metacognition in LLMs: Foundations, Progress, and Opportunities. arXiv:2607.11881
- Gabrielle Kaili-May Liu, Gal Yona, Avi Caciularu, Idan Szpektor, Tim G. J. Rudner, and Arman Cohan. 2025d. MetaFaith: Faithful Natural Language Uncertainty Expression in LLMs. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2505.24858
- Hongfu Liu, Hengguan Huang, Xiangming Gu, Hao Wang, and Ye Wang. 2024b. On Calibration of LLM-based Guard Models for Reliable Content Moderation. In *International Conference on Learning Representations (ICLR)*. arXiv:2410.10414
- Haotian Liu, Shuo Wang, and Hongteng Xu. 2025b. C²GSPG: Confidence-calibrated Group Sequence Policy Gradient towards Self-aware Reasoning. arXiv:2509.23129
- Jiayu Liu, Qing Zong, Weiqi Wang, and Yangqiu Song. 2025e. Revisiting Epistemic Markers in Confidence Estimation: Can Markers Accurately Reflect Large Language Models’ Uncertainty?. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. arXiv:2505.24778
- Kevin Liu, Stephen Casper, Dylan Hadfield-Menell, and Jacob Andreas. 2023a. Cognitive Dissonance: Why Do Language Model Outputs Disagree with Internal Representations of Truthfulness?. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2023.emnlp-main.291> arXiv:2312.03729
- Keliang Liu, Dingkan Yang, Ziyun Qian, Weijie Yin, Yuchi Wang, Hongsheng Li, Jun Liu, Peng Zhai, Yang Liu, and Lihua Zhang. 2025c. Reinforcement Learning Meets Large Language Models: A Survey of Advancements and Applications Across the LLM Lifecycle. arXiv:2509.16679
- Linyu Liu, Yu Pan, Xiaocheng Li, and Quanting Chen. 2024d. Uncertainty Estimation and Quantification for LLMs: A Simple Supervised Approach. arXiv:2404.15993
- Shudong Liu, Zhaocong Li, Xuebo Liu, Runzhe Zhan, Derek F. Wong, Lidia S. Chao, and Min Zhang. 2024c. Can LLMs Learn Uncertainty on Their Own? Expressing Uncertainty Effectively in A Self-Training Manner. In *EMNLP 2024*. <https://doi.org/10.18653/v1/2024.emnlp-main.1205>

- Terrance Liu and Zhiwei Steven Wu. 2024. Multi-group Uncertainty Quantification for Long-form Text Generation. In *UAI 2025*. arXiv:2407.21057
- Xin Liu, Farima Fatahi Bayat, and Lu Wang. 2024a. Enhancing Language Model Factuality via Activation-Based Confidence Calibration and Guided Decoding. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2024.emnlp-main.583> arXiv:2406.13230
- Xiaou Liu, Tiejun Chen, Longchao Da, Chacha Chen, Zhen Lin, and Hua Wei. 2025a. Uncertainty Quantification and Confidence Calibration in Large Language Models: A Survey. arXiv:2503.15850 KDD 2025.
- Xin Liu, Muhammad Khalifa, and Lu Wang. 2023c. LitCab: Lightweight Language Model Calibration over Short- and Long-form Responses. In *International Conference on Learning Representations (ICLR)*. arXiv:2310.19208
- Xin Liu and Lu Wang. 2026. Think Through Uncertainty: Improving Long-Form Generation Factuality via Reasoning Calibration. arXiv:2604.12046
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023b. G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. arXiv:2303.16634
- Yuxuan Liu, Tianchi Yang, Shaohan Huang, Zihan Zhang, Haizhen Huang, Furu Wei, Weiwei Deng, Feng Sun, and Qi Zhang. 2023d. Calibrating LLM-Based Evaluator. arXiv:2309.13308
- Adian Liusie, Potsawee Manakul, and Mark J. F. Gales. 2023. Mitigating Word Bias in Zero-shot Prompt-based Classifiers. In *Findings of the Association for Computational Linguistics: IJCNLP-AACL*. <https://doi.org/10.18653/v1/2023.findings-ijcnlp.29> arXiv:2309.04992
- Xingzhou Lou, Dong Yan, Wei Shen, Yuzi Yan, Jian Xie, and Junge Zhang. 2024. Uncertainty-aware Reward Model: Teaching Reward Models to Know What is Unknown. arXiv:2410.00847
- João Loula, Benjamin LeBrun, Li Du, Ben Lipkin, Clemente Pasti, Gabriel Grand, Tianyu Liu, Yahya Emara, Marjorie Freedman, Jason Eisner, Ryan Cotterell, Vikash Mansinghka, Alexander K. Lew, Tim Vieira, and Timothy J. O'Donnell. 2025. Syntactic and Semantic Control of Large Language Models via Sequential Monte Carlo. arXiv:2504.13139
- Charles Lovering, Michael Krumdtick, Viet Dac Lai, Seth Ebner, Nilesh Kumar, Varshini Reddy, Rik Koncel-Kedziorski, and Chris Tanner. 2025. Language Model Probabilities are Not Calibrated in Numeric Contexts. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2025.acl-long.1417> arXiv:2410.16007
- Yi-Long Lu. 2026. Diagnosing Correctness Probes under Self-Judgement Confounding. arXiv:2607.16799
- Beier Luo, Cheng Wang, Hongxin Wei, Sharon Li, and Xuefeng Du. 2026. Unlocking the Pre-Trained Model as a Dual-Alignment Calibrator for Post-Trained LLMs. arXiv:2601.04277
- Beier Luo, Shuoyuan Wang, Sharon Li, and Hongxin Wei. 2025. Your Pre-trained LLM is Secretly an Unsupervised Confidence Calibrator. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2505.16690
- Qing Lyu, Kumar Shridhar, Chaitanya Malaviya, Li Zhang, Yanai Elazar, Niket Tandon, Marianna Apidianaki, Mrinmaya Sachan, and Chris Callison-Burch. 2025. Calibrating Large Language Models with Sample Consistency. In *AAAI Conference on Artificial Intelligence*. arXiv:2402.13904
- Huan Ma, Jiadong Pan, Jing Liu, Yan Chen, Joey Tianyi Zhou, Guangyu Wang, Qinghua Hu, Hua Wu, Changqing Zhang, and Haifeng Wang. 2025. Semantic Energy: Detecting LLM Hallucination Beyond Entropy. arXiv:2508.14496
- Rachel Ma, Dylan Hadfield-Menell, and Kristjan Greenewald. 2026a. Distributional Process Reward Models: Calibrated Prediction of Future Rewards via Conditional Optimal Transport. arXiv:2605.06785
- Zhengzhao Ma, Xueru Wen, Boxi Cao, Yaojie Lu, Hongyu Lin, Jinglin Yang, Min He, Xianpei Han, and Le Sun. 2026b. Decoupling Reasoning and Confidence: Resurrecting Calibration in Reinforcement Learning from Verifiable Rewards. In *ICML 2026*. arXiv:2603.09117
- Matéo Mahaut, Laura Aina, Paula Czarnowska, Momchil Hardalov, Thomas Müller, and Lluís Màrquez. 2024. Factual Confidence of LLMs: on Reliability and Robustness of Current Estimators. In *Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2024.acl-long.250> arXiv:2406.13415
- Andrey Malinin and Mark Gales. 2020. Uncertainty Estimation in Autoregressive Structured Prediction. arXiv:2002.07650
- Potsawee Manakul, Adian Liusie, and Mark J. F. Gales. 2023. SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2303.08896
- Ali Marashian, Alexis Palmer, and Katharina von der Wense. 2026. Speaking in Self-Assessing Tongues: On the Verbalized Confidence of LLMs in Machine Translation. arXiv:2606.17234
- Zhiteng Mei, Christina Zhang, Tenny Yin, Justin Lidard, Ola Shorinwa, and Anirudha Majumdar. 2025. Reasoning about Uncertainty: Do Reasoning Models Know When They Don't Know? arXiv:2506.18183
- Sabrina J. Mielke, Arthur Szlam, Emily Dinan, and Y-Lan Boureau. 2022. Reducing conversational agents' overconfidence through linguistic calibration. *Transactions of the Association for Computational Linguistics* (2022). https://doi.org/10.1162/tacl_a_00494 arXiv:2012.14983

- Mohamad Amin Mohamadi, Tianhao Wang, and Zhiyuan Li. 2025. Honesty over Accuracy: Trustworthy Language Models through Reinforced Hesitation. arXiv:2511.11500
- Christopher Mohri and Tatsunori Hashimoto. 2024. Language Models with Conformal Factuality Guarantees. In *ICML 2024*. arXiv:2402.10978
- Yasmin Moslem and John D. Kelleher. 2026. Dynamic Model Routing and Cascading for Efficient LLM Inference: A Survey. arXiv:2603.04445 TMLR.
- multimodalart. 2026. Jev Decision Index. Hugging Face Space, <https://huggingface.co/spaces/multimodalart/jev-decision-index>. Release 0.2.1, data of 27 September 2026.
- Niels Mündler, Jingxuan He, Slobodan Jenko, and Martin Vechev. 2023. Self-contradictory Hallucinations of Large Language Models: Evaluation, Detection and Mitigation. arXiv:2305.15852
- Niels Mündler, Jingxuan He, Hao Wang, Koushik Sen, Dawn Song, and Martin Vechev. 2025. Type-Constrained Code Generation with Language Models. <https://doi.org/10.1145/3729274> arXiv:2504.09246 PACMPL (PLDI 2025).
- Nariman Naderi, Zahra Atf, Peter R Lewis, Aref Mahjoub far, Seyed Amir Ahmad Safavi-Naini, and Ali Soroush. 2025. Evaluating Prompt Engineering Techniques for Accuracy and Confidence Elicitation in Medical LLMs. In *International Workshop on EXplainable, Trustworthy, and Responsible AI and Multi-Agent Systems (EXTRAAMAS)*. arXiv:2506.00072
- Preetum Nakkiran, Arwen Bradley, Adam Goliński, Eugene Ndiaye, Michael Kirchhof, and Sinead Williamson. 2025. Trained on Tokens, Calibrated on Concepts: The Emergence of Semantic Calibration in LLMs. arXiv:2511.04869
- Harikrishna Narasimhan, Wittawat Jitkrittum, Ankit Singh Rawat, Seungyeon Kim, Neha Gupta, Aditya Krishna Menon, and Sanjiv Kumar. 2024. Faster Cascades via Speculative Decoding. arXiv:2405.19261
- Dang Nguyen, Ali Payani, and Baharan Mirzasoleiman. 2025. Beyond Semantic Entropy: Boosting LLM Uncertainty Quantification with Pairwise Semantic Similarity. In *Findings of the Association for Computational Linguistics: ACL 2025*. arXiv:2506.00245
- Shiyu Ni, Keping Bi, Jiafeng Guo, and Xueqi Cheng. 2024a. When Do LLMs Need Retrieval Augmentation? Mitigating LLMs’ Overconfidence Helps Retrieval Augmentation. In *Findings of the Association for Computational Linguistics: ACL*. arXiv:2402.11457
- Shiyu Ni, Keping Bi, Lulu Yu, and Jiafeng Guo. 2024b. Are Large Language Models More Honest in Their Probabilistic or Verbalized Confidence? arXiv:2408.09773
- Alexander Nikitin, Jannik Kossen, Yarin Gal, and Pekka Marttinen. 2024. Kernel Language Entropy: Fine-grained Uncertainty Quantification for LLMs from Semantic Similarities. arXiv:2405.20003
- Ruijia Niu, Dongxia Wu, Rose Yu, and Yi-An Ma. 2024. Functional-level Uncertainty Quantification for Calibrated Fine-tuning on LLMs. arXiv:2410.06431
- Sima Noorani, Shayan Kiyani, George Pappas, and Hamed Hassani. 2025. Conformal Prediction Beyond the Seen: A Missing Mass Perspective for Uncertainty Quantification in Generative Models. In *NeurIPS 2025*. arXiv:2506.05497
- Stephen Obadinma and Xiaodan Zhu. 2025. On the Robustness of Verbal Confidence of LLMs in Adversarial Attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2507.06489
- Changdae Oh, Seongheon Park, To Eun Kim, Jiatong Li, Wendi Li, Samuel Yeh, Xuefeng Du, Hamed Hassani, Paul Bogdan, Dawn Song, and Sharon Li. 2026. Uncertainty Quantification in LLM Agents: Foundations, Emerging Challenges, and Opportunities. arXiv:2602.05073 ACL 2026.
- Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M Waleed Kadous, and Ion Stoica. 2024. RouteLLM: Learning to Route LLMs with Preference Data. arXiv:2406.18665
- OpenAI. 2023. GPT-4 Technical Report. arXiv:2303.08774
- Hadas Orgad, Michael Tokar, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Kotek, and Yonatan Belinkov. 2024. LLMs Know More Than They Show: On the Intrinsic Representation of LLM Hallucinations. arXiv:2410.02707
- Myle Ott, Michael Auli, David Grangier, and Marc’Aurelio Ranzato. 2018. Analyzing Uncertainty in Neural Machine Translation. In *Proceedings of the International Conference on Machine Learning (ICML)*. arXiv:1803.00047
- William Overman and Mohsen Bayati. 2025. Conformal Arbitrage: Risk-Controlled Balancing of Competing Objectives in Language Models. In *NeurIPS 2025*. arXiv:2506.00911
- Licheng Pan, Haocheng Yang, Haoxuan Li, Yichen Sun, Yunsheng Lu, Shijian Wang, Lei Shen, Yuan Lu, Zhixuan Chu, and Hao Wang. 2026a. Uncertainty-Aware Reward Modeling for Stable RLHF. arXiv:2606.19818
- Muyu Pan, Shu Zhao, Nan Zhang, Philip Shin, Varun Parekh, Vijaykrishnan Narayanan, and Rui Zhang. 2026b. TIAR: Trajectory-Informed Advantage Reweighting for LLM Abstention Learning. arXiv:2605.25850
- Kanghee Park, Jiayu Wang, Taylor Berg-Kirkpatrick, Nadia Polikarpova, and Loris D’Antoni. 2024. Grammar-Aligned Decoding. arXiv:2405.21047 NeurIPS 2024.
- Kanghee Park, Timothy Zhou, and Loris D’Antoni. 2025b. Flexible and Efficient Grammar-Constrained Decoding. arXiv:2502.05111
- Young-Jin Park, Kristjan Greenewald, Kaveh Alim, Hao Wang, and Navid Azizan. 2025a. Know What You Don’t Know: Uncertainty Calibration of Process Reward Models. In *NeurIPS 2025*. arXiv:2506.09338

- Yudi Pawitan and Chris Holmes. 2024. Confidence in the Reasoning of Large Language Models. arXiv:2412.15296
- Benjamin Plaut, Nguyen X. Khanh, and Tu Trinh. 2024. Probabilities of Chat LLMs Are Miscalibrated but Still Predict Correctness on Multiple-Choice Q&A. *Transactions on Machine Learning Research* (2024). arXiv:2402.13213
- Gabriel Poesia, Oleksandr Polozov, Vu Le, Ashish Tiwari, Gustavo Soares, Christopher Meek, and Sumit Gulwani. 2022. SynchroMesh: Reliable code generation from pre-trained language models. arXiv:2201.11227
- Gwenyth Portillo Wightman, Alexandra Delucia, and Mark Dredze. 2023. Strength in Numbers: Estimating Confidence of Large Language Models by Prompt Agreement. In *Proceedings of the 3rd Workshop on Trustworthy Natural Language Processing (TrustNLP 2023)*. Association for Computational Linguistics, 326–362. <https://doi.org/10.18653/v1/2023.trustnlp-1.28>
- Irina Proskurina, Mayank Kumar, and Oyindolapo O. Komolafe. 2026. Are You Sure You’re Sure? On the Impact of Instruction Tuning on Confidence and Lexical Diversity. arXiv:2608.13430
- Xin Qiu and Risto Miikkulainen. 2024. Semantic Density: Uncertainty Quantification for Large Language Models through Confidence Measurement in Semantic Space. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2405.13845
- Victor Quach, Adam Fisch, Tal Schuster, Adam Yala, Jae Ho Sohn, Tommi S. Jaakkola, and Regina Barzilay. 2023. Conformal Language Modeling. In *ICLR 2024*. arXiv:2306.10193
- Bhaktipriya Radharapu, Eshika Saxena, Kenneth Li, Chenxi Whitehouse, Adina Williams, and Nicola Cancedda. 2025. Calibrating LLM Judges: Linear Probes for Fast and Reliable Uncertainty Estimation. arXiv:2512.22245
- Guillem Ramírez, Alexandra Birch, and Ivan Titov. 2024. Optimising Calls to Large Language Models with Uncertainty-Based Two-Tier Selection. arXiv:2405.02134 COLM 2024.
- Neil Rathi, Dan Jurafsky, and Kaitlyn Zhou. 2025. Humans overrely on overconfident language models, across languages. arXiv:2507.06306
- Shauli Ravfogel, Yoav Goldberg, and Jacob Goldberger. 2023. Conformal Nucleus Sampling. In *Findings ACL 2023*. <https://doi.org/10.18653/v1/2023.findings-acl.3> arXiv:2305.02633
- Srijith Ravikumar. 2026. The Missing “I Don’t Know”: Why Three Reasoning-Reliability Findings Converge on Calibrated Abstention. arXiv:2609.17686
- Avinash Reddy, Thayne T. Walker, James S. Ide, and Amrit Singh Bedi. 2026. The Hidden Cost of Structured Generation in LLMs: Draft-Conditioned Constrained Decoding. arXiv:2603.03305
- Yuval Reif and Roy Schwartz. 2024. Beyond Performance: Quantifying and Mitigating Label Bias in LLMs. In *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*. <https://doi.org/10.18653/v1/2024.naacl-long.378> arXiv:2405.02743
- Allen Z. Ren, Jaden Clark, Anushri Dixit, Masha Itkina, Anirudha Majumdar, and Dorsa Sadigh. 2024. Explore until Confident: Efficient Exploration for Embodied Question Answering. In *RSS 2024*. arXiv:2403.15941
- Allen Z. Ren, Anushri Dixit, Alexandra Bodrova, Sumeet Singh, Stephen Tu, Noah Brown, Peng Xu, Leila Takayama, Fei Xia, Jake Varley, Zhenjia Xu, Dorsa Sadigh, Andy Zeng, and Anirudha Majumdar. 2023a. Robots That Ask For Help: Uncertainty Alignment for Large Language Model Planners. In *CoRL 2023*. arXiv:2307.01928
- Jie Ren, Yao Zhao, Tu Vu, Peter J. Liu, and Balaji Lakshminarayanan. 2023b. Self-Evaluation Improves Selective Generation in Large Language Models. arXiv:2312.09300
- Mauricio Rivera, Jean-François Godbout, Reihaneh Rabbany, and Kellin Pelrine. 2024. Combining Confidence Elicitation and Sample-based Methods for Uncertainty Quantification in Misinformation Mitigation. arXiv:2401.08694
- Maxon Rubin-Toles, Maya Gambhir, Keshav Ramji, Aaron Roth, and Surbhi Goel. 2025. Conformal Language Model Reasoning with Coherent Factuality. In *ICLR 2025*. arXiv:2505.17126
- Hyun Ryu and Eric Kim. 2024. Closer Look at Efficient Inference Methods: A Survey of Speculative Decoding. arXiv:2411.13157
- Mario Sanz-Guerrero, Manuel Mager, and Katharina von der Wense. 2026. Large Language Models Are Overconfident in Their Own Responses. In *Findings of the Association for Computational Linguistics: ACL*. arXiv:2606.03437
- Torsten Scholach, Nathan Schucher, and Dzmitry Bahdanau. 2021. PICARD: Parsing Incrementally for Constrained Auto-Regressive Decoding from Language Models. arXiv:2109.05093 EMNLP 2021.
- Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Metzler. 2022. Confident Adaptive Language Modeling. arXiv:2207.07061 NeurIPS 2022.
- Tal Schuster, Adam Fisch, Tommi Jaakkola, and Regina Barzilay. 2021. Consistent Accelerated Inference via Confident Adaptive Transformers. In *EMNLP 2021*. <https://doi.org/10.18653/v1/2021.emnlp-main.406> arXiv:2104.08803
- Hooman Shahrokhi, Devjeet Raj Roy, Yan Yan, Venera Arnaudova, and Janaradhan Rao Doppa. 2025. Conformal Prediction Sets for Deep Generative Models via Reduction to Conformal Regression. In *UAI 2025*. arXiv:2503.10512
- Arindam Sharma and Cristina David. 2025. Assessing Correctness in LLM-Based Code Generation via Uncertainty Estimation. arXiv:2502.11620
- Maohao Shen, Subhro Das, Kristjan Greenewald, Prasanna Sattigeri, Gregory Wornell, and Soumya Ghosh. 2024. Thermometer: Towards Universal Calibration for Large Language Models. In *ICML*. arXiv:2403.08819

- Huanxin Sheng, Xinyi Liu, Hangfeng He, Jieyu Zhao, and Jian Kang. 2025. Analyzing Uncertainty of LLM-as-a-Judge: Interval Evaluations with Conformal Prediction. In *EMNLP 2025*. <https://doi.org/10.18653/v1/2025.emnlp-main.569> arXiv:2509.18658
- Ola Shorinwa, Zhiting Mei, Justin Lidard, Allen Z. Ren, and Anirudha Majumdar. 2025. A Survey on Uncertainty Quantification of Large Language Models: Taxonomy, Open Research Challenges, and Future Directions. arXiv:2412.05563 ACM Computing Surveys.
- Connor Shorten, Charles Piere, Thomas Benjamin Smith, Erika Cardenas, Akanksha Sharma, John Trengrove, and Bob van Luijt. 2024. StructuredRAG: JSON Response Formatting with Large Language Models. arXiv:2408.11061
- Chenglei Si, Navita Goyal, Sherry Tongshuang Wu, Chen Zhao, Shi Feng, Hal Daumé, and Jordan Boyd-Graber. 2024. Large Language Models Help Humans Verify Truthfulness – Except When They Are Convincingly Wrong. arXiv:2310.12558 NAACL 2024.
- Chenglei Si, Chen Zhao, Sewon Min, and Jordan Boyd-Graber. 2022. Re-Examining Calibration: The Case of Question Answering. In *Findings of the Association for Computational Linguistics: EMNLP*. <https://doi.org/10.18653/v1/2022.findings-emnlp.204> arXiv:2205.12507
- Anthony Sicilia, Mert Inan, and Malihe Alikhani. 2024. Accounting for Sycophancy in Language Model Uncertainty Estimation. arXiv:2410.14746
- Damien Sileo and Marie-Francine Moens. 2022. Probing neural language models for understanding of words of estimative probability. In *Proceedings of the Joint Conference on Lexical and Computational Semantics ("SEM")*. arXiv:2211.03358
- Aniket Kumar Singh, Suman Devkota, Bishal Lamichhane, Uttam Dhakal, and Chandra Dhakal. 2023. The Confidence-Competence Gap in Large Language Models: A Cognitive Study. arXiv:2309.16145
- Abhinav Kumar Singh, Harsha Vardhan Khurdula, Yoeven D Khemlani, and Vineet Agarwal. 2026a. The Structured Output Benchmark: A Multi-Source Benchmark for Evaluating Structured Output Quality in Large Language Models. arXiv:2604.25359
- Sadanand Singh, Allam Reddy, and Manan Chopra. 2026b. Verifiable Rewards for Calibrated Probabilistic Forecasting. arXiv:2607.00164
- Aviv Slobodkin, Omer Goldman, Avi Caciularu, Ido Dagan, and Shauli Ravfogel. 2023. The Curious Case of Hallucinatory (Un)answerability: Finding Truths in the Hidden States of Over-Confident Large Language Models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2023.emnlp-main.220> arXiv:2310.11877
- Duncan Soiffer, Steven Kolawole, and Virginia Smith. 2025. Semantic Agreement Enables Efficient Open-Ended LLM Cascades. arXiv:2509.21837 EMNLP 2025 Industry.
- Junru Song, Yimeng Hu, Yijing Chen, Huining Li, Qian Li, Lizhen Cui, and Yuntao Du. 2026. Delineating Knowledge Boundaries for Honest Large Vision-Language Models. In *LNCS*. https://doi.org/10.1007/978-3-032-37174-4_3 arXiv:2604.26419
- Linxin Song, Taiwei Shi, and Jieyu Zhao. 2025. The Hallucination Tax of Reinforcement Finetuning. In *Findings EMNLP 2025*. <https://doi.org/10.18653/v1/2025.findings-emnlp.112> arXiv:2505.13988
- Claudio Spiess, David Gros, Kunal Suresh Pai, Michael Pradel, Md Rafiqul Islam Rabin, Amin Alipour, Susmit Jha, Prem Devanbu, and Toufique Ahmed. 2024. Calibration and Correctness of Language Models for Code. In *Proceedings of the IEEE/ACM International Conference on Software Engineering (ICSE)*. arXiv:2402.02047
- Tejas Srinivasan, Jack Hessel, Tanmay Gupta, Bill Yuchen Lin, Yejin Choi, Jesse Thomason, and Khyathi Raghavi Chandu. 2024. Selective "Selective Prediction": Reducing Unnecessary Abstention in Vision-Language Reasoning. arXiv:2402.15610 Findings ACL 2024.
- Elias Stengel-Eskin, Peter Hase, and Mohit Bansal. 2024. LACIE: Listener-Aware Finetuning for Confidence Calibration in Large Language Models. arXiv:2405.21028
- Elias Stengel-Eskin and Benjamin Van Durme. 2023. Calibrated Interpretation: Confidence Estimation in Semantic Parsing. *Transactions of the Association for Computational Linguistics* (2023). https://doi.org/10.1162/tacl_a_00598 arXiv:2211.07443
- Mark Steyvers, Catarina Belem, and Padhraic Smyth. 2025a. Generalization of Fine-Tuned Uncertainty Communication and Metacognition in Large Language Models. arXiv:2510.05126
- Mark Steyvers, Heliodoro Tejeda, Aakriti Kumar, Catarina Belem, Sheer Karny, Xinyue Hu, Lukas Mayer, and Padhraic Smyth. 2025b. What Large Language Models Know and What People Think They Know. <https://doi.org/10.1038/s42256-024-00976-7> arXiv:2401.13835 Nat Mach Intell.
- Alessandro Stolfo, Ben Wu, Wes Gurnee, Yonatan Belinkov, Xingyi Song, Mrinmaya Sachan, and Neel Nanda. 2024. Confidence Regulation Neurons in Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2406.16254
- Joshua Strong, Emma Sun, Harry Rogers, Helen Higham, and Alison Noble. 2025. Learning to Defer: A Survey. <https://doi.org/10.5281/zenodo.17843044> Zenodo preprint (not peer reviewed).
- Jiayuan Su, Jing Luo, Hongwei Wang, and Lu Cheng. 2024a. API Is Enough: Conformal Prediction for Large Language Models Without Logit-Access. In *Findings EMNLP 2024*. <https://doi.org/10.18653/v1/2024.findings-emnlp.54> arXiv:2403.01216

- Weihang Su, Changyue Wang, Qingyao Ai, Yiran HU, Zhijing Wu, Yujia Zhou, and Yiqun Liu. 2024b. Unsupervised Real-Time Hallucination Detection based on the Internal States of Large Language Models. [arXiv:2403.06448](#)
- Yang Sui, Yu-Neng Chuang, Guanchu Wang, Jiamu Zhang, Tianyi Zhang, Jiayi Yuan, Hongyi Liu, Andrew Wen, Shaochen Zhong, Na Zou, Hanjie Chen, and Xia Hu. 2025. Stop Overthinking: A Survey on Efficient Reasoning for Large Language Models. [arXiv:2503.16419](#) TMLR 2025 (per ML Anthology).
- Fengfei Sun, Ningke Li, Kailong Wang, and Lorenz Goette. 2025. Large Language Models are overconfident and amplify human bias. [arXiv:2505.02151](#)
- Qiyao Sun, Xingming Li, Xixiang He, Ao Cheng, Xuanyu Ji, Hailun Lu, Runke Huang, and Qingyong Hu. 2026a. Efficient Hallucination Detection: Adaptive Bayesian Estimation of Semantic Entropy with Guided Semantic Exploration. In *AAAI Conference on Artificial Intelligence*. [arXiv:2603.22812](#)
- Xin Sun, Yue Su, Yifan Mo, Qingyu Meng, Yuxuan Li, Min Chen, Mengyuan Zhang, Saku Sugawara, Charlotte Gerritsen, Sander L. Koole, Koen Hindriks, and Jiahuan Pei. 2026b. Trust Stack for Mental Health AI: A Survey of Calibration across Human, Interaction, and Model Levels. [arXiv:2604.20166](#)
- Yi-Jyun Sun, Suvodip Dey, Dilek Hakkani-Tur, and Gokhan Tur. 2024. Confidence Estimation for LLM-Based Dialogue State Tracking. In *IEEE Spoken Language Technology Workshop (SLT)*. [arXiv:2409.09629](#)
- David Smith Sundarsingh, Jun Wang, Jyotirmoy V. Deshmukh, and Yiannis Kantaros. 2025. ConformalNL2LTL: Translating Natural Language Instructions into Temporal Logic Formulas with Conformal Correctness Guarantees. [arXiv:2504.21022](#)
- Zhi Rui Tam, Cheng-Kuang Wu, Yi-Lin Tsai, Chieh-Yen Lin, Hung yi Lee, and Yun-Nung Chen. 2024. Let Me Speak Freely? A Study on the Impact of Format Restrictions on Performance of Large Language Models. In *EMNLP Industry*. [arXiv:2408.02442](#)
- Chee Heng Tan, Zhuoyi Lin, Mehul Motani, and Wee Sun Lee. 2026. On the effectiveness of reward functions in reinforcement learning for confidence calibration of large language models. [arXiv:2607.04332](#)
- Hexiang Tan, Fei Sun, Sha Liu, Du Su, Qi Cao, Xin Chen, Jingang Wang, Xunliang Cai, Yuanzhuo Wang, Huawei Shen, and Xueqi Cheng. 2025. Too Consistent to Detect: A Study of Self-Consistent Errors in LLMs. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. [arXiv:2505.17656](#)
- Xiangru Tang, Yiming Zong, Jason Phang, Yilun Zhao, Wangchunshu Zhou, Arman Cohan, and Mark Gerstein. 2023. Struc-Bench: Are Large Language Models Really Good at Generating Complex Structured Data? [arXiv:2309.08963](#)
- Zhisheng Tang, Ke Shen, and Mayank Kejriwal. 2024. An Evaluation of Estimative Uncertainty in Large Language Models. [arXiv:2405.15185](#)
- Sree Harsha Tanneru, Chirag Agarwal, and Himabindu Lakkaraju. 2023. Quantifying Uncertainty in Natural Language Explanations of Large Language Models. In *Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS)*. [arXiv:2311.03533](#)
- Linwei Tao, Yi-Fan Yeh, Mingjing Dong, Tao Huang, Philip Torr, and Chang Xu. 2025a. Revisiting Uncertainty Estimation and Calibration of Large Language Models. [arXiv:2505.23854](#)
- Linwei Tao, Yi-Fan Yeh, Bo Kai, Mingjing Dong, Tao Huang, Tom A. Lamb, Jialin Yu, Philip H. S. Torr, and Chang Xu. 2025b. Can Large Language Models Express Uncertainty Like Human? [arXiv:2509.24202](#)
- Shuchang Tao, Liuyi Yao, Hanxing Ding, Yuexiang Xie, Qi Cao, Fei Sun, Jinyang Gao, Huawei Shen, and Bolin Ding. 2024. When to Trust LLMs: Aligning Confidence with Response Quality. In *Findings ACL 2024*. <https://doi.org/10.18653/v1/2024.findings-acl.357> [arXiv:2404.17287](#)
- Amir Taubenfeld, Tom Sheffer, Eran Ofek, Amir Feder, Ariel Goldstein, Zorik Gekhman, and Gal Yona. 2025. Confidence Improves Self-Consistency in LLMs. <https://doi.org/10.18653/v1/2025.findings-acl.1030> [arXiv:2502.06233](#) Findings ACL 2025.
- Sina Tayebati, Divake Kumar, Nastaran Darabi, Dinithi Jayasuriya, Ranganath Krishnan, and Amit Ranjan Trivedi. 2025. Learning Conformal Abstention Policies for Adaptive Risk Management in Large Language and Vision-Language Models. [arXiv:2502.06884](#)
- Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D. Manning. 2023. Just Ask for Calibration: Strategies for Eliciting Calibrated Confidence Scores from Language Models Fine-Tuned with Human Feedback. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2023.emnlp-main.330> [arXiv:2305.14975](#)
- Zailong Tian, Zhuoheng Han, Yanzhe Chen, Haozhe Xu, Xi Yang, Richeng Xuan, Houfeng Wang, and Lizi Liao. 2025. Overconfidence in LLM-as-a-Judge: Diagnosis and Confidence-Driven Solution. [arXiv:2508.06225](#)
- Christian Tomani, Kamalika Chaudhuri, Ivan Evtimov, Daniel Cremers, and Mark Ibrahim. 2024. Uncertainty-Based Abstention in LLMs Improves Safety and Reduces Hallucinations. [arXiv:2404.10960](#)
- Francesco Tonolini, Nikolaos Aletras, Jordan Massiah, and Gabriella Kazai. 2024. Bayesian Prompt Ensembles: Model Uncertainty Estimation for Black-Box Large Language Models. In *Findings of the Association for Computational Linguistics: ACL 2024*. Association for Computational Linguistics, 12229–12272. <https://doi.org/10.18653/v1/2024.findings-acl.728>

- Benjamin Turtel, Danny Franklin, Kris Skotheim, Luke Hewitt, and Philipp Schoenegger. 2025. Outcome-based Reinforcement Learning to Predict the Future. [arXiv:2505.17989](#)
- Benjamin Turtel, Paul Wilczewski, Kris Skotheim, Ville A. Satopää, and Philip E. Tetlock. 2026. How Proper Scoring Rules Shape LLM Forecasting. [arXiv:2608.28482](#)
- Shubham Ugare, Tarun Suresh, Hangoo Kang, Sasa Misailovic, and Gagandeep Singh. 2024. SynCode: LLM Generation with Grammar Augmentation. [arXiv:2403.01632](#)
- Dennis Ulmer, Martin Gubri, Hwaran Lee, Sangdoo Yun, and Seong Joon Oh. 2024a. Calibrating Large Language Models Using Their Generations Only. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2024.acl-long.824> [arXiv:2403.05973](#)
- Dennis Ulmer, Alexandra Lorson, Ivan Titov, and Christian Hardmeier. 2025. Anthropomorphic Uncertainty: What Verbalized Uncertainty in Language Models is Missing. [arXiv:2507.10587](#)
- Dennis Ulmer, Chrysoula Zerva, and André F. T. Martins. 2024b. Non-Exchangeable Conformal Language Generation with Nearest Neighbors. In *Findings EACL 2024*. <https://doi.org/10.18653/v1/2024.findings-eacl.129> [arXiv:2402.00707](#)
- Carel van Niekerk, Renato Vukovic, Benjamin Ruppik, Hsien-chin Lin, Shutong Feng, and Milica Gašić. 2025. Post-Training Large Language Models via Reinforcement Learning from Self-Feedback. [arXiv:2507.21931](#)
- Neeraj Varshney and Chitta Baral. 2022. Model Cascading: Towards Jointly Improving Efficiency and Accuracy of NLP Systems. [arXiv:2210.05528](#) EMNLP 2022.
- Roman Vashurin, Ekaterina Fadeeva, Artem Vazhentsev, Lyudmila Rvanova, Akim Tsvigun, Daniil Vasilev, Rui Xing, Abdelrahman Boda Sadallah, Kirill Grishchenkov, Sergey Petrakov, Alexander Panchenko, Timothy Baldwin, Preslav Nakov, Maxim Panov, and Artem Shelmanov. 2025. Benchmarking Uncertainty Quantification Methods for Large Language Models with LM-Polygraph. *Transactions of the Association for Computational Linguistics* (2025). [arXiv:2406.15627](#)
- Artem Vazhentsev, Lyudmila Rvanova, Gleb Kuzmin, Ekaterina Fadeeva, Ivan Lazichny, Alexander Panchenko, Maxim Panov, Mrinmaya Sachan, Preslav Nakov, Timothy Baldwin, and Artem Shelmanov. 2026. Efficient Hallucination Detection for LLMs Using Uncertainty-Aware Attention Heads. In *International Conference on Machine Learning (ICML)*. [arXiv:2505.20045](#)
- Mauricio Villavicencio, Sitong Pan, and Qianwen Wang. 2026. Not All Uncertainty Is Equal: How Uncertainty Granularity Shapes Human Verification in LLM-Assisted Decision Making. <https://doi.org/10.1145/3805689.3812329> [arXiv:2605.28571](#) FAccT 2026.
- Harit Vishwakarma, Alan Mishler, Thomas Cook, Niccolò Dalmaso, Natraj Raman, and Sumitra Ganesh. 2024. Prune 'n Predict: Optimizing LLM Decision-making with Conformal Prediction. In *ICML 2025*. [arXiv:2501.00555](#)
- Ante Wang, Weizhi Ma, and Yang Liu. 2025d. Let the Model Distribute Its Doubt: Confidence Estimation through Verbalized Probability Distribution. [arXiv:2511.14275](#)
- Cheng Wang. 2023. Calibration in Deep Learning: A Survey of the State-of-the-Art. [arXiv:2308.01222](#)
- Cheng Wang, Gyuri Szarvas, Georges Balazs, Pavel Danchenko, and Patrick Ernst. 2024f. Calibrating Verbalized Probabilities for Large Language Models. [arXiv:2410.06707](#)
- Darren Yow-Bang Wang, Zhengyuan Shen, Soumya Smruti Mishra, Zhichao Xu, Yifei Teng, and Haibo Ding. 2025e. SLOT: Structuring the Output of Large Language Models. [arXiv:2505.04016](#)
- Jun Wang, Guocheng He, and Yiannis Kantaros. 2025c. Probabilistically Correct Language-based Multi-Robot Planning using Conformal Prediction. *IEEE Robotics and Automation Letters* (2025). <https://doi.org/10.1109/LRA.2024.3504233> [arXiv:2402.15368](#)
- Jun Wang, Jiaming Tong, Kaiyuan Tan, Yevgeniy Vorobeychik, and Yiannis Kantaros. 2025f. Conformal Temporal Logic Planning using Large Language Models. *ACM Transactions on Cyber-Physical Systems* (2025). <https://doi.org/10.1145/3769111> [arXiv:2309.10092](#)
- Jiawei Wang, Yanfei Zhou, Siddhartha Devic, and Deqing Fu. 2026c. Are LLM Decisions Faithful to Verbal Confidence? [arXiv:2601.07767](#)
- Liaoyaqi Wang, Chunsheng Zuo, William Jurayj, Benjamin Van Durme, and Anqi Liu. 2026d. Process Supervision of Confidence Margin for Calibrated LLM Reasoning. [arXiv:2604.23333](#)
- Peiyi Wang, Lei Li, Liang Chen, Zefan Cai, Dawei Zhu, Binghuai Lin, Yunbo Cao, Qi Liu, Tianyu Liu, and Zhifang Sui. 2023a. Large Language Models are not Fair Evaluators. [arXiv:2305.17926](#)
- Qingni Wang, Yue Fan, and Xin Eric Wang. 2025b. SAFER: Risk-Constrained Sample-then-Filter in Large Language Models. [arXiv:2510.10193](#)
- Qingni Wang, Tiantian Geng, Zhiyuan Wang, Teng Wang, Bo Fu, and Feng Zheng. 2024b. Sample then Identify: A General Framework for Risk Control and Assessment in Multimodal Large Language Models. In *ICLR 2025*. [arXiv:2410.08174](#)
- Shuo Wang, Zhaopeng Tu, Shuming Shi, and Yang Liu. 2020. On the Inference Calibration of Neural Machine Translation. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2020.acl-main.278> [arXiv:2005.00963](#)

- Victor Wang and Elias Stengel-Eskin. 2025. Calibrating Verbalized Confidence with Self-Generated Distractors. In *International Conference on Learning Representations (ICLR)*. arXiv:2509.25532
- Xinpeng Wang, Chengzhi Hu, Bolei Ma, Paul Röttger, and Barbara Plank. 2024c. Look at the Text: Instruction-Tuned Language Models are More Robust Multiple Choice Selectors than You Think. arXiv:2404.08382 COLM 2024.
- Xinpeng Wang, Bolei Ma, Chengzhi Hu, Leon Weber-Genzel, Paul Röttger, Frauke Kreuter, Dirk Hovy, and Barbara Plank. 2024d. "My Answer is C": First-Token Probabilities Do Not Match Text Answers in Instruction-Tuned Language Models. arXiv:2402.14499 Findings ACL 2024.
- Xi Wang, Anushri Suresh, Alvin Zhang, Rishi More, William Jurayj, Benjamin Van Durme, Mehrdad Farajtabar, Daniel Khashabi, and Eric Nalisnick. 2026b. Conformal Thinking: Risk Control for Reasoning on a Compute Budget. In *ICML 2026*. arXiv:2602.03814
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023b. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *International Conference on Learning Representations (ICLR)*. arXiv:2203.11171
- Yibin Wang, Haizhou Shi, Ligong Han, Dimitris Metaxas, and Hao Wang. 2024e. BLoB: Bayesian Low-Rank Adaptation by Backpropagation for Large Language Models. In *NeurIPS 2024*. arXiv:2406.11675
- Zhiyuan Wang, Jinhao Duan, Lu Cheng, Yue Zhang, Qingni Wang, Xiaoshuang Shi, Kaidi Xu, Hengtao Shen, and Xiaofeng Zhu. 2024a. ConU: Conformal Uncertainty in Large Language Models with Correctness Coverage Guarantees. In *Findings EMNLP 2024*. <https://doi.org/10.18653/v1/2024.findings-emnlp.404> arXiv:2407.00499
- Zhiyuan Wang, Jinhao Duan, Chenxi Yuan, Qingyu Chen, Tianlong Chen, Yue Zhang, Ren Wang, Xiaoshuang Shi, and Kaidi Xu. 2025a. Word-Sequence Entropy: Towards Uncertainty Estimation in Free-Form Medical Question Answering Applications and Beyond. *Engineering Applications of Artificial Intelligence* (2025). arXiv:2402.14259
- Ziqi Wang, Xingzhou Lou, Meiqi Wu, Zhenqiang Wen, and Junge Zhang. 2026a. Calibration-Aware Policy Optimization for Reasoning LLMs. In *ACL 2026*. <https://doi.org/10.18653/v1/2026.acl-long.836> arXiv:2604.12632
- Zhiyuan Wang, Qingni Wang, Yue Zhang, Tianlong Chen, Xiaofeng Zhu, Xiaoshuang Shi, and Kaidi Xu. 2025g. SConU: Selective Conformal Uncertainty in Large Language Models. In *ACL 2025*. <https://doi.org/10.18653/v1/2025.acl-long.934> arXiv:2504.14154
- Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese, John Schulman, and William Fedus. 2024. Measuring short-form factuality in large language models. arXiv:2411.04368
- Zhepei Wei, Xiao Yang, Kai Sun, Jiaqi Wang, Rulin Shao, Jingxiang Chen, Mohammad Kachuee, Teja Gollapudi, Yiwei Liao, Nicolas Scheffer, Rakesh Wang, Anuj Kumar, Yu Meng, Wen-tau Yih, and Xin Luna Dong. 2025. TruthRL: Incentivizing Truthful LLMs via Reinforcement Learning. In *ICML 2026*. arXiv:2509.25760
- Bingbing Wen, Jihan Yao, Shangbin Feng, Chenjun Xu, Yulia Tsvetkov, Bill Howe, and Lucy Lu Wang. 2025. Know Your Limits: A Survey of Abstention in Large Language Models. arXiv:2407.18418 TACL.
- Brandon T. Willard and Rémi Louf. 2023. Efficient Guided Generation for Large Language Models. arXiv:2307.09702
- Jiayun Wu, Jiashuo Li, Zhiyuan Zeng, Tianyang Zhan, Tianle Cai, and Wenhao Huang. 2025a. Mitigating LLM Hallucination via Behaviorally Calibrated Reinforcement Learning. arXiv:2512.19920
- Yu Wu, Shuo Wu, Ye Tao, Yansong Li, and Anand D. Sarwate. 2025b. From Deferral to Learning: Online In-Context Knowledge Distillation for LLM Cascades. arXiv:2509.22984
- Boming Xia, Liming Zhu, Erdun Gao, Qinghua Lu, Minhui Xue, and Dino Sejdinovic. 2026. Uncertainty Propagation in LLM-Based Systems. arXiv:2604.23505
- Congying Xia, Chen Xing, Jiangshu Du, Xinyi Yang, Yihao Feng, Ran Xu, Wenpeng Yin, and Caiming Xiong. 2024a. FOFO: A Benchmark to Evaluate LLMs' Format-Following Capability. arXiv:2402.18667
- Heming Xia, Zhe Yang, Qingxiu Dong, Peiyi Wang, Yongqi Li, Tao Ge, Tianyu Liu, Wenjie Li, and Zhifang Sui. 2024b. Unlocking Efficiency in Large Language Model Inference: A Comprehensive Survey of Speculative Decoding. arXiv:2401.07851 Findings of ACL 2024.
- Zhiqiu Xia, Jinxuan Xu, Yuqian Zhang, and Hang Liu. 2025. A Survey of Uncertainty Estimation Methods on Large Language Models. <https://doi.org/10.18653/v1/2025.findings-acl.1101> Findings of ACL 2025.
- Jiancong Xiao, Bojian Hou, Zhanliang Wang, Ruochen Jin, Qi Long, Weijie J. Su, and Li Shen. 2025. Restoring Calibration for Aligned Large Language Models: A Calibration-Aware Fine-Tuning Approach. In *ICML 2025*. arXiv:2505.01997
- Wenyi Xiao, Xinchu Xu, and Leilei Gan. 2026. VL-Calibration: Decoupled Confidence Calibration for Large Vision-Language Models Reasoning. In *ACL 2026*. <https://doi.org/10.18653/v1/2026.acl-long.2074> arXiv:2604.09529
- Yuxin Xiao, Paul Pu Liang, Umang Bhatt, Willie Neiswanger, Ruslan Salakhutdinov, and Louis-Philippe Morency. 2022. Uncertainty Quantification with Pre-trained Language Models: A Large-Scale Empirical Analysis. In *Findings of the Association for Computational Linguistics: EMNLP*. <https://doi.org/10.18653/v1/2022.findings-emnlp.538> arXiv:2210.04714
- Yijun Xiao and William Yang Wang. 2021. On Hallucination and Predictive Uncertainty in Conditional Language Generation. In *Conference of the European Chapter of the Association for Computational Linguistics (EACL)*. arXiv:2103.15025

- Yisheng Xiao, Lijun Wu, Junliang Guo, Juntao Li, Min Zhang, Tao Qin, and Tie yan Liu. 2023. A Survey on Non-Autoregressive Generation for Neural Machine Translation and Beyond. arXiv:2204.09269 IEEE TPAMI.
- Johnathan Xie, Annie S. Chen, Yoonho Lee, Eric Mitchell, and Chelsea Finn. 2024a. Calibrating Language Models with Adaptive Temperature Scaling. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2024.emnlp-main.1007> arXiv:2409.19817
- Liangru Xie, Hui Liu, Jingying Zeng, Xianfeng Tang, Yan Han, Chen Luo, Jing Huang, Zhen Li, Suhang Wang, and Qi He. 2024b. A Survey of Calibration Process for Black-Box LLMs. arXiv:2412.12767
- Xiaohu Xie, Xiaohu Liu, and Benjamin Yao. 2026. Know When You're Wrong: Aligning Confidence with Correctness for LLM Error Detection. arXiv:2603.06604
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2023. Can LLMs Express Their Uncertainty? An Empirical Evaluation of Confidence Elicitation in LLMs. In *International Conference on Learning Representations (ICLR)*. arXiv:2306.13063
- Hongshen Xu, Zichen Zhu, Situo Zhang, Da Ma, Shuai Fan, Lu Chen, and Kai Yu. 2024b. Rejection Improves Reliability: Training LLMs to Refuse Unknown Questions Using RL from Knowledge Feedback. arXiv:2403.18349
- Tianyang Xu, Shujin Wu, Shizhe Diao, Xiaozhe Liu, Xingyao Wang, Yangyi Chen, and Jing Gao. 2024a. SaySelf: Teaching LLMs to Express Confidence with Self-Reflective Rationales. In *EMNLP 2024*. <https://doi.org/10.18653/v1/2024.emnlp-main.343> arXiv:2405.20974
- Xiaogang Xu, Jiaqi Tang, Jianmin Chen, Yingying Yan, Zhenchao Tang, Xiangxin Zhou, Xiaobin Hu, Wei Wei, Jinfeng Wu, Qifeng Chen, Lu Zhou, Jiafei Wu, Zhe Liu, Jianwei Yin, and Weimin Zheng. 2026b. LLM-based Agents for Forecasting and Prediction: Methods, Training, Evaluation, and Applications. arXiv:2608.23058
- Yinuo Xu, Yuwei Liang, Jianjie Cheng, Meng Wang, Yongcan Yu, Shuo Lu, and Jian Liang. 2026a. DualStake: Dual-Path Confidence Calibration in Deep Research Agents. In *EMNLP 2026*. arXiv:2609.00935
- Weihao Xuan, Qingcheng Zeng, Heli Qi, Yunze Xiao, Junjue Wang, and Naoto Yokoya. 2026. The Confidence Dichotomy: Analyzing and Mitigating Miscalibration in Tool-Use Agents. In *ACL 2026*. <https://doi.org/10.18653/v1/2026.acl-long.520> arXiv:2601.07264
- Boyang Xue, Fei Mi, Qi Zhu, Hongru Wang, Rui Wang, Sheng Wang, Erxin Yu, Xuming Hu, and Kam-Fai Wong. 2025. UAlign: Leveraging Uncertainty Estimations for Factuality Alignment on Large Language Models. In *ACL 2025*. <https://doi.org/10.18653/v1/2025.acl-long.299> arXiv:2412.11803
- Boyang Xue, Hongru Wang, Rui Wang, Sheng Wang, Zezhong Wang, Yiming Du, Bin Liang, Wenxuan Zhang, and Kam-Fai Wong. 2024. MlingConf: A Comprehensive Study of Multilingual Confidence Estimation on Large Language Models. In *Findings of the Association for Computational Linguistics: ACL*. arXiv:2402.13606
- Yasin Abbasi Yadkori, Ilja Kuzborskij, András György, and Csaba Szepesvári. 2024a. To Believe or Not to Believe Your LLM. arXiv:2406.02543
- Yasin Abbasi Yadkori, Ilja Kuzborskij, David Stutz, András György, Adam Fisch, Arnaud Doucet, Iuliya Beloshapka, Wei-Hung Weng, Yao-Yuan Yang, Csaba Szepesvári, Ali Taylan Cemgil, and Nenad Tomasev. 2024b. Mitigating LLM Hallucinations via Conformal Abstention. arXiv:2405.01563
- Duygu Nur Yaldiz, Evangelia Spiliopoulou, Zheng Qi, Siddharth Varia, Srikanth Doss, and Nikolaos Pappas. 2026. Balancing Classification and Calibration Performance in Decision-Making LLMs via Calibration Aware Reinforcement Learning. In *Findings ACL 2026*. <https://doi.org/10.18653/v1/2026.findings-acl.610> arXiv:2601.13284
- Adam Yang, Chen Chen, and Konstantinos Pitas. 2024a. Just rephrase it! Uncertainty estimation in closed-source language models via multiple rephrased queries. In *NeurIPS 2024 Workshop*. arXiv:2405.13907
- Chenxu Yang, Qingyi Si, Yongjie Duan, Zheliang Zhu, Chenyu Zhu, Qiaowei Li, Minghui Chen, Zheng Lin, and Weiping Wang. 2025b. Dynamic Early Exit in Reasoning Models. arXiv:2504.15895
- Daniel Yang, Yao-Hung Hubert Tsai, and Makoto Yamada. 2024d. On Verbalized Confidence Scores for LLMs. arXiv:2412.14737
- Guang Yang and Xinyang Liu. 2025. Conformal Sets in Multiple-Choice Question Answering under Black-Box Settings with Provable Coverage Guarantees. arXiv:2508.05544
- Jisoo Yang, Jaeho Han, Trung X. Pham, and Junyeong Kim. 2026a. The Mirage of Calibrated Confidence: Trajectory-Independence of Verbalized Confidence in Vision-Language Models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2609.18453
- Jialin Yang, Dongfu Jiang, Lipeng He, Sherman Siu, Yuxuan Zhang, Disen Liao, Zhuofeng Li, Huaye Zeng, Yiming Jia, Haozhe Wang, Benjamin Schneider, Chi Ruan, Wentao Ma, Zhiheng Lyu, Yifei Wang, Yi Lu, Quy Duc Do, Ziyang Jiang, Ping Nie, and Wenhu Chen. 2025a. StructEval: Benchmarking LLMs' Capabilities to Generate Structural Outputs. arXiv:2505.20139
- Ruixin Yang, Dheeraj Rajagopal, Shirley Anugrah Hayati, Bin Hu, and Dongyeop Kang. 2024c. Confidence Calibration and Rationalization for LLMs via Multi-Agent Deliberation. In *ICLR Workshop on Reliable and Responsible Foundation Models*. arXiv:2404.09127
- Xuqing Yang, Yi Yuan, Shanzhe Lei, and Xuhong Wang. 2026b. Scaling with Confidence: Calibrating Confidence of LLMs for Adaptive Test Time Scaling. arXiv:2607.01612

- Yuqing Yang, Ethan Chern, Xipeng Qiu, Graham Neubig, and Pengfei Liu. 2024b. Alignment for Honesty. In *NeurIPS 2024*. <https://doi.org/10.52202/079017-2030> arXiv:2312.07000
- Fanghua Ye, Mingming Yang, Jianhui Pang, Longyue Wang, Derek F. Wong, Emine Yilmaz, Shuming Shi, and Zhaopeng Tu. 2024. Benchmarking LLMs via Uncertainty Quantification. In *NeurIPS 2024 D&B*. arXiv:2401.12794
- Xi Ye and Greg Durrett. 2022. Can Explanations Be Useful for Calibrating Black Box Models?. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*. <https://doi.org/10.18653/v1/2022.acl-long.429> arXiv:2110.07586
- Yuanchang Ye. 2025. Conformal P-Value in Multiple-Choice Question Answering Tasks with Provable Risk Control. arXiv:2508.10022
- Zhangyue Yin, Qiushi Sun, Qipeng Guo, Jiawen Wu, Xipeng Qiu, and Xuanjing Huang. 2023. Do Large Language Models Know What They Don't Know?. In *Findings of the Association for Computational Linguistics: ACL*. arXiv:2305.18153
- Zhenyun Yin, Shujie Wang, Xuhong Wang, Xingjun Ma, and Yinchun Wang. 2026. Deliberative Searcher: Improving LLM Reliability via Reinforcement Learning with constraints. In *ACL 2026*. <https://doi.org/10.18653/v1/2026.acl-long.199> arXiv:2507.16727
- Gal Yona, Roei Aharoni, and Mor Geva. 2024. Can Large Language Models Faithfully Express Their Intrinsic Uncertainty in Words?. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2405.16908
- Dongkeun Yoon, Seungone Kim, Sohee Yang, Sunkyoung Kim, Soyeon Kim, Yongil Kim, Eunbi Choi, Yireun Kim, and Minjoon Seo. 2025. Reasoning Models Better Express Their Confidence. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2505.14489
- Fengfei Yu, Ruijia Niu, Dongxia Wu, Yian Ma, and Rose Yu. 2026b. Calibrating LLMs with Semantic-level Reward. arXiv:2605.15588
- Runpeng Yu, Qi Li, and Xinchao Wang. 2025a. Discrete Diffusion in Large Language and Multimodal Models: A Survey. arXiv:2506.13759
- Tianjiao Yu, Vedant Shah, Muntasir Wahed, Kiet A. Nguyen, Adheesh Juvekar, Tal August, and Ismini Lourentzou. 2025b. Uncertainty in Action: Confidence Elicitation in Embodied Agents. arXiv:2503.10628
- Zhiyin Yu, Yuchen Mou, Juncheng Yan, Junyu Luo, Chunchun Chen, Xing Wei, Yunhui Liu, Hongru Sun, Yuxing Zhang, Jun Xu, Yatao Bian, Ming Zhang, Wei Ye, Tieke He, Jie Yang, Guanjie Zheng, Zhonghai Wu, Bo Zhang, Lei Bai, and Xiao Luo. 2026a. A Survey of Reinforcement Learning for Large Language Models under Data Scarcity: Challenges and Solutions. arXiv:2604.17312
- Murong Yue, Jie Zhao, Min Zhang, Liang Du, and Ziyu Yao. 2024. Large Language Model Cascades with Mixture of Thoughts Representations for Cost-efficient Reasoning. arXiv:2310.03094 ICLR 2024.
- Michael J. Zellinger and Matt Thomson. 2025. Rational Tuning of LLM Cascades via Probabilistic Modeling. arXiv:2501.09345
- Chrysoula Zerva and André F. T. Martins. 2024. Conformalizing Machine Translation Evaluation. *Transactions of the Association for Computational Linguistics* (2024). https://doi.org/10.1162/tacl_a_00711 arXiv:2306.06221
- Skylar Zhai, Jingcheng Liang, and Dongyeop Kang. 2026. Abstain-R1: Calibrated Abstention and Post-Refusal Clarification via Verifiable RL. In *Findings ACL 2026*. <https://doi.org/10.18653/v1/2026.findings-acl.985> arXiv:2604.17073
- Botong Zhang, Shuo Li, and Osbert Bastani. 2024b. Conformal Structured Prediction. In *ICLR 2025*. arXiv:2410.06296
- Caiqi Zhang, Fangyu Liu, Marco Basaldella, and Nigel Collier. 2024c. LUQ: Long-text Uncertainty Quantification for LLMs. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2403.20279
- Caiqi Zhang, Ruihan Yang, Zhisong Zhang, Xinting Huang, Sen Yang, Dong Yu, and Nigel Collier. 2025d. Atomic Calibration of LLMs in Long-Form Generations. In *Asian Chapter of the Association for Computational Linguistics (AACL-IJCNLP)*. arXiv:2410.13246
- Caiqi Zhang, Xiaochen Zhu, Chengzu Li, Nigel Collier, and Andreas Vlachos. 2025e. LoVeC: Reinforcement Learning for Better Verbalized Confidence in Long-Form Generations. In *ACL 2026*. arXiv:2505.23912
- Glenn Zhang, Treasure Mayowa, Jason Fan, Yicheng Fu, Aaron Sandoval, Sean O'Brien, and Kevin Zhu. 2025c. Direct Confidence Alignment: Aligning Verbalized Confidence with Internal Confidence In Large Language Models. In *ACL-SRW 2025*. arXiv:2512.11998
- Hanning Zhang, Shizhe Diao, Yong Lin, Yi R. Fung, Qing Lian, Xingyao Wang, Yangyi Chen, Heng Ji, and Tong Zhang. 2023a. R-Tuning: Instructing Large Language Models to Say 'I Don't Know'. In *NAACL 2024*. arXiv:2311.09677
- Hanlin Zhang, Yi-Fan Zhang, Yaodong Yu, Dhruv Madeka, Dean Foster, Eric Xing, Himabindu Lakkaraju, and Sham Kakade. 2023d. A Study on the Calibration of In-context Learning. In *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*. arXiv:2312.04021
- Jiaxin Zhang, Wendi Cui, Zhuohang Li, Lifu Huang, Bradley Malin, Caiming Xiong, and Chien-Sheng Wu. 2026a. From Passive Metric to Active Signal: The Evolving Role of Uncertainty Quantification in Large Language Models. arXiv:2601.15690
- Jiaxin Zhang, Zhuohang Li, Kamalika Das, Bradley A. Malin, and Sricharan Kumar. 2023b. SAC3: Reliable Hallucination Detection in Black-Box Language Models via Semantic-aware Cross-check Consistency. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. arXiv:2311.01740

- Kaiyan Zhang, Yuxin Zuo, Bingxiang He, Youbang Sun, Runze Liu, Che Jiang, Yuchen Fan, Kai Tian, Guoli Jia, Pengfei Li, Yu Fu, Xingtai Lv, Yuchen Zhang, Sihang Zeng, Shang Qu, Haozhan Li, Shijie Wang, Yuru Wang, Xinwei Long, Fangfu Liu, Xiang Xu, Jiaze Ma, Xuekai Zhu, Ermo Hua, Yihao Liu, Zonglin Li, Huayu Chen, Xiaoye Qu, Yafu Li, Weize Chen, Zhenzhao Yuan, Junqi Gao, Dong Li, Zhiyuan Ma, Ganqu Cui, Zhiyuan Liu, Biqing Qi, Ning Ding, and Bowen Zhou. 2025f. A Survey of Reinforcement Learning for Large Reasoning Models. [arXiv:2509.08827](#)
- Mozhi Zhang, Mianqiu Huang, Rundong Shi, Linsen Guo, Chong Peng, Peng Yan, Yaqian Zhou, and Xipeng Qiu. 2024a. Calibrating the Confidence of Large Language Models by Eliciting Fidelity. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. [arXiv:2404.02655](#)
- Min-Ling Zhang and Deng-Bao Wang. 2026. Uncertainty Calibration in Deep Learning: Methods, Emerging Challenges, and LLM Frontiers. *Journal of Computer Science and Technology* 41, 1 (2026), 318–340. <https://doi.org/10.1007/s11390-026-6426-z>
- Ruiyang Zhang, Hu Zhang, and Zhedong Zheng. 2024e. VL-Uncertainty: Detecting Hallucination in Large Vision-Language Model via Uncertainty Estimation. [arXiv:2411.11919](#)
- Shujian Zhang, Chengyue Gong, and Eunsol Choi. 2021. Knowing More About Questions Can Help: Improving Calibration in Question Answering. In *Findings of the Association for Computational Linguistics: ACL*. <https://doi.org/10.18653/v1/2021.findings-acl.172> [arXiv:2106.01494](#)
- Tianhang Zhang, Lin Qiu, Qipeng Guo, Cheng Deng, Yue Zhang, Zheng Zhang, Chenghu Zhou, Xinbing Wang, and Luoyi Fu. 2023c. Enhancing Uncertainty-Based Hallucination Detection with Stronger Focus. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. [arXiv:2311.13230](#)
- Xiaoying Zhang, Baolin Peng, Ye Tian, Jingyan Zhou, Lifeng Jin, Linfeng Song, Haitao Mi, and Helen Meng. 2024d. Self-Alignment for Factuality: Mitigating Hallucinations in LLMs via Self-Evaluation. In *ACL 2024*. <https://doi.org/10.18653/v1/2024.acl-long.107> [arXiv:2402.09267](#)
- Xiaopan Zhang, Zejin Wang, Zhixu Li, Jianpeng Yao, and Jiachen Li. 2026b. CommCP: Efficient Multi-Agent Coordination via LLM-Based Communication with Conformal Prediction. In *ICRA 2026*. [arXiv:2602.06038](#)
- Yue Zhang, Yafu Li, Leyang Cui, Deng Cai, Lema Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Chen Xu, Yulong Chen, Longyue Wang, Anh Tuan Luu, Wei Bi, Freda Shi, and Shuming Shi. 2025a. Siren’s Song in the AI Ocean: A Survey on Hallucination in Large Language Models. [arXiv:2309.01219](#)
- Zhaohan Zhang, Ziquan Liu, and Ioannis Patras. 2025b. GrACE: A Generative Approach to Better Confidence Elicitation and Efficient Test-Time Scaling in Large Language Models. [arXiv:2509.09438](#)
- Tony Z. Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. 2021. Calibrate Before Use: Improving Few-Shot Performance of Language Models. In *Proceedings of the International Conference on Machine Learning (ICML)*. [arXiv:2102.09690](#)
- Xingtao Zhao, Hao Peng, Dingli Su, Xianghua Zeng, Chunyang Liu, Jinzhi Liao, and Philip S. Yu. 2026. SeSE: Black-Box Uncertainty Quantification for Large Language Models Based on Structural Information Theory. In *Conference on Uncertainty in Artificial Intelligence (UAI)*. [arXiv:2511.16275](#)
- Xinran Zhao, Hongming Zhang, Xiaoman Pan, Wenlin Yao, Dong Yu, Tongshuang Wu, and Jianshu Chen. 2024b. Fact-and-Reflection (FaR) Improves Confidence Calibration of Large Language Models. In *Findings of the Association for Computational Linguistics: ACL*. [arXiv:2402.17124](#)
- Yukun Zhao, Lingyong Yan, Weiwei Sun, Guoliang Xing, Chong Meng, Shuaiqiang Wang, Zhicong Cheng, Zhaochun Ren, and Dawei Yin. 2024a. Knowing What LLMs DO NOT Know: A Simple Yet Effective Self-Detection Method. In *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*. [arXiv:2310.17918](#)
- Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. 2023. Large Language Models Are Not Robust Multiple Choice Selectors. In *International Conference on Learning Representations (ICLR)*. [arXiv:2309.03882](#)
- Han Zhou, Xingchen Wan, Lev Proleev, Diana Mincu, Jilin Chen, Katherine Heller, and Subhrajit Roy. 2023b. Batch Calibration: Rethinking Calibration for In-Context Learning and Prompt Engineering. In *International Conference on Learning Representations (ICLR)*. [arXiv:2309.17249](#)
- Kaitlyn Zhou, Jena D. Hwang, Xiang Ren, Nouha Dziri, Dan Jurafsky, and Maarten Sap. 2024b. Rel-A.I.: An Interaction-Centered Approach To Measuring Human-LM Reliance. [arXiv:2407.07950](#)
- Kaitlyn Zhou, Jena D. Hwang, Xiang Ren, and Maarten Sap. 2024a. Relying on the Unreliable: The Impact of Language Models’ Reluctance to Express Uncertainty. [arXiv:2401.06730](#) *ACL* 2024.
- Kaitlyn Zhou, Dan Jurafsky, and Tatsunori Hashimoto. 2023a. Navigating the Grey Area: How Expressions of Uncertainty and Overconfidence Affect Language Models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. [arXiv:2302.13439](#)
- Ziang Zhou, Tianyuan Jin, Jieming Shi, and Qing Li. 2025. SteerConf: Steering LLMs for Confidence Elicitation. [arXiv:2503.02863](#)
- Zixuan Zhou, Xuefei Ning, Ke Hong, Tianyu Fu, Jiaming Xu, Shiyao Li, Yuming Lou, Luning Wang, Zhihang Yuan, Xiuhong Li, Shengen Yan, Guohao Dai, Xiao-Ping Zhang, Yuhan Dong, and Yu Wang. 2024c. A Survey on Efficient Inference for Large Language Models. [arXiv:2404.14294](#)

- Chiwei Zhu, Benfeng Xu, Quan Wang, Yongdong Zhang, and Zhendong Mao. 2023. On the Calibration of Large Language Models and Alignment. In *Findings of the Association for Computational Linguistics: EMNLP*. <https://doi.org/10.18653/v1/2023.findings-emnlp.654> arXiv:2311.13240
- Yinglun Zhu. 2026. Closing the Reflection Gap: A Free Calibration Bonus for Agentic RL. arXiv:2606.14211
- Thomas P. Zollo, Todd Morrill, Zhun Deng, Jake C. Snell, Toniann Pitassi, and Richard Zemel. 2023. Prompt Risk Control: A Rigorous Framework for Responsible Deployment of Large Language Models. In *ICLR 2024*. arXiv:2311.13628
- Haotian Zong, Binze Li, Yufei Long, Sinyin Chang, Jialong Wu, and Gillian K. Hadfield. 2026. I-CALM: Incentivizing Confidence-Aware Abstention for LLM Selective Answering. arXiv:2604.03904
- Chelsea Zou, Yiheng Yao, and Basant Khalil. 2025. Thinking, Faithful and Stable: Mitigating Hallucinations in LLMs. arXiv:2511.15921

A OPERATIONAL DEFINITIONS

Table A1 gives the rule by which each work is placed: for every stage and property, what counts and the nearest case that does not.

Table A1. **Operational definitions used to place each work.** The stages (S0 to S4, plus the consumers X; theory and analysis works are listed but sit outside the core) are defined by *where the calibration signal enters the decision pipeline*, so a work is placed by what its main contribution changes, not by the task it is evaluated on. The first row defines a term, not an inclusion rule: works that report only discrimination are included and marked by their metric. The last three rows define properties recorded for every work: the Contract, Action and Guar. columns of Tables C1 and F1, where Action records what the rule or trained behaviour does when the confidence is low.

Construct	Counts if and only if...	Does not count
Calibrated confidence	the confidence is scored against empirical accuracy with a calibration metric (ECE, Brier, NLL or a reliability diagram)	a confidence that is only reported, or only used to rank answers (AUROC measures discrimination, not calibration)
S0 read-off	the confidence is the model's own token or sequence probability, unchanged	any map applied on top, including a probe on the hidden states (S1)
S1 post-hoc map	a map fitted on held-out data [Guo et al. 2017], on top of the read-off, as a probe on the hidden states of the frozen model, or as an auxiliary model reading the input and output text, or a fixed map (such as pooling the probability of synonymous labels or reweighting token probabilities), changes the confidence but not the model's parameters	re-training or fine-tuning the model (S4)
S2 prompt elicitation	the model is asked, in its input, to state a confidence in words or numbers [Tian et al. 2023]; a method whose main contribution is the prompt stays here even if it also aggregates several prompted answers	a confidence computed outside the model's own text (S0, S3)
S3 extra inference	the confidence needs extra inference, repeated sampling or perturbation (including internal-state signals compared across samples), or carries a distribution-free calibration: a prediction set or a threshold calibrated by conformal or risk control [Campos et al. 2024; Farquhar et al. 2024]	a single forward pass, read off (S0) or probed (S1)
S4 training	a calibration term appears in the model's own loss or reward: a confidence target or calibration loss, a calibrated reward model, a calibration reward, an abstention target, or a step- or meaning-level confidence reward (families T1 to T5) [Bani-Harouni et al. 2026b; Leng et al. 2025]	calibration that changes only as a side effect of training for something else; theory and analyses of such training, which are listed separately
X consumers and contract	the work uses a confidence to act (abstain, hand off, retrieve, filter, adapt compute, set), fixes the output contract, calibrates a model used as a judge against human labels, studies how people rely on a model's output or its stated confidence, or evaluates or benchmarks confidence-based abstention, without changing how the confidence is produced	a work whose main contribution is the confidence itself (S0 to S4)
Theory and analysis	the work analyses when or why training changes calibration, without proposing a training objective	a training method, even one motivated by such analysis (S4)

Table A1, continued

Construct	Counts if and only if...	Does not count
Typed decision output	the options are declared before inference and the model assigns a probability to each option, so the decision and its confidence come from the same distribution	a single output generated under a grammar or schema (constrained output); free text parsed after the fact
Action policy	an explicit rule, or a behaviour the model is trained into (such as refusing when unsure), maps the confidence to an action: answer when it is high; when it is low, abstain, hand the case off to another model or a human (which the papers call routing, cascading, deferral or escalation), trigger retrieval, filter the output (remove, rewrite or revise low-confidence claims, or keep the most confident of several sampled answers), adapt compute (keep sampling, reasoning, querying or exploring, or spend a larger budget; the same rule stops early when the confidence is high), or return a prediction set	a confidence reported with the answer or displayed to a person, with no rule or trained behaviour attached, or a selective-prediction curve reported only as an evaluation (recorded as none; the Metric column records the metrics of the main text)
Coverage guarantee	the method states a finite-sample, distribution-free bound, such as conformal coverage of at least $1 - \alpha$ [Campos et al. 2024]	good empirical calibration on a benchmark, without a stated bound

B PRIOR SURVEYS

Table B1 compares the 39 closest prior surveys on seven topics (Section 1.1).

Table B1. **Coverage of the 39 closest prior surveys:** every verified survey found by the search or by citation searching whose topic is calibration and uncertainty, reinforcement learning for language models, closed-set decisions, acting on confidence, hallucination or conformal prediction. CAL confidence estimation and calibration; RLC calibration-aware training of any kind; RLR reinforcement learning whose reward is a calibration or proper-scoring-rule objective; DEC decisions over a predefined option set; TYP typed, per-option decision outputs; ESC acting on confidence (abstention, deferral, routing); HAL hallucination. ✓ covered, ~ partial, – absent, ? not determinable. The RLR marks were checked in the full text or full reference list of every survey; the other marks come from abstracts and contents, corrected from a full-text reading for the surveys marked †. A 0 marks a column that no prior survey fully covers (TYP). Each full RLR mark is one subsection scoped to a single slice: 2024 methods [Li et al. 2025e], LLM agents [Lin et al. 2026] and forecasting [Xu et al. 2026b]. The last row marks the dimensions within this survey’s scope (●); hallucination is in scope only through abstention training and the theory works (○).

Survey	Venue	CAL	RLC	RLR	DEC	TYP	ESC	HAL
<i>Calibration and uncertainty quantification</i>								
Geng et al. [2024a]†	NAACL	✓	~	–	✓	–	~	~
Liu et al. [2025a]†	KDD	✓	~	–	–	–	~	~
Shorinwa et al. [2025]†	CSUR	✓	~	~	–	–	~	✓
Xia et al. [2025]	ACL Findings	✓	–	–	–	–	–	~
Kang et al. [2025a]†	arXiv	✓	–	–	–	–	~	✓
Zhang et al. [2026a]†	arXiv	✓	~	–	–	–	~	~
Wang [2023]†	arXiv	✓	–	–	✓	–	~	–
Li et al. [2025e]†	TMLR	✓	✓	✓	?	–	~	~
Lin et al. [2026]†	arXiv	✓	~	✓	?	~	✓	?
Li et al. [2025g]	ACL	✓	~	~	?	–	~	~
Xie et al. [2024b]	arXiv	✓	–	~	?	–	?	–
Abbasli et al. [2025]	IEEE TNNLS	✓	?	~	?	–	?	~
Huang et al. [2024c]	arXiv	✓	?	–	?	–	?	?
Beigi et al. [2024b]	arXiv	✓	?	–	?	–	?	?
Oh et al. [2026]	ACL	✓	~	–	?	–	~	?
Xia et al. [2026]	arXiv	✓	–	–	?	–	?	?
Liu et al. [2026b]	arXiv	✓	~	~	?	–	?	?
Sun et al. [2026b]	arXiv	✓	–	–	?	–	?	?
Xu et al. [2026b]†	arXiv	~	✓	✓	~	–	–	–
Ulmer et al. [2025]†	arXiv	✓	~	~	–	–	–	–
Zhang and Wang [2026]	J. Comput. Sci. Technol.	✓	?	–	~	–	?	?
He et al. [2026b]	Inf. Fusion	✓	?	–	–	–	?	?
<i>Reinforcement learning for LLMs</i>								
Zhang et al. [2025f]†	arXiv	–	–	~	–	–	–	–
Liu et al. [2025c]†	arXiv	–	–	–	–	–	–	–
Kaufmann et al. [2023]†	TMLR	–	?	–	–	–	–	–
Yu et al. [2026a]†	arXiv	–	?	–	–	–	–	–
<i>Closed-set decisions and judging</i>								
Gu et al. [2024]	arXiv	~	–	–	✓	–	–	–
Li et al. [2024c]	arXiv	~	–	~	✓	–	–	–
Li et al. [2024b]	arXiv	~	–	–	✓	–	–	–
<i>Acting on confidence (abstain, defer, route)</i>								
Wen et al. [2025]†	TACL	~	~	~	–	–	✓	~

Table B1, continued

Survey	Venue	CAL	RLC	RLR	DEC	TYP	ESC	HAL
Hendrickx et al. [2024]	Mach. Learn.	~	–	–	✓	–	✓	–
Strong et al. [2025]	Zenodo (preprint)	~	–	–	✓	–	✓	–
Moslem and Kelleher [2026][†]	TMLR	~	–	–	–	–	✓	–
Boudiaf et al. [2026][†]	arXiv	✓	–	~	~	–	✓	?
Hallucination								
Huang et al. [2025d]	ACM TOIS	~	–	–	–	–	–	✓
Alansari and Luqman [2026]	Comput. Sci. Rev.	~	–	~	–	–	–	✓
Zhang et al. [2025a][†]	arXiv	~	~	~	–	–	–	✓
Conformal prediction								
Campos et al. [2024][†]	TACL	✓	–	–	✓	–	~	–
Ashby et al. [2026]	Phil. Trans. A	✓	–	–	~	–	~	–
Surveys covering (✓)		24	2	3	8	0	6	5
This survey (scope)		●	●	●	●	●	●	○

C THE TRAINING CORE

Table C1 lists the 80 core works by training family, followed by the 5 theory and analysis works (Section 3.1).

Table C1. **The core: 80 works with a calibration term in the model’s own training loss or reward (S4)**, grouped by what the training signal scores; each work was assigned to one family by reading it. T1, the model is fitted to confidence targets or a calibration loss or regulariser, without a reward; T2, the reward model, a separate model in the training loop, is calibrated or made uncertainty-aware; T3, a reward or preference scores the stated confidence against the outcome, including proper scoring rules and ranking-based (AUC) confidence rewards, which target discrimination as well as calibration; T4, the signal targets the action: abstaining or refusing when unsure (supervised or rewarded); T5, the confidence reward is given per reasoning step or over meaning-equivalent answers. Family sizes: T1 23, T2 5, T3 31, T4 17, T5 4. The 5 theory and analysis works that follow are not counted in the core. * marks the 14 general-domain calibration-training works of 2025 and 2026 of which no prior survey cites more than one. Yr: year of the first public version (a citation may show a later venue year); Contract: the form of the output the confidence is attached to: typed (a probability per declared option), set (a prediction set), constrained (one output under a grammar or schema), score (one confidence number per answer, stated, read off or estimated by a separate model) or free-text (free text or a label with no per-answer probability; any confidence is a verbal hedge, an abstention, or agreement across sampled answers); Action: what an explicit rule or a trained behaviour does with the confidence when it is low (abstain; hand off = pass the case to another model or a human, which the papers call routing, cascading, deferral or escalation; retrieve = trigger retrieval; filter = remove, rewrite or revise low-confidence claims, or keep the most confident of several sampled answers; adapt compute = keep sampling, reasoning, querying or exploring, or spend a larger compute budget (the same rule stops early when the confidence is high); set = return a prediction set; none = no rule or trained behaviour acts on the confidence, including a selective-prediction curve reported only as an evaluation); Guar.: a stated distribution-free guarantee; Metric: the calibration or confidence metric reported in the main text (ECE, Brier, NLL, AUROC, selective = selective risk or accuracy at a coverage, coverage = achieved coverage of a prediction set or achieved risk under a controlled threshold, multiple = more than one of these (a reliability diagram counts with ECE; PRR with AUROC), other = a different metric, none = no confidence metric). ? = not determinable from the paper.

Work	Yr	Contribution	Contract	Action	Guar.	Metric
T1 supervised calibration objectives (23)						
Jiang et al. [2021]	2020	Finds T5, BART and GPT-2 QA probabilities are poorly calibrated and compares fine-tuning, temperature scaling and paraphrase fixes	typed	none	no	ECE
Kong et al. [2020]	2020	on- and off-manifold regularisers added to fine-tuning loss for in- and out-of-distribution calibration	typed	none	no	ECE
Lin et al. [2022]	2022	Fine-tunes GPT-3 to state calibrated numeric or word-level confidence without logits and introduces the CalibratedMath suite	score	none	no	multiple
Chaudhry et al. [2024]	2024	supervised fine-tuning on uncertainty-augmented answers to emit calibrated linguistic hedges	free-text	none	no	multiple
Chuang et al. [2024]	2024	Model emits a learned confidence token with each answer and low confidence triggers routing to a larger LLM	score	hand off	no	multiple
Han et al. [2024]	2024	LePe: fine-tunes on records of past answer correctness so the model learns to express its own confidence	score	none	no	ECE

Table C1, continued

Work	Yr	Contribution	Contract	Action	Guar.	Metric
Kapoor et al. [2024]	2024	calibration tuning on a small graded dataset of correct and incorrect answers yields generalising uncertainty estimates	score	none	no	multiple
Krishnan et al. [2024]	2024	uncertainty-aware causal LM loss grounded in decision theory raises uncertainty on wrong tokens	score	none	no	multiple
Liu et al. [2024c]	2024	self-training scheme in which the LLM learns to express uncertainty from its own consistency signals	score	retrieve	no	AUROC
Niu et al. [2024]	2024	UQ4CT: calibration loss over mixture-of-LoRA-experts functional space during PEFT	typed	none	no	ECE
Wang et al. [2024e]	2024	BLoB: jointly learns mean and covariance of LoRA weights by backpropagation for calibrated fine-tuning	typed	none	no	multiple
Zhang et al. [2024d]	2024	self-evaluation calibration tuning (SK-Tuning) then DPO on self-annotated factuality preferences	score	none	no	multiple
Eikema et al. [2025]	2025	Faithful Uncertainty Tuning: SFT on self-samples augmented with hedges matched to sample consistency	free-text	none	no	other
Hager et al. [2025]	2025	uncertainty distillation: SFT on verbalised probabilities derived from Monte Carlo semantic sampling	score	none	no	multiple
Huang et al. [2025a]	2025	Self-Calibration: distills self-consistency confidence into a single-pass confidence for efficient test-time scaling	score	adapt compute	no	multiple
Jang et al. [2025]	2025	SFT on scalar confidence labels alone elicits self-verification in chain-of-thought	score	adapt compute	no	multiple
Li et al. [2025d]*	2025	ConfTuner: tokenised Brier score loss (a proper scoring rule) for verbalised confidence fine-tuning	score	hand off	no	multiple
Steyvers et al. [2025a]	2025	tests cross-domain transfer of supervised fine-tuning for numeric confidence and pairwise confidence comparison	score	none	no	multiple
Xiao et al. [2025]	2025	calibration-aware fine-tuning plus EM-based ECE regulariser to restore calibration lost in preference alignment	typed	none	no	ECE
Zhang et al. [2025c]	2025	Direct Confidence Alignment: DPO aligns verbalised confidence with internal token-probability confidence	score	none	no	other
Zhang et al. [2025b]	2025	GrACE: fine-tunes a special confidence token whose hidden-state similarity gives calibrated confidence	score	adapt compute	no	multiple
Guo et al. [2026]	2026	post-trains end-of-reasoning verbalised confidence and an in-reasoning <uncertain> marker used to trigger RAG	score	retrieve	no	multiple
Xie et al. [2026]	2026	post-RL SFT with self-distillation restores anchor-token confidence reliability degraded by GRPO and DPO	typed	retrieve	no	multiple
T2 reward-model calibration (5)						
Leng et al. [2025]	2024	PPO-M calibrated reward modelling with confidence-augmented pairs and PPO-C reward adjustment by running mean	score	none	no	multiple

Table C1, continued

Work	Yr	Contribution	Contract	Action	Guar.	Metric
Lou et al. [2024]	2024	URM: reward model outputs mean and variance of attribute rewards to expose reward uncertainty	score	none	no	selective
Park et al. [2025a]	2025	quantile-regression calibration of process reward models used for instance-adaptive compute	score	adapt compute	no	multiple
Ma et al. [2026a]	2026	conditional optimal transport over PRM hidden states yields calibrated success quantiles	score	adapt compute	no	multiple
Pan et al. [2026a]	2026	UARM: quantile conformal uncertainty for reward models plus variance-reweighted GRPO advantages	set	none	yes	selective
T3 calibration rewards on the stated confidence (31)						
Band et al. [2024]	2024	linguistic calibration: SFT then RL rewarding generations that let a simulated reader make calibrated forecasts	free-text	none	no	ECE
Stengel-Eskin et al. [2024]	2024	listener-aware DPO where preferences reward answers a simulated listener accepts only when correct	free-text	abstain	no	multiple
Tao et al. [2024]	2024	CONQORD: PPO with quality reward plus order-preserving reward aligning verbalised confidence with response quality	score	retrieve	no	ECE
Xu et al. [2024a]	2024	SaySelf: SFT on self-reflective rationales then RL with a reward penalising overconfident wrong answers	score	none	no	ECE
Bani-Harouni et al. [2026b]*	2025	Rewarding Doubt: log scoring rule reward on verbalised confidence	score	none	no	multiple
Damani et al. [2025]*	2025	RLCR: binary correctness reward plus Brier score reward on reasoned confidence	score	none	no	multiple
Kale and Dhami [2025]	2025	KnowRL: introspection plus consensus-based self-reward to strengthen feasibility self-knowledge	free-text	none	no	other
Liu et al. [2025b]	2025	C2GSPG: group sequence policy gradient with a cross-entropy regulariser aligning sequence confidence with reward	score	none	no	multiple
Turtel et al. [2025]	2025	outcome-based RLVR on resolved prediction-market events that improves forecast calibration	score	none	no	multiple
van Niekerk et al. [2025]	2025	RLSF: preference optimisation using the frozen model own answer-span confidence as self-feedback reward	score	none	no	ECE
Wu et al. [2025a]*	2025	behaviourally calibrated RL with strictly proper scoring rule rewards enabling abstention or claim flagging	score	abstain	no	multiple
Yin et al. [2026]	2025	Deliberative Searcher: search agent RL maximising accuracy under a soft reliability (confidence) constraint	score	none	no	other
Zhang et al. [2025e]	2025	LoVeC: GRPO and DPO with calibration reward for sentence-level confidence tags in long-form output	score	none	no	multiple
Zou et al. [2025]	2025	GRPO with composite reward penalising unjustified high confidence and token-entropy spikes	score	none	no	multiple

Table C1, continued

Work	Yr	Contribution	Contract	Action	Guar.	Metric
Bani-Harouni et al. [2026a]	2026	ConRad: RL with calibration reward for report- and sentence-level verbalised confidence in radiology LVLMS	score	none	no	multiple
Cui et al. [2026]	2026	RAC: ranking-aware group loss and clean-corrupted pairwise confidence loss inside multimodal RL	score	none	no	multiple
Li et al. [2026d]	2026	CoCA: confidence-first GRPO with segmented Brier-style confidence reward and answer reward	score	none	no	multiple
Li et al. [2026a]	2026	ORCE: rank-based RL reward aligning post-hoc verbalised confidence order with sampled correctness surrogate	score	none	no	multiple
Liu et al. [2026a]	2026	RLMF: preference optimisation whose completion rankings use metacognitive feedback on self-judgement quality	score	none	no	Brier
Liu and Wang [2026]	2026	CURE: claim-level confidence reasoning trained with DPO confidence correction and RL	score	filter	no	multiple
Ma et al. [2026b]*	2026	DCPO: decouples reasoning and calibration objectives with separate advantages to avoid gradient conflict	score	none	no	multiple
Singh et al. [2026b]	2026	studies verifiable Brier rewards for aleatoric forecasting (NFL win probability) and why per-outcome reward fails	score	none	no	multiple
Tan et al. [2026]*	2026	characterises non-hackable confidence reward schemes (Brier, log hybrids) against confidence reward hacking	score	none	no	multiple
Turtel et al. [2026]	2026	compares five proper scoring rules (Brier, log, others) as terminal RL rewards for LLM forecasters	score	none	no	multiple
Wang et al. [2026a]*	2026	CAPO: logistic AUC surrogate loss for uncertainty-aware advantage estimation in GRPO	score	none	no	multiple
Xiao et al. [2026]	2026	VL-Calibration: RL rewarding separate visual and reasoning confidence in LVLMS	score	none	no	multiple
Xu et al. [2026a]	2026	DualStake: margin-clipped confidence-dependent stake rewards for evidence and answer confidence in deep research agents	score	none	no	multiple
Xuan et al. [2026]	2026	RL fine-tuning jointly rewarding task accuracy and verbalised calibration in tool-use agents	score	none	no	multiple
Yaldiz et al. [2026]*	2026	calibration-aware RLVR that directly adjusts decision-token probabilities to curb overconfidence	typed	none	no	ECE
Yang et al. [2026b]	2026	C3RL: correctness reward plus calibration reward plus dataset reference-accuracy reward for adaptive test-time scaling	score	adapt compute	no	multiple
Zhu [2026]	2026	RefGRPO: free calibration bonus contrasting agent reflection score with actual environment outcome	score	filter	no	selective

Table C1, continued

Work	Yr	Contribution	Contract	Action	Guar.	Metric
T4 abstention and refusal training (17)						
Yang et al. [2024b]	2023	alignment for honesty: fine-tuning to answer known questions and say I do not know otherwise, with honesty metrics	free-text	abstain	no	other
Zhang et al. [2023a]	2023	R-Tuning: builds refusal-aware SFT data by appending uncertainty expressions to questions the model gets wrong	score	abstain	no	other
Chen et al. [2025a]	2024	CoKE: probes knowledge boundary via internal confidence then trains the model to express it	free-text	abstain	no	other
Cheng et al. [2024]	2024	Idk dataset built per model and SFT plus preference optimisation to refuse unknown questions	free-text	abstain	no	other
Kang et al. [2025b]	2024	shows unfamiliar finetuning examples shape hallucination and proposes more reliable reward models for RL factuality finetuning	free-text	abstain	no	other
Li et al. [2025c]	2024	US-Tuning: two-stage instruction tuning for knowledge-boundary awareness then instruction adherence	free-text	abstain	no	other
Liang et al. [2024a]	2024	RLKF: knowledge-state preference data from DreamCatcher used as reward for honesty-oriented RL	free-text	abstain	no	other
Xu et al. [2024b]	2024	RLKF with a reliability reward model trained on knowledge-feedback preferences to refuse out-of-knowledge questions	free-text	abstain	no	other
Xue et al. [2025]	2024	UAlign: uncertainty estimates (confidence, semantic entropy) as features for SFT and PPO factuality alignment	free-text	abstain	no	AUROC
Mohamadi et al. [2025]*	2025	Reinforced Hesitation: ternary RLVR reward (+1, 0, -lambda) with penalty-dependent abstention frontier	free-text	abstain	no	other
Song et al. [2025]	2025	shows reinforcement finetuning erodes refusal and restores it by mixing synthetic unanswerable math	free-text	abstain	no	other
Wei et al. [2025]*	2025	TruthRL: ternary reward (correct +1, abstain 0, hallucination -1) under GRPO	free-text	abstain	no	other
Che et al. [2026]	2026	proves error-penalised abstention rewards can collapse KL-anchored RL to refuse-all and proposes a repair	score	abstain	no	Brier
Gao et al. [2026]*	2026	KARL: knowledge-boundary-aware reward from within-group statistics plus two-stage RL	free-text	abstain	no	other
Pan et al. [2026b]	2026	TIAR: trajectory-informed reweighting of the abstention advantage within GRPO ternary reward	free-text	abstain	no	other
Song et al. [2026]	2026	Visual-Idk dataset and SFT plus preference optimisation so VLMs refuse unknown visual facts	free-text	abstain	no	other
Zhai et al. [2026]*	2026	Abstain-R1: clarification-aware RLVR reward for explicit abstention plus post-refusal clarification	free-text	abstain	no	other

Table C1, continued

Work	Yr	Contribution	Contract	Action	Guar.	Metric
T5 process- and meaning-level confidence rewards (4)						
An and Xu [2025]	2025	FiSCoRe: fine-grained sample-specific semantic confidence reward guiding abstention in GRPO	free-text	abstain	no	other
He et al. [2025]	2025	MMBoundary: SFT on step-level self-rewarded confidence then RL with step calibration rewards for MLLMs	free-text	none	no	ECE
Wang et al. [2026d]*	2026	RLCM: margin-enhanced process reward on confidence at intermediate reasoning budgets plus Brier term	score	abstain	no	multiple
Yu et al. [2026b]*	2026	CSR: correctness reward plus semantic calibration reward on agreement among rollouts, no verbal confidence	free-text	none	no	multiple
Theory and analysis of training-stage calibration (not counted in the core) (5)						
Blasiok et al. [2023]	2023	relates proper-loss local optimality to smooth calibration, explaining when training yields calibration	?	none	no	none
Kalai and Vempala [2024]	2023	proves calibrated LMs must hallucinate arbitrary facts at a rate near the Good-Turing fraction of facts seen once	free-text	none	no	none
Bereket and Leskovec [2025]	2025	shows GRPO group-std normalisation causes overconfidence on stochastic outcomes; PPO and RLOO stay calibrated	score	none	no	multiple
Kalai et al. [2025]	2025	argues binary-graded training and evaluation reward guessing and calls for scoring that credits admitting uncertainty	free-text	none	no	none
Nakkiran et al. [2025]	2025	theory that semantic calibration emerges from next-token loss and breaks under RL tuning and CoT	free-text	none	no	ECE

D THE OUTPUT CONTRACT

Table D1 lists every typed work and the consumers on grammar- or schema-constrained output (Section 4).

Table D1. **The output contract: 93 works**, selected by a stated rule: every work whose confidence is attached to a typed decision (contract “typed”, 65, at any stage), and the X-stage works on grammar- or schema-constrained output (28). Cal.: ✓ when the main text reports a calibration metric (ECE, Brier score, NLL or a reliability diagram); 41 of the 65 typed works do, all 7 in the training core. The constrained-output works mostly measure validity, speed or accuracy under the constraint. Each work’s contribution is in Table C1 or Table F1. Yr: year of the first public version; Family: the training family (S4) or method family; Metric: the calibration or confidence metric reported in the main text (ECE, Brier, NLL, AUROC, selective = selective risk or accuracy at a coverage, coverage = achieved coverage of a prediction set or achieved risk under a controlled threshold, multiple = more than one of these (a reliability diagram counts with ECE; PRR with AUROC), other = a different metric, none = no confidence metric). ? = not determinable from the paper.

Work	Yr	Family	Metric	Cal.
Typed output, S0 read-off (16)				
OpenAI [2023]	2023	Calibration after alignment	ECE	✓
Huang et al. [2026]	2026	Calibration after alignment	ECE	✓
Proskurina et al. [2026]	2026	Calibration after alignment	ECE	✓
Sanz-Guerrero et al. [2026]	2026	Calibration after alignment	multiple	✓
Desai and Durrett [2020]	2020	Empirical studies of read-off calibration	ECE	✓
Dan and Roth [2021]	2021	Empirical studies of read-off calibration	ECE	✓
Ahuja et al. [2022]	2022	Empirical studies of read-off calibration	ECE	✓
Chen et al. [2023a]	2022	Empirical studies of read-off calibration	ECE	✓
Xiao et al. [2022]	2022	Empirical studies of read-off calibration	multiple	✓
Li et al. [2025h]	2023	Empirical studies of read-off calibration	multiple	✓
Zhang et al. [2023d]	2023	Empirical studies of read-off calibration	ECE	✓
Mahaut et al. [2024]	2024	Empirical studies of read-off calibration	other	–
Joshi et al. [2025]	2025	Empirical studies of read-off calibration	ECE	✓
Lovering et al. [2025]	2024	Likelihood and logit read-off	other	–
Plaut et al. [2024]	2024	Likelihood and logit read-off	multiple	✓
Wang et al. [2024d]	2024	Likelihood and logit read-off	none	–
Typed output, S1 post-hoc map (24)				
Badhe et al. [2026]	2026	Fixed maps over the read-off	multiple	✓
Zheng et al. [2023]	2023	Label-bias and option-prior debiasing	none	–
Reif and Schwartz [2024]	2024	Label-bias and option-prior debiasing	none	–
Han et al. [2022]	2022	Learned calibrators	none	–
Luo et al. [2026]	2026	Learned calibrators	multiple	✓
Liu et al. [2024b]	2024	Multicalibration and task-specific recalibration	ECE	✓
Burns et al. [2022]	2022	Probes and learned confidence heads	other	–
Azaria and Mitchell [2023]	2023	Probes and learned confidence heads	AUROC	–
Liu et al. [2023a]	2023	Probes and learned confidence heads	ECE	✓
Holtzman et al. [2021]	2021	Prompt and context calibration	none	–
Zhao et al. [2021]	2021	Prompt and context calibration	none	–
Fei et al. [2023]	2023	Prompt and context calibration	none	–
Jiang et al. [2023b]	2023	Prompt and context calibration	none	–
Liusie et al. [2023]	2023	Prompt and context calibration	none	–
Zhou et al. [2023b]	2023	Prompt and context calibration	none	–
Abbas et al. [2024]	2024	Prompt and context calibration	ECE	✓
Cho et al. [2025]	2024	Prompt and context calibration	none	–
Gundem et al. [2025]	2025	Prompt and context calibration	none	–
Guo et al. [2017]	2017	Scaling and binning	multiple	✓

Table D1, continued

Work	Yr	Family	Metric	Cal.
Kull et al. [2019]	2019	Scaling and binning	multiple	✓
Kumar et al. [2019]	2019	Scaling and binning	multiple	✓
Gupta and Ramdas [2021]	2021	Scaling and binning	ECE	✓
Shen et al. [2024]	2024	Scaling and binning	multiple	✓
Luo et al. [2025]	2025	Scaling and binning	multiple	✓
Typed output, S3 extra inference (11)				
Ling et al. [2024]	2024	Benchmarks and empirical studies	AUROC	–
Jung et al. [2024]	2024	Conformal abstention, selection and cascades	multiple	✓
Tonolini et al. [2024]	2024	Ensembles and cross-model disagreement	multiple	✓
Kruse et al. [2025]	2025	Ensembles and cross-model disagreement	multiple	✓
Schuster et al. [2021]	2021	Risk control and early exit	coverage	–
Jazbec et al. [2024]	2024	Risk control and early exit	multiple	✓
Liu et al. [2023b]	2023	Sampling consistency	other	–
Portillo Wightman et al. [2023]	2023	Sampling consistency	ECE	✓
Becker and Soatto [2024]	2024	Sampling consistency	multiple	✓
Yang et al. [2024a]	2024	Sampling consistency	multiple	✓
Zhang et al. [2024a]	2024	Sampling consistency	ECE	✓
Typed output, S4 training (7)				
Yaldiz et al. [2026]	2026	T3 calibration rewards on the stated confidence	ECE	✓
Kong et al. [2020]	2020	T1 supervised calibration objectives	ECE	✓
Niu et al. [2024]	2024	T1 supervised calibration objectives	ECE	✓
Wang et al. [2024e]	2024	T1 supervised calibration objectives	multiple	✓
Xiao et al. [2025]	2025	T1 supervised calibration objectives	ECE	✓
Xie et al. [2026]	2026	T1 supervised calibration objectives	multiple	✓
Jiang et al. [2021]	2020	T1 supervised calibration objectives	ECE	✓
Typed output, X consumers and contract (7)				
Jitkrittum et al. [2023]	2023	Deferral	none	–
Kondadadi and Ortega [2026]	2026	Deferral	none	–
Li et al. [2026c]	2026	Judge calibration	multiple	✓
Varshney and Baral [2022]	2022	Routing/cascade	none	–
Ramirez et al. [2024]	2024	Routing/cascade	none	–
Zellinger and Thomson [2025]	2025	Routing/cascade	ECE	✓
Wang et al. [2024c]	2024	Typed classification	none	–
Constrained output, X consumers and contract (28)				
Scholak et al. [2021]	2021	Constrained decoding	none	
Beurer-Kellner et al. [2023]	2022	Constrained decoding	none	
Poesia et al. [2022]	2022	Constrained decoding	none	
Geng et al. [2023]	2023	Constrained decoding	none	
Lew et al. [2023]	2023	Constrained decoding	none	
Willard and Louf [2023]	2023	Constrained decoding	none	
Beurer-Kellner et al. [2024]	2024	Constrained decoding	none	
Dong et al. [2025]	2024	Constrained decoding	none	
Geng et al. [2024b]	2024	Constrained decoding	none	
Koo et al. [2024]	2024	Constrained decoding	none	
Park et al. [2024]	2024	Constrained decoding	none	
Ugare et al. [2024]	2024	Constrained decoding	none	
Banerjee et al. [2025]	2025	Constrained decoding	none	
Li et al. [2026b]	2025	Constrained decoding	none	
Loula et al. [2025]	2025	Constrained decoding	other	
Mündler et al. [2025]	2025	Constrained decoding	none	
Park et al. [2025b]	2025	Constrained decoding	none	
Reddy et al. [2026]	2026	Constrained decoding	none	

Table D1, continued

Work	Yr	Family	Metric	Cal.
Tang et al. [2023]	2023	Structured-output benchmark	none	
Shorten et al. [2024]	2024	Structured-output benchmark	none	
Tam et al. [2024]	2024	Structured-output benchmark	none	
Xia et al. [2024a]	2024	Structured-output benchmark	none	
Geng et al. [2025]	2025	Structured-output benchmark	none	
Yang et al. [2025a]	2025	Structured-output benchmark	none	
Fan [2026]	2026	Structured-output benchmark	none	
Galeone et al. [2026]	2026	Structured-output benchmark	none	
Lee et al. [2026a]	2026	Structured-output benchmark	none	
Singh et al. [2026a]	2026	Structured-output benchmark	none	

E PROVENANCE OF THE GALLERY PANELS

Table E1 gives, for every panel of Figures 1, 8 and 9, the source figure and page in the cited paper and the licence under which it is reproduced, together with the rule by which panels were chosen.

Table E1. **Provenance of the gallery panels** in Figures 1, 8 and 9: the source figure and page in the cited arXiv paper, and the licence under which it is reproduced. Panels are chosen by a fixed rule: only papers whose licence permits reuse (CC BY); per stage, one panel per method family (per training family at S4), taking the largest families first and, in each, the newest figure that legibly shows the method or a result, at most 3 per stage; families with no reusable figure that legibly shows the method or a result are skipped. Figure 8: No reusable figure passed the visual check for S0 read-off. At S4, T5 had a figure but was cut by the cap of 3. T2 has no reusable figure that passed the check. Figure 9: At S4, T2, T4, T5 have no reusable figure that passed the check.

Work	Source	Panel	Licence	Used in Fig.
S0 read-off				
Sanz-Guerrero et al. [2026]	arXiv:2606.03437	Fig. 4, p. 6	CC BY 4.0	9
Stengel-Eskin and Van Durme [2023]	arXiv:2211.07443	Fig. 1, p. 6	CC BY 4.0	1, 9
S1 post-hoc map				
Gundem et al. [2025]	arXiv:2505.23783	Fig. 2, p. 5	CC BY 4.0	8
Liu et al. [2023a]	arXiv:2312.03729	Fig. 2, p. 3	CC BY 4.0	9
Lu [2026]	arXiv:2607.16799	Fig. 1, p. 2	CC BY 4.0	1, 8
Luo et al. [2025]	arXiv:2505.16690	Fig. 1, p. 3	CC BY 4.0	9
Luo et al. [2026]	arXiv:2601.04277	Fig. 1, p. 2	CC BY 4.0	8
Luo et al. [2026]	arXiv:2601.04277	Fig. 5, p. 15	CC BY 4.0	9
S2 prompt elicitation				
Ferrer et al. [2026]	arXiv:2603.29559	Fig. 1, p. 10	CC BY 4.0	9
Kumar et al. [2024]	arXiv:2405.16282	Fig. 1, p. 1	CC BY 4.0	8
Marashian et al. [2026]	arXiv:2606.17234	Fig. 1, p. 1	CC BY 4.0	1, 8
Marashian et al. [2026]	arXiv:2606.17234	Fig. 2, p. 5	CC BY 4.0	9
Yang et al. [2024c]	arXiv:2404.09127	Fig. 3, p. 11	CC BY 4.0	9
S3 extra inference				
Chen et al. [2025b]	arXiv:2502.20285	Fig. 3, p. 8	CC BY 4.0	8
Cherian et al. [2024]	arXiv:2406.09714	Fig. 8, p. 18	CC BY 4.0	9
Cohen et al. [2023]	arXiv:2305.13281	Fig. 2, p. 3	CC BY 4.0	8
Doula et al. [2025]	arXiv:2505.09427	Fig. 1, p. 3	CC BY 4.0	1, 8
Ulmer et al. [2024b]	arXiv:2402.00707	Fig. 2, p. 6	CC BY 4.0	9
Wang et al. [2026b]	arXiv:2602.03814	Fig. 4, p. 7	CC BY 4.0	9
S4 training				
Gao et al. [2026]	arXiv:2604.22779	Fig. 4, p. 4	CC BY 4.0	1, 8
Guo et al. [2026]	arXiv:2604.05306	Fig. 1, p. 2	CC BY 4.0	1, 8
Steyvers et al. [2025a]	arXiv:2510.05126	Fig. 3, p. 7	CC BY 4.0	9
Tan et al. [2026]	arXiv:2607.04332	Fig. 4, p. 8	CC BY 4.0	9
Yang et al. [2026b]	arXiv:2607.01612	Fig. 1, p. 3	CC BY 4.0	1, 8
X consumers and contract				
Aggarwal et al. [2023]	arXiv:2305.11860	Fig. 1, p. 2	CC BY 4.0	8
Fanconi and van der Schaar [2025]	arXiv:2506.11887	Fig. 3, p. 8	CC BY 4.0	9
Soiffer et al. [2025]	arXiv:2509.21837	Fig. 1, p. 1	CC BY 4.0	8
Soiffer et al. [2025]	arXiv:2509.21837	Fig. 2, p. 5	CC BY 4.0	9
Srinivasan et al. [2024]	arXiv:2402.15610	Fig. 3, p. 5	CC BY 4.0	9
Wang et al. [2025e]	arXiv:2505.04016	Fig. 1, p. 1	CC BY 4.0	8

F THE LANDSCAPE

Table F1 lists the 323 works outside the training core, grouped by the stage at which calibration enters and then by method family. The 80 core works and the 5 theory and analysis works are in Table C1.

Table F1. **The landscape: the 323 works outside the training core**, grouped by the stage at which calibration enters (Table A1) and then by method family, oldest first; the 80 core works and the 5 theory and analysis works are listed in Table C1. Yr: year of the first public version (a citation may show a later venue year); Contract: the form of the output the confidence is attached to: typed (a probability per declared option), set (a prediction set), constrained (one output under a grammar or schema), score (one confidence number per answer, stated, read off or estimated by a separate model) or free-text (free text or a label with no per-answer probability; any confidence is a verbal hedge, an abstention, or agreement across sampled answers); Access: white (needs logits or weights), black (text only) or either; Samples: the method needs repeated inference; Guar.: a stated distribution-free guarantee; Trained: yes when calibration is a training objective of the model itself, partial when a separate model (a probe, calibrator, router or reward model) is trained; Metric: the calibration or confidence metric reported in the main text (ECE, Brier, NLL, AUROC, selective = selective risk or accuracy at a coverage, coverage = achieved coverage of a prediction set or achieved risk under a controlled threshold, multiple = more than one of these (a reliability diagram counts with ECE; PRR with AUROC), other = a different metric, none = no confidence metric); Task: the task evaluated; Action: what an explicit rule or a trained behaviour does with the confidence when it is low (abstain; hand off = pass the case to another model or a human, which the papers call routing, cascading, deferral or escalation; retrieve = trigger retrieval; filter = remove, rewrite or revise low-confidence claims, or keep the most confident of several sampled answers; adapt compute = keep sampling, reasoning, querying or exploring, or spend a larger compute budget (the same rule stops early when the confidence is high); set = return a prediction set; none = no rule or trained behaviour acts on the confidence, including a selective-prediction curve reported only as an evaluation). ? = not determinable from the paper.

Work	Yr	Contract	Access	Samples	Guar.	Trained	Metric	Task	Action
S0 read-off (25)									
<i>Calibration after alignment</i>									
OpenAI [2023]	2023	typed	white	no	no	no	ECE	QA	none
Zhu et al. [2023]	2023	score	white	no	no	no	ECE	QA	none
Huang et al. [2026]	2026	typed	white	no	no	no	ECE	QA	none
Proskurina et al. [2026]	2026	typed	white	no	no	no	ECE	QA	none
Sanz-Guerrero et al. [2026]	2026	typed	white	no	no	no	multiple	QA	none
<i>Empirical studies of read-off calibration</i>									
Desai and Durrett [2020]	2020	typed	white	no	no	no	ECE	classification	none
Wang et al. [2020]	2020	score	white	no	no	no	ECE	generation	none
Dan and Roth [2021]	2021	typed	white	no	no	no	ECE	classification	none
Xiao and Wang [2021]	2021	score	white	no	no	no	other	generation	none
Ahuja et al. [2022]	2022	typed	white	no	no	no	ECE	classification	none
Chen et al. [2023a]	2022	typed	white	no	no	no	ECE	classification	none
Si et al. [2022]	2022	score	white	no	no	no	ECE	QA	none
Stengel-Eskin and Van Durme [2023]	2022	constrained	white	no	no	no	ECE	code	none
Xiao et al. [2022]	2022	typed	white	no	no	no	multiple	classification	none
Li et al. [2025h]	2023	typed	white	no	no	no	multiple	classification	none
Zhang et al. [2023d]	2023	typed	white	no	no	no	ECE	classification	none
Mahaut et al. [2024]	2024	typed	either	no	no	no	other	QA	none
Stolfo et al. [2024]	2024	score	white	no	no	no	other	generation	none
Joshi et al. [2025]	2025	typed	white	no	no	no	ECE	QA	none
<i>Likelihood and logit read-off</i>									
Ott et al. [2018]	2018	score	white	yes	no	no	other	generation	none
Lovering et al. [2025]	2024	typed	white	no	no	no	other	reasoning	none
Plaut et al. [2024]	2024	typed	white	no	no	no	multiple	QA	abstain
Spieß et al. [2024]	2024	score	white	no	no	no	multiple	code	none
Wang et al. [2024d]	2024	typed	white	no	no	no	none	QA	none
<i>Self-evaluation and read-off calibration</i>									
Kadavath et al. [2022]	2022	score	white	no	no	partial	multiple	QA	none

Table F1, continued

Work	Yr	Contract	Access	Samples	Guar.	Trained	Metric	Task	Action
S1 post-hoc map (52)									
<i>Fixed maps over the read-off</i>									
Zhang et al. [2023c]	2023	score	white	no	no	no	other	generation	none
Bakman et al. [2024]	2024	score	white	no	no	partial	AUROC	QA	none
Fadeeva et al. [2024]	2024	score	white	no	no	no	AUROC	generation	none
Li et al. [2025b]	2025	score	white	no	no	no	multiple	reasoning	none
Vazhentsev et al. [2026]	2025	score	white	no	no	no	AUROC	generation	none
Badhe et al. [2026]	2026	typed	white	no	no	no	multiple	classification	none
<i>Label-bias and option-prior debiasing</i>									
Zheng et al. [2023]	2023	typed	white	no	no	no	none	QA	none
Reif and Schwartz [2024]	2024	typed	white	no	no	no	none	classification	none
<i>Learned calibrators</i>									
Ye and Durrett [2022]	2021	score	black	no	no	partial	multiple	QA	none
Zhang et al. [2021]	2021	score	white	no	no	partial	selective	QA	filter
Dhuliawala et al. [2022]	2022	score	white	no	no	partial	multiple	QA	none
Han et al. [2022]	2022	typed	white	no	no	partial	none	classification	none
Liu et al. [2024a]	2024	score	white	no	no	partial	multiple	QA	none
Liu et al. [2024d]	2024	score	either	no	no	partial	AUROC	QA	none
Ulmer et al. [2024a]	2024	score	black	no	no	partial	multiple	QA	none
Radharapu et al. [2025]	2025	score	white	no	no	partial	ECE	judging	none
Luo et al. [2026]	2026	typed	white	no	no	no	multiple	QA	none
<i>Multicalibration and task-specific recalibration</i>									
Detommaso et al. [2024]	2024	score	either	no	no	no	multiple	QA	none
Liu et al. [2024b]	2024	typed	white	no	no	no	ECE	classification	none
<i>Probes and learned confidence heads</i>									
Kumar and Sarawagi [2019]	2019	score	white	no	no	partial	ECE	generation	none
Mielke et al. [2022]	2020	free-text	either	no	no	partial	multiple	QA	none
Burns et al. [2022]	2022	typed	white	no	no	partial	other	classification	none
Azaria and Mitchell [2023]	2023	typed	white	no	no	partial	AUROC	classification	none
CH-Wang et al. [2024]	2023	score	white	no	no	partial	other	generation	none
Liu et al. [2023a]	2023	typed	white	no	no	partial	ECE	classification	none
Liu et al. [2023c]	2023	score	white	no	no	partial	multiple	generation	none
Slobodkin et al. [2023]	2023	free-text	white	no	no	partial	other	QA	none
Beigi et al. [2024a]	2024	score	white	no	no	partial	ECE	QA	none
Du et al. [2024]	2024	score	white	no	no	partial	AUROC	QA	none
Gottesman and Geva [2024]	2024	score	white	no	no	partial	other	QA	none
Ji et al. [2024]	2024	score	white	no	no	partial	other	generation	none
Kossen et al. [2024]	2024	score	white	no	no	partial	AUROC	QA	none
Orgad et al. [2024]	2024	score	white	no	no	partial	AUROC	QA	none
Su et al. [2024b]	2024	score	white	no	no	partial	AUROC	generation	none
Khanmohammadi et al. [2025]	2025	score	white	no	no	partial	multiple	QA	none
Lu [2026]	2026	score	white	no	no	partial	AUROC	reasoning	none
<i>Prompt and context calibration</i>									
Holtzman et al. [2021]	2021	typed	white	no	no	no	none	classification	none
Zhao et al. [2021]	2021	typed	white	no	no	no	none	classification	none
Fei et al. [2023]	2023	typed	white	no	no	no	none	classification	none
Jiang et al. [2023b]	2023	typed	white	yes	no	no	none	classification	none
Liusie et al. [2023]	2023	typed	white	no	no	no	none	classification	none
Zhou et al. [2023b]	2023	typed	white	no	no	no	none	classification	none
Abbas et al. [2024]	2024	typed	white	no	no	no	ECE	classification	none
Cho et al. [2025]	2024	typed	white	no	no	no	none	classification	none
Gundem et al. [2025]	2025	typed	white	no	no	no	none	classification	none
<i>Scaling and binning</i>									
Guo et al. [2017]	2017	typed	white	no	no	no	multiple	classification	none
Kull et al. [2019]	2019	typed	white	no	no	no	multiple	classification	none
Kumar et al. [2019]	2019	typed	either	no	yes	no	multiple	classification	none
Gupta and Ramdas [2021]	2021	typed	either	no	yes	no	ECE	classification	none
Shen et al. [2024]	2024	typed	white	no	no	partial	multiple	QA	none
Xie et al. [2024a]	2024	score	white	no	no	partial	multiple	QA	none
Luo et al. [2025]	2025	typed	white	no	no	no	multiple	QA	none
S2 prompt elicitation (54)									
<i>Abstention prompting</i>									
Zong et al. [2026]	2026	score	black	no	no	no	multiple	QA	abstain
<i>Introspection, deliberation and task-specific elicitation</i>									
Sun et al. [2024]	2024	constrained	either	no	no	partial	multiple	other	none
Yang et al. [2024c]	2024	score	black	yes	no	no	multiple	QA	none
Mei et al. [2025]	2025	score	black	no	no	no	ECE	reasoning	none
<i>Linguistic hedges and epistemic markers</i>									
Sileo and Moens [2022]	2022	free-text	white	no	no	no	none	reasoning	none

Table F1, continued

Work	Yr	Contract	Access	Samples	Guar.	Trained	Metric	Task	Action
Ren et al. [2023b]	2023	score	white	yes	no	no	multiple	QA	none
Singh et al. [2023]	2023	score	black	no	no	no	other	other	none
Yin et al. [2023]	2023	free-text	black	no	no	no	other	QA	none
Ghafouri et al. [2024]	2024	free-text	either	no	no	partial	other	QA	none
Huang et al. [2024b]	2024	score	black	yes	no	no	multiple	generation	hand off
Kumar et al. [2024]	2024	score	white	no	no	no	other	QA	none
Li et al. [2024a]	2024	free-text	black	no	no	no	other	reasoning	filter
Li et al. [2024e]	2024	score	either	?	no	no	multiple	QA	none
Pawitan and Holmes [2024]	2024	score	black	no	no	no	ECE	reasoning	none
Sicilia et al. [2024]	2024	score	either	no	no	partial	Brier	QA	none
Tang et al. [2024]	2024	free-text	black	no	no	no	other	other	none
Xue et al. [2024]	2024	score	either	?	no	no	multiple	QA	none
Zhao et al. [2024b]	2024	score	black	no	no	no	ECE	QA	retrieve
Liu et al. [2025e]	2025	free-text	black	no	no	no	ECE	QA	none
Sun et al. [2025]	2025	score	black	no	no	no	ECE	reasoning	none
Liu and Cohan [2026]	2026	free-text	black	no	no	no	other	QA	none
<i>Self-evaluation prompts</i>									
Tian et al. [2025]	2025	score	black	no	no	no	multiple	judging	none
<i>Stated confidence that triggers an action</i>									
Ni et al. [2024a]	2024	free-text	black	no	no	no	other	QA	retrieve
Ferrer et al. [2026]	2026	score	black	no	no	no	multiple	judging	hand off
<i>Verbalised numeric confidence</i>									
Tanneru et al. [2023]	2023	score	black	yes	no	no	other	QA	none
Tian et al. [2023]	2023	score	black	no	no	no	multiple	QA	none
Xiong et al. [2023]	2023	score	black	yes	no	no	multiple	reasoning	none
Zhou et al. [2023a]	2023	free-text	black	no	no	no	ECE	QA	none
Belem et al. [2024]	2024	free-text	black	no	no	no	other	other	none
Groot and Valdenegro-Toro [2024]	2024	score	black	no	no	no	ECE	other	none
Lee et al. [2024a]	2024	free-text	black	no	no	no	none	judging	none
Ni et al. [2024b]	2024	score	either	no	no	no	other	QA	none
Rivera et al. [2024]	2024	score	black	yes	no	no	multiple	classification	none
Wang et al. [2024f]	2024	score	black	no	no	no	multiple	classification	none
Wei et al. [2024]	2024	score	black	yes	no	no	ECE	QA	none
Yang et al. [2024d]	2024	score	black	no	no	no	ECE	QA	none
Yona et al. [2024]	2024	free-text	black	yes	no	no	other	QA	none
Chhikara [2025]	2025	score	black	no	no	no	ECE	QA	none
Liu et al. [2025d]	2025	free-text	black	yes	no	no	other	QA	none
Naderi et al. [2025]	2025	score	black	no	no	no	multiple	QA	none
Obadinma and Zhu [2025]	2025	score	black	no	no	no	other	QA	none
Tao et al. [2025b]	2025	free-text	black	no	no	partial	multiple	QA	none
Tao et al. [2025a]	2025	score	either	no	no	no	multiple	QA	none
Wang and Stengel-Eskin [2025]	2025	score	black	yes	no	no	multiple	QA	none
Wang et al. [2025d]	2025	score	black	no	no	no	multiple	QA	none
Yoon et al. [2025]	2025	score	black	no	no	no	multiple	reasoning	none
Yu et al. [2025b]	2025	score	black	yes	no	no	multiple	agentic	none
Zhou et al. [2025]	2025	score	black	yes	no	no	multiple	QA	none
Cacioli [2026]	2026	score	black	no	no	no	AUROC	QA	none
ffrench Constant et al. [2026]	2026	score	black	no	no	no	Brier	QA	none
Hsiao [2026]	2026	score	black	no	no	no	ECE	judging	none
Kumaran et al. [2026]	2026	score	white	no	no	no	multiple	QA	none
Marashian et al. [2026]	2026	score	black	no	no	no	ECE	generation	none
Yang et al. [2026a]	2026	score	black	no	no	no	multiple	QA	none
S3 extra inference (121)									
<i>Benchmarks and empirical studies</i>									
Fadeeva et al. [2023]	2023	free-text	either	?	no	no	AUROC	generation	none
Huang et al. [2025c]	2023	free-text	either	yes	no	no	AUROC	other	none
Huang et al. [2024a]	2024	free-text	either	yes	no	no	AUROC	QA	none
Ling et al. [2024]	2024	typed	white	yes	no	no	AUROC	classification	none
Vashurin et al. [2025]	2024	free-text	either	?	no	no	multiple	generation	none
<i>Conformal abstention, selection and cascades</i>									
Fisch et al. [2020]	2020	set	white	no	yes	no	coverage	QA	set
Lee et al. [2024b]	2023	free-text	either	no	yes	partial	coverage	generation	abstain
Yadkori et al. [2024b]	2024	free-text	black	yes	yes	no	coverage	QA	abstain
Gui et al. [2024]	2024	free-text	either	yes	yes	partial	coverage	generation	abstain
Jung et al. [2024]	2024	typed	either	yes	yes	no	multiple	judging	hand off
Tayebati et al. [2025]	2025	set	white	no	no	partial	multiple	QA	abstain
<i>Conformal factuality and claim filtering</i>									
Cherian et al. [2024]	2024	free-text	either	yes	yes	partial	coverage	generation	filter
Mohri and Hashimoto [2024]	2024	free-text	either	yes	yes	no	coverage	generation	filter
Feng et al. [2025]	2025	free-text	black	yes	yes	no	coverage	QA	filter

Table F1, continued

Work	Yr	Contract	Access	Samples	Guar.	Trained	Metric	Task	Action
Jiang et al. [2025]	2025	free-text	black	yes	yes	partial	coverage	QA	filter
Lee et al. [2025]	2025	free-text	either	?	yes	no	coverage	QA	abstain
Rubin-Toles et al. [2025]	2025	free-text	black	yes	yes	no	coverage	reasoning	filter
<i>Conformal prediction sets</i>									
Dey et al. [2022]	2021	set	white	no	yes	no	coverage	classification	set
Fisch et al. [2021]	2021	set	white	no	yes	partial	coverage	classification	set
Choubey et al. [2022]	2022	set	white	no	yes	no	coverage	classification	set
Deuschmann et al. [2023]	2023	set	white	no	yes	no	coverage	generation	set
Kumar et al. [2023]	2023	set	white	no	yes	no	coverage	QA	set
Li et al. [2024d]	2023	set	either	yes	yes	no	coverage	QA	set
Quach et al. [2023]	2023	set	either	yes	yes	no	coverage	generation	set
Ravfogel et al. [2023]	2023	set	white	no	yes	no	coverage	generation	set
Ren et al. [2023a]	2023	set	white	no	yes	no	coverage	robotics	hand off
Wang et al. [2025f]	2023	set	white	no	yes	no	coverage	robotics	hand off
Zerva and Martins [2024]	2023	set	white	?	yes	no	coverage	judging	set
Kaur et al. [2024]	2024	set	black	yes	yes	no	multiple	QA	set
Kostumov et al. [2024]	2024	set	white	no	yes	no	multiple	QA	set
Liang et al. [2024b]	2024	set	white	no	yes	no	coverage	robotics	hand off
Liu and Wu [2024]	2024	free-text	black	yes	yes	partial	multiple	generation	filter
Ren et al. [2024]	2024	set	white	no	yes	no	other	robotics	adapt compute
Su et al. [2024a]	2024	set	black	yes	yes	no	coverage	QA	set
Ulmer et al. [2024b]	2024	set	white	no	yes	no	coverage	generation	set
Vishwakarma et al. [2024]	2024	set	white	no	yes	partial	coverage	QA	set
Wang et al. [2024a]	2024	set	black	yes	yes	no	multiple	QA	set
Wang et al. [2025c]	2024	set	white	no	yes	no	coverage	robotics	hand off
Ye et al. [2024]	2024	set	white	no	yes	no	multiple	QA	set
Zhang et al. [2024b]	2024	set	either	no	yes	no	coverage	QA	set
Chan et al. [2025]	2025	set	?	?	yes	no	coverage	QA	adapt compute
Doula et al. [2025]	2025	set	black	yes	yes	no	coverage	robotics	hand off
Huang et al. [2025b]	2025	set	white	no	yes	no	multiple	QA	hand off
Kaur et al. [2025]	2025	set	black	yes	yes	no	none	reasoning	set
Kuwahara et al. [2025]	2025	free-text	either	no	yes	no	coverage	generation	filter
Noorani et al. [2025]	2025	set	black	yes	yes	no	coverage	generation	set
Shahrokhi et al. [2025]	2025	set	black	yes	yes	no	coverage	code	set
Sheng et al. [2025]	2025	set	either	no	yes	no	coverage	judging	set
Sundarsingh et al. [2025]	2025	set	?	?	yes	no	coverage	robotics	hand off
Wang et al. [2025g]	2025	set	black	yes	yes	no	coverage	QA	abstain
Yang and Liu [2025]	2025	set	black	yes	yes	no	multiple	QA	set
Emmenegger et al. [2026]	2026	free-text	?	yes	yes	no	coverage	generation	abstain
Zhang et al. [2026b]	2026	set	?	?	yes	no	none	robotics	adapt compute
<i>Ensembles and cross-model disagreement</i>									
Dong et al. [2018]	2018	constrained	white	yes	no	partial	selective	code	none
Hou et al. [2024]	2023	free-text	black	yes	no	no	AUROC	QA	hand off
Tonolini et al. [2024]	2024	typed	black	yes	no	no	multiple	classification	none
Kruse et al. [2025]	2025	typed	?	yes	no	no	multiple	classification	none
Hamidieh et al. [2026]	2026	free-text	black	yes	no	no	multiple	generation	none
<i>Internal-state signals</i>									
Fomicheva et al. [2020]	2020	score	white	yes	no	no	other	generation	none
Chen et al. [2024]	2024	free-text	white	yes	no	no	AUROC	QA	none
Li et al. [2026e]	2025	free-text	black	yes	no	no	AUROC	QA	none
<i>Long-form and claim-level</i>									
Da et al. [2024]	2024	free-text	black	yes	no	no	multiple	QA	none
Jiang et al. [2024]	2024	free-text	black	yes	no	no	AUROC	generation	filter
Zhang et al. [2025d]	2024	free-text	either	yes	no	no	multiple	generation	filter
Zhang et al. [2024c]	2024	free-text	black	yes	no	no	selective	generation	filter
Fan et al. [2026]	2026	free-text	black	yes	no	no	AUROC	generation	none
<i>Risk control and early exit</i>									
Angelopoulos et al. [2025]	2021	set	either	no	yes	no	coverage	other	set
Schuster et al. [2021]	2021	typed	white	no	yes	partial	coverage	classification	adapt compute
Angelopoulos et al. [2022]	2022	set	either	no	yes	no	coverage	other	set
Schuster et al. [2022]	2022	score	white	no	yes	partial	other	generation	adapt compute
Zollo et al. [2023]	2023	free-text	black	no	yes	no	coverage	generation	none
Jazbec et al. [2024]	2024	typed	white	no	yes	no	multiple	other	adapt compute
Kang et al. [2024]	2024	set	black	yes	yes	no	coverage	generation	set
Kladny et al. [2024]	2024	set	black	yes	yes	no	coverage	generation	set
Wang et al. [2024b]	2024	set	black	yes	yes	no	coverage	QA	set
Chen et al. [2025b]	2025	free-text	black	yes	yes	no	coverage	generation	filter
Ke et al. [2025]	2025	set	black	yes	yes	no	multiple	QA	set
Overman and Bayati [2025]	2025	score	black	no	yes	no	coverage	QA	hand off
Wang et al. [2025b]	2025	set	black	yes	yes	no	coverage	QA	abstain

Table F1, continued

Work	Yr	Contract	Access	Samples	Guar.	Trained	Metric	Task	Action
Ye [2025]	2025	set	black	yes	yes	no	coverage	QA	set
Feng et al. [2026]	2026	constrained	either	?	yes	no	other	agentic	hand off
Wang et al. [2026b]	2026	score	white	no	yes	no	coverage	reasoning	abstain
<i>Sampling consistency</i>									
Wang et al. [2023b]	2022	free-text	black	yes	no	no	other	reasoning	none
Agrawal et al. [2023]	2023	free-text	black	yes	no	no	multiple	other	none
Chen and Mueller [2023]	2023	free-text	black	yes	no	no	AUROC	QA	filter
Cohen et al. [2023]	2023	free-text	black	yes	no	no	other	QA	none
Cole et al. [2023]	2023	free-text	black	yes	no	no	multiple	QA	abstain
Friel and Sanyal [2023]	2023	free-text	black	yes	no	no	AUROC	judging	none
Liu et al. [2023b]	2023	typed	either	yes	no	no	other	judging	none
Manakul et al. [2023]	2023	free-text	black	yes	no	no	other	generation	none
Mündler et al. [2023]	2023	free-text	black	yes	no	no	other	generation	filter
Portillo Wightman et al. [2023]	2023	typed	either	yes	no	no	ECE	QA	none
Zhao et al. [2024a]	2023	free-text	black	yes	no	no	other	QA	none
Becker and Soatto [2024]	2024	typed	white	yes	no	no	multiple	QA	none
Gao et al. [2024]	2024	free-text	black	yes	no	no	ECE	QA	none
Lyu et al. [2025]	2024	free-text	black	yes	no	no	multiple	reasoning	none
Yadkori et al. [2024a]	2024	free-text	black	yes	no	no	selective	QA	abstain
Yang et al. [2024a]	2024	typed	black	yes	no	no	multiple	QA	none
Zhang et al. [2024a]	2024	typed	either	yes	no	no	ECE	QA	none
Tan et al. [2025]	2025	free-text	white	yes	no	partial	AUROC	QA	none
<i>Semantic entropy and meaning clustering</i>									
Kuhn et al. [2023]	2023	free-text	white	yes	no	no	AUROC	QA	none
Lin et al. [2024]	2023	free-text	black	yes	no	no	multiple	QA	none
Zhang et al. [2023b]	2023	free-text	black	yes	no	no	AUROC	QA	none
Aichberger et al. [2025]	2024	free-text	white	yes	no	no	AUROC	QA	none
Ao et al. [2024]	2024	free-text	black	yes	no	no	multiple	QA	none
Farquhar et al. [2024]	2024	free-text	either	yes	no	no	multiple	QA	none
Grewal et al. [2024]	2024	free-text	either	yes	no	partial	AUROC	QA	none
Nikitin et al. [2024]	2024	free-text	either	yes	no	no	multiple	QA	none
Qiu and Mäikkulainen [2024]	2024	free-text	white	yes	no	no	AUROC	QA	none
Zhang et al. [2024e]	2024	free-text	black	yes	no	no	other	QA	none
Ciosek et al. [2025]	2025	free-text	?	yes	no	no	AUROC	QA	adapt compute
Ma et al. [2025]	2025	free-text	white	yes	no	no	AUROC	QA	none
Nguyen et al. [2025]	2025	free-text	black	yes	no	no	AUROC	generation	none
Sharma and David [2025]	2025	free-text	?	yes	no	no	other	code	abstain
Zhao et al. [2026]	2025	free-text	black	yes	no	no	multiple	QA	none
Bouchard et al. [2026]	2026	free-text	black	yes	no	no	AUROC	code	none
He et al. [2026a]	2026	free-text	black	yes	no	no	AUROC	code	abstain
Sun et al. [2026a]	2026	free-text	?	yes	no	no	AUROC	QA	adapt compute
<i>Token-level aggregation</i>									
Malinin and Gales [2020]	2020	free-text	white	yes	no	no	AUROC	generation	none
Duan et al. [2024]	2023	free-text	white	yes	no	no	AUROC	QA	none
Wang et al. [2025a]	2024	free-text	white	yes	no	no	AUROC	QA	filter
S4 training (80): listed in Table C1									
Theory and analysis (5): listed in Table C1									
X consumers and contract (71): output-contract works detailed in Table D1									
<i>Abstention</i>									
Kamath et al. [2020]	2020	score	white	no	no	partial	selective	QA	abstain
Feng et al. [2024]	2024	free-text	black	no	no	no	ECE	QA	abstain
Srinivasan et al. [2024]	2024	score	white	no	no	no	multiple	QA	adapt compute
<i>Adaptive compute</i>									
Aggarwal et al. [2023]	2023	free-text	black	yes	no	no	none	reasoning	adapt compute
Jiang et al. [2023a]	2023	score	white	no	no	no	none	QA	retrieve
Damani et al. [2024]	2024	score	white	yes	no	partial	multiple	reasoning	hand off
Li et al. [2024f]	2024	free-text	black	yes	no	no	none	reasoning	adapt compute
Fu et al. [2025]	2025	score	white	yes	no	no	other	reasoning	filter
Yang et al. [2025b]	2025	score	white	no	no	no	other	reasoning	adapt compute
<i>Confidence-weighted selection</i>									
Taubenfeld et al. [2025]	2025	score	either	yes	no	no	multiple	reasoning	filter
<i>Constrained decoding</i>									
Scholak et al. [2021]	2021	constrained	white	no	no	no	none	code	none
Beurer-Kellner et al. [2023]	2022	constrained	white	no	no	no	none	generation	none
Poesia et al. [2022]	2022	constrained	white	no	no	partial	none	code	none
Geng et al. [2023]	2023	constrained	white	no	no	no	none	other	none
Lew et al. [2023]	2023	constrained	white	yes	no	no	none	generation	none
Willard and Louf [2023]	2023	constrained	white	no	no	no	none	generation	none

Table F1, continued

Work	Yr	Contract	Access	Samples	Guar.	Trained	Metric	Task	Action
Beurer-Kellner et al. [2024]	2024	constrained	white	no	no	no	none	generation	none
Dong et al. [2025]	2024	constrained	white	no	no	no	none	generation	none
Geng et al. [2024b]	2024	constrained	black	no	no	no	none	other	none
Koo et al. [2024]	2024	constrained	white	no	no	no	none	generation	none
Park et al. [2024]	2024	constrained	white	yes	no	no	none	code	none
Ugare et al. [2024]	2024	constrained	white	no	no	no	none	code	none
Banerjee et al. [2025]	2025	constrained	white	no	no	no	none	reasoning	none
Li et al. [2026b]	2025	constrained	white	no	no	no	none	code	none
Loula et al. [2025]	2025	constrained	white	yes	no	no	other	code	none
Mündler et al. [2025]	2025	constrained	white	no	no	no	none	code	none
Park et al. [2025b]	2025	constrained	white	no	no	no	none	generation	none
Wang et al. [2025e]	2025	free-text	either	no	no	partial	none	generation	none
Reddy et al. [2026]	2026	constrained	white	no	no	no	none	reasoning	none
<i>Deferral</i>									
Jitkrittum et al. [2023]	2023	typed	white	no	no	partial	none	classification	hand off
Fanconi and van der Schaar [2025]	2025	score	either	?	no	partial	none	QA	hand off
Wu et al. [2025b]	2025	score	black	?	no	no	other	reasoning	hand off
Kondadadi and Ortega [2026]	2026	typed	either	?	no	partial	none	classification	hand off
<i>Human reliance on model output and stated confidence</i>									
Si et al. [2024]	2023	free-text	black	no	no	no	other	QA	none
Kim et al. [2024]	2024	free-text	black	no	no	no	other	QA	none
Steyvers et al. [2025b]	2024	free-text	either	no	no	no	multiple	QA	none
Zhou et al. [2024b]	2024	free-text	black	no	no	no	other	QA	none
Zhou et al. [2024a]	2024	free-text	black	no	no	no	other	QA	none
Rathi et al. [2025]	2025	free-text	black	no	no	no	other	QA	none
Villavicencio et al. [2026]	2026	score	either	?	no	no	other	QA	none
<i>Judge calibration</i>									
Liu et al. [2023d]	2023	score	black	no	no	no	none	judging	none
Wang et al. [2023a]	2023	free-text	black	yes	no	no	none	judging	hand off
Lee et al. [2026b]	2025	free-text	black	no	no	no	coverage	judging	none
Li et al. [2026c]	2026	typed	black	no	no	no	multiple	judging	hand off
<i>Routing/cascade</i>									
Varshney and Baral [2022]	2022	typed	white	no	no	no	none	classification	hand off
Aggarwal et al. [2024]	2023	score	black	yes	no	partial	none	QA	hand off
Chen et al. [2023b]	2023	score	black	no	no	partial	none	QA	hand off
Yue et al. [2024]	2023	free-text	black	yes	no	no	other	reasoning	hand off
Ding et al. [2024]	2024	score	black	no	no	partial	none	generation	hand off
Gupta et al. [2024]	2024	score	white	no	no	partial	none	generation	hand off
Hu et al. [2024]	2024	score	black	no	no	partial	none	QA	hand off
Narasimhan et al. [2024]	2024	score	white	no	no	no	none	generation	hand off
Ong et al. [2024]	2024	score	black	no	no	partial	none	generation	hand off
Ramírez et al. [2024]	2024	typed	white	no	no	no	none	classification	hand off
Soiffer et al. [2025]	2025	free-text	black	yes	no	no	none	generation	hand off
Zellinger and Thomson [2025]	2025	typed	white	no	no	no	ECE	QA	hand off
Bouchard [2026]	2026	score	either	?	no	no	other	QA	hand off
<i>Structured-output benchmark</i>									
Tang et al. [2023]	2023	constrained	either	no	no	no	none	generation	none
Shorten et al. [2024]	2024	constrained	black	no	no	no	none	QA	none
Tam et al. [2024]	2024	constrained	either	no	no	no	none	reasoning	none
Xia et al. [2024a]	2024	constrained	black	no	no	no	none	generation	none
Geng et al. [2025]	2025	constrained	white	no	no	no	none	generation	none
Yang et al. [2025a]	2025	constrained	black	no	no	no	none	generation	none
Fan [2026]	2026	constrained	black	no	no	no	none	reasoning	none
Galeone et al. [2026]	2026	constrained	black	no	no	no	none	reasoning	none
Lee et al. [2026a]	2026	constrained	either	no	no	no	none	reasoning	none
Singh et al. [2026a]	2026	constrained	black	no	no	no	none	QA	none
<i>Studies and benchmarks of abstention</i>									
Tomani et al. [2024]	2024	free-text	either	yes	no	no	multiple	QA	none
Kirichenko et al. [2025]	2025	free-text	black	no	no	no	other	reasoning	none
Wang et al. [2026c]	2026	score	black	no	no	no	multiple	QA	abstain
<i>Typed classification</i>									
Wang et al. [2024c]	2024	typed	white	no	no	no	none	classification	none